



- A linear boundary separates two classes.

Perceptron Learning





- Randomly Initialize the weights
- For each sample:
 - Feed a sample and find the output (forward pass)
 - Find the difference between actual and desired outputs (cost function)
 - Find the effect of each weight on the cost (derivative)
 - Update the weights with a learning rate (GD)

Multi-layer Perceptron




MLP in Computer Vision



- 30x32 "Input Retina"
- 5 hidden units
- 30 output units
 - Sharp Left to Sharp Right

Convolution Layer

- $32 \times 32 \times 3$ image



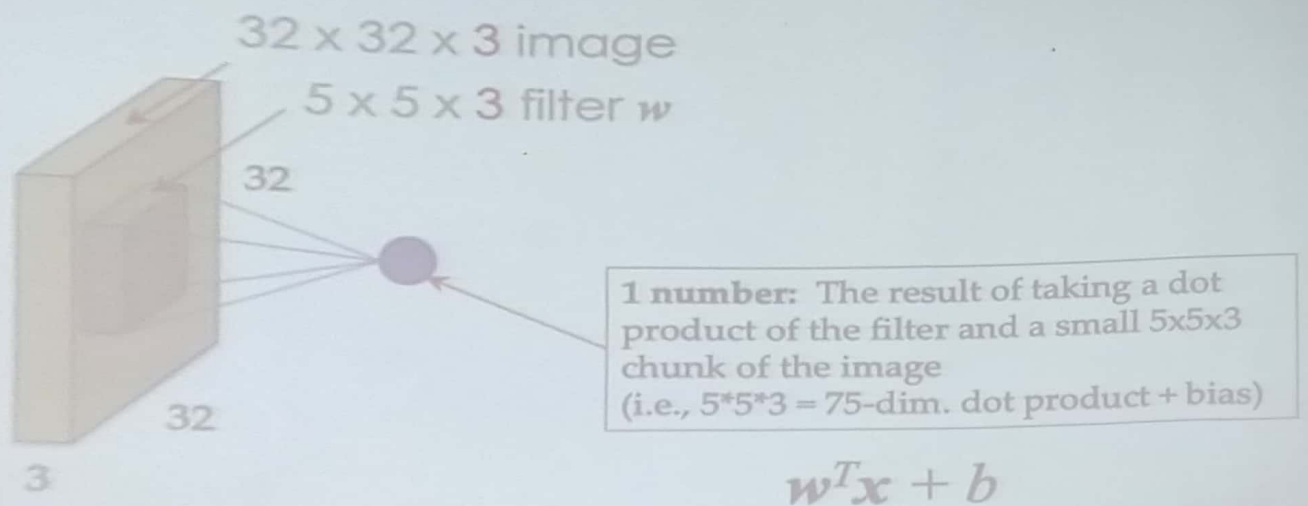
$5 \times 5 \times 3$
filter



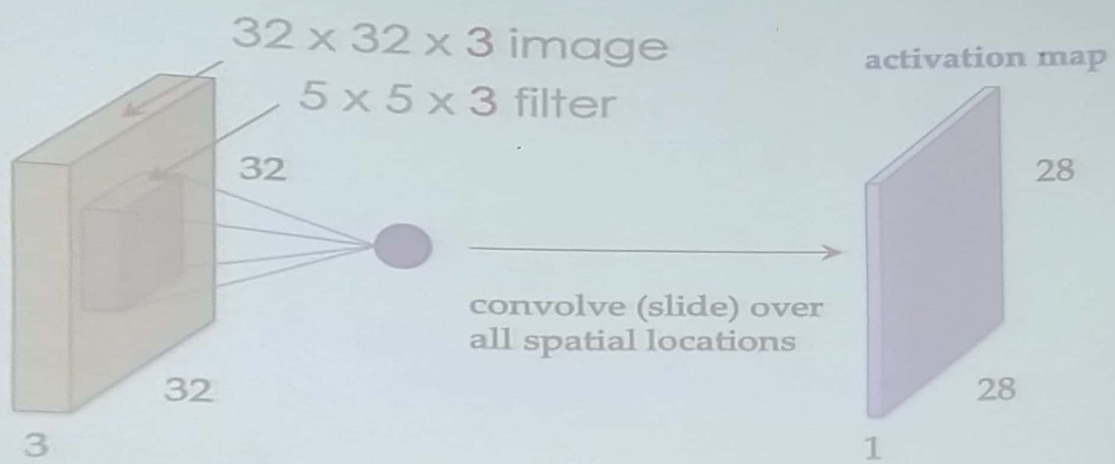
Filters always extend the full depth of the input volume

Convolve the filter with the image. i.e. "Slide over the image spatially, computing dot products"

Convolution Layer



Convolution Layer




$$f[x,y] * g[x,y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1,n_2] \cdot g[x-n_1,y-n_2]$$



Convolutional Neural Net


is a sequence of Conv. Layers, interspersed with activation functions





Convolutional Neural Net


Is a sequence of Conv. Layers, interspersed with activation functions





Convolutional Neural Net


Is a sequence of Conv. Layers, interspersed with activation functions



Understanding a Conv. Net.



Understanding a Conv. Net.



A closer look at spatial dimensions:



7

7x7 input (spatially)
assume 3x3 filter
applied with stride 3?

Doesn't fit !
Cannot apply 3x3 filter
on a 7x7 input with
stride 3

Output Dimensions:

N						
			F			
	F					

N

Output Size:
 $(N-F)/\text{stride} + 1$

e.g., $N = 7, F = 3$

- stride 1: $(7-3)/1 + 1 = 5$
- stride 2: $(7-3)/2 + 1 = 3$
- stride 3: $(7-3)/3 + 1 = 2.33!$

Common to Zero-pad the border in practice

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

e.g., input 7x7
3x3 filter applied with stride 1
pad with 1 pixel border

What is the output size?

Size = 7x7

Note: output Size:
 $(N-F)/\text{stride} + 1$

Common to Zero-pad the border in practice

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Output Size = 7x7

In general, it is common to have conv layers with **stride 1**, **filter size $F \times F$** , and **zero padding $(F-1)/2$** , preserving spatial size


- $F=3 \rightarrow$ zero pad with 1
- $F=5 \rightarrow$ zero pad with 2
- $F=7 \rightarrow$ zero pad with 3

Example:


- Input Volume: **32 x 32 x 3**
- 10 **5x5** filters with stride 1, pad 2
- Number of parameters in this layer?

each filter has $5*5*3 + 1 = \mathbf{76 \text{ params}}$ (1 for bias)

→ $\mathbf{76*10 = 760 \text{ parameters}}$ in the layer




Note: 1x1 convolutions are perfectly fine




1x1 CONV
with 32 filters

(each filter has size
1x1x64, and performs
a 64-dimensional dot
product)




Pooling Layer

- Makes representations smaller and manageable
- Later filters have larger support
- Operates over each activation map independently



Max Pooling (2D)



Summary

- CNNs are a series of CONV, ReLU, Pool, FC layers
- CNNs are computationally efficient and compact
- Parallels to human/animal visual system.
- Learnt features can be used for classification
- Recent Trends:
 - Stick with 3x3 filters, make the network deeper
 - Improve connectivity
 - Several innovations for specific applications