



CS131

Tracking people

A. Alahi
November 19th 2014



A. Alahi

Topic of the day



Understanding human behavior

Motivation: Elderly monitoring



Courtesy of Ph.D Guido Pusiol



A. Alahi

Motivation: Path-to-purchase understanding



Motivation: Space analytics



Interaction

- Number of social interaction



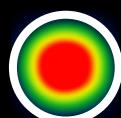
Distances

- Walked distance



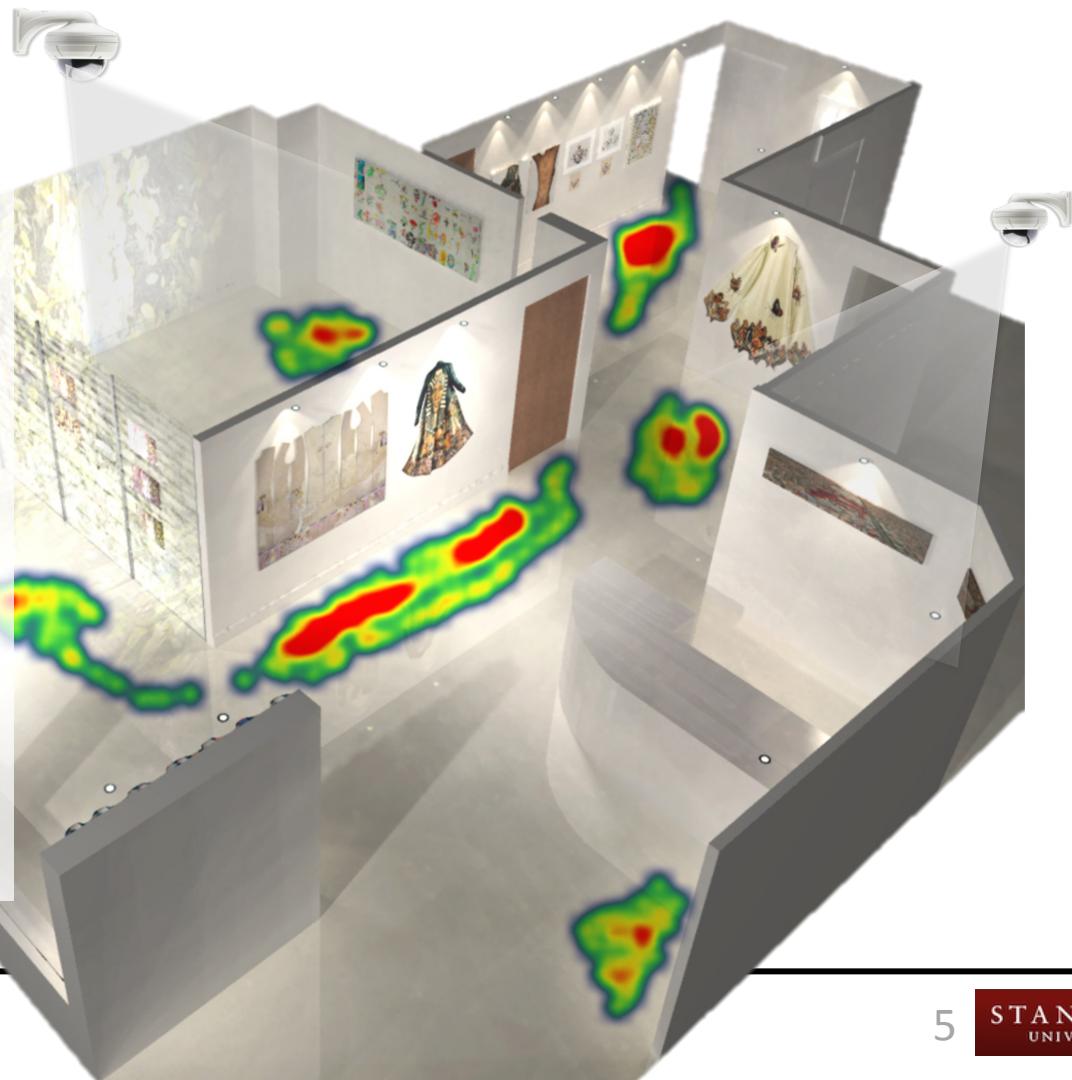
Durations

- Duration in each room



Heatmaps

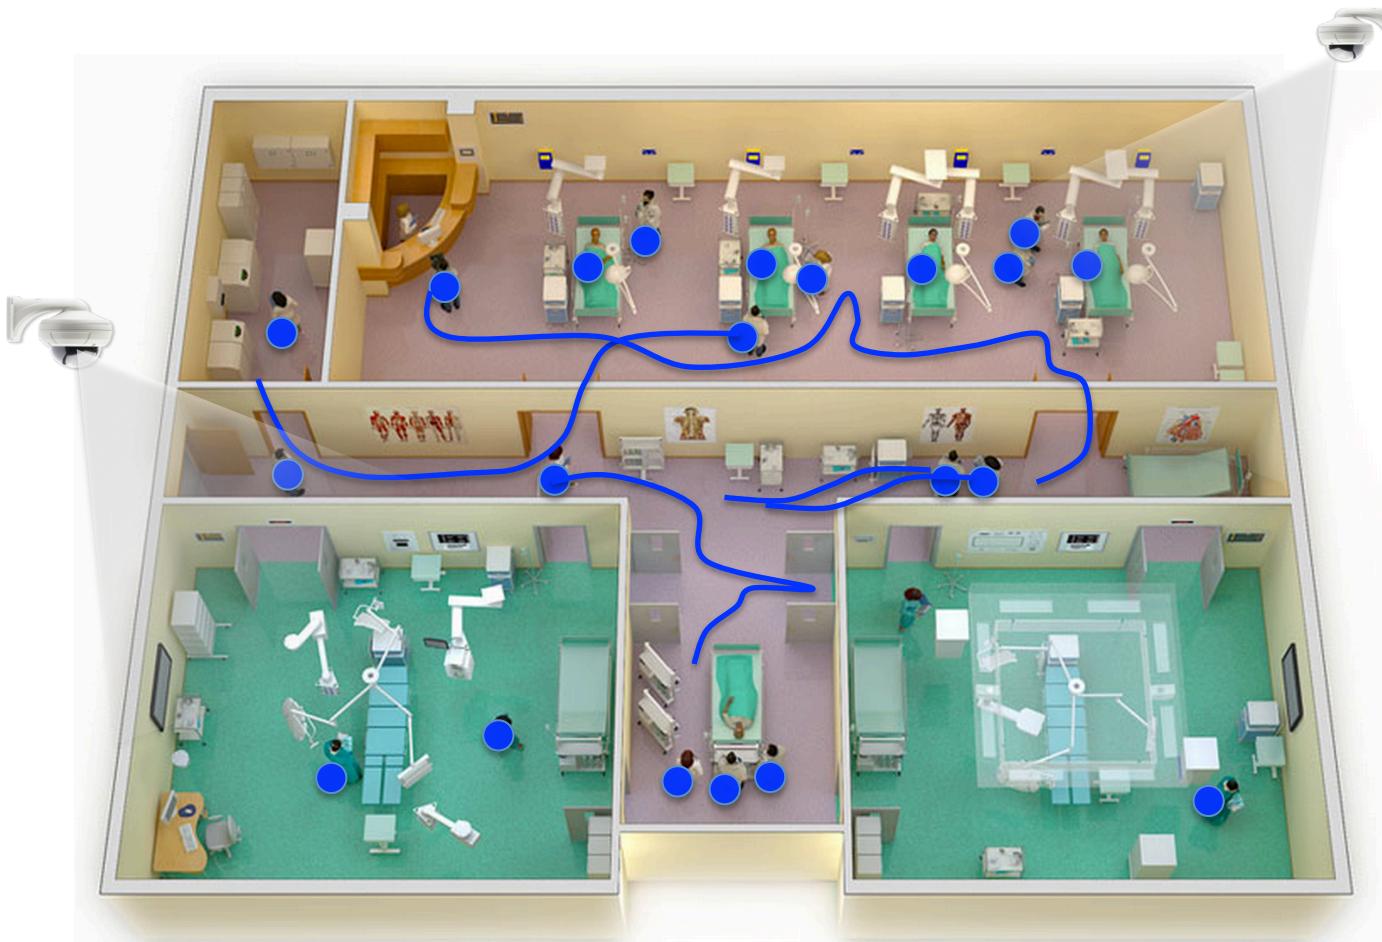
- Hot spots



Motivation: Performance analysis



Motivation: Behavior monitoring



Motivation: Large-space analytics



Motivation: Large-space analytics



- Number of visitors



- Path of visitors



- Duration



- Rank spots



Motivation: Large-space analytics



- Number of visitors



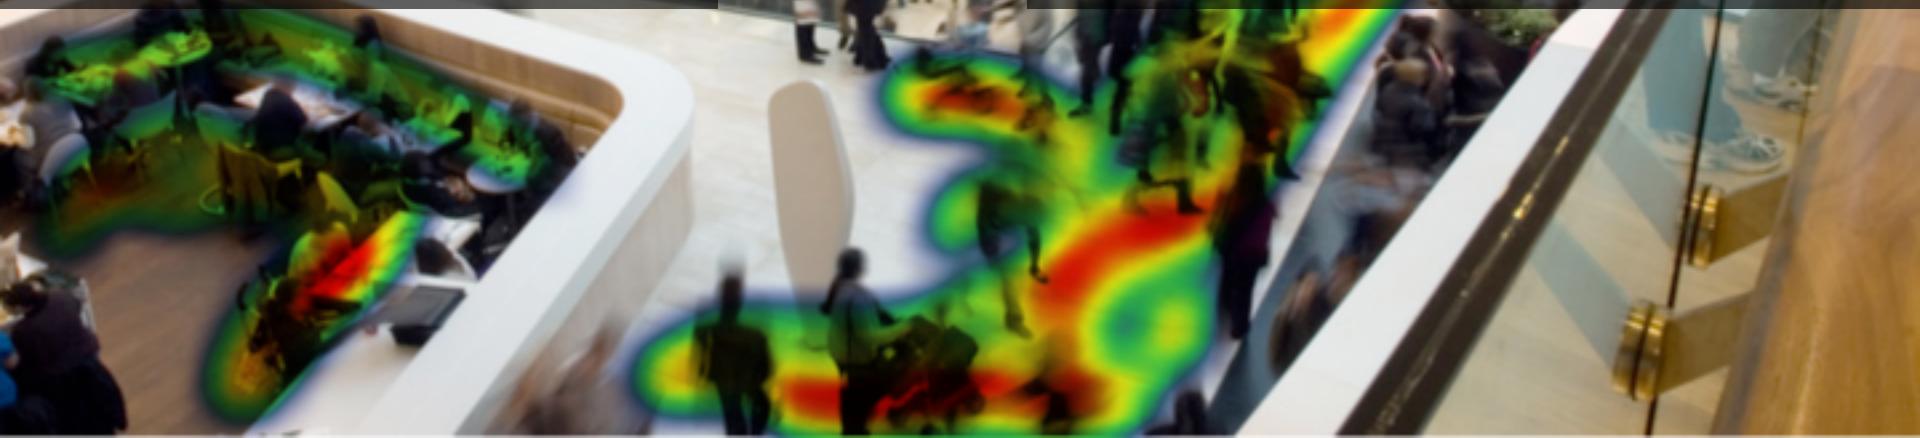
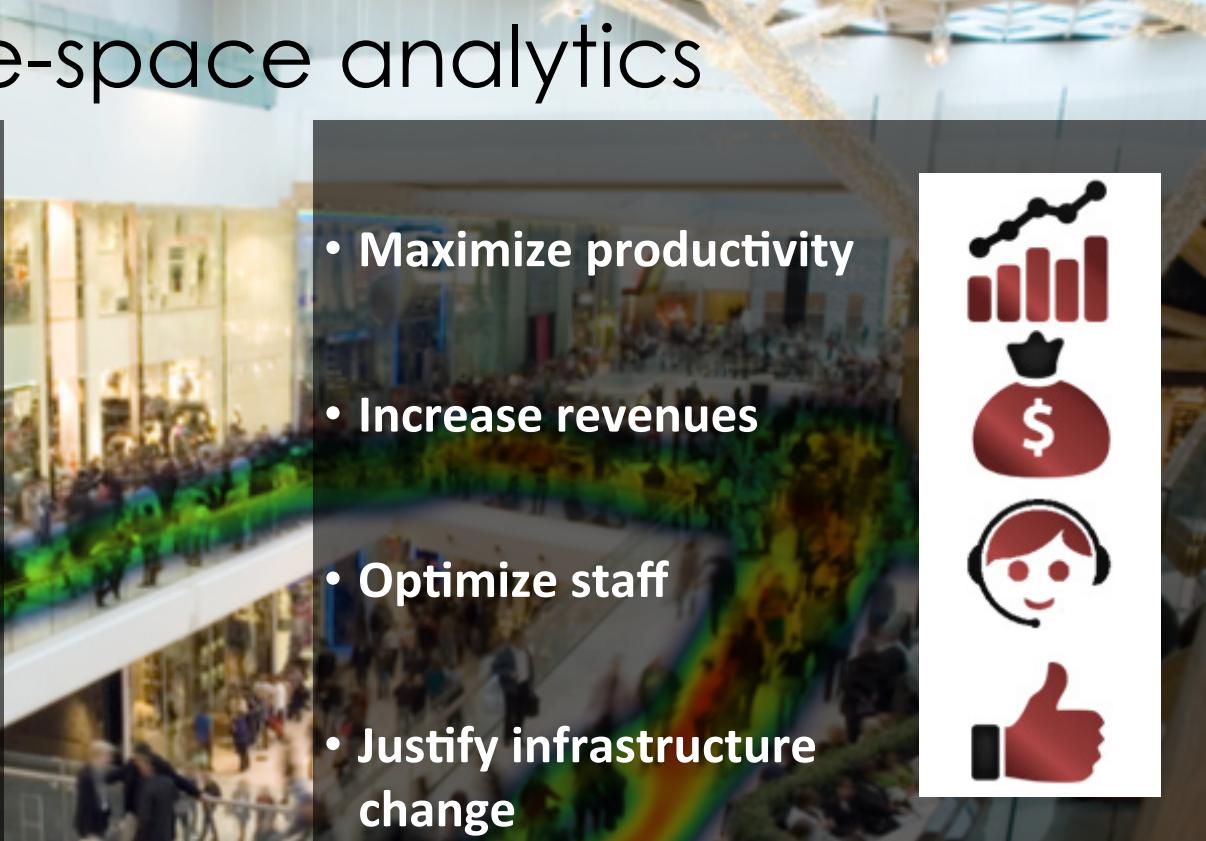
- Path of visitors



- Duration



- Rank spots



Understand human mobility in a large terminal over a year

18.05	SBZ	Thalwil	Zug	Lucern	8	18.12	18.18	Stadelhofen	Winterthur	Seuzach	28	18.24	18.32	Enge	Thalwil	Kreuz Säntis	Wetzikon	01
18.06	SBZ	Büren Brugg	Bremgarten	Lucern	94	18.13	18.19	Wetkiss Egg	Thalwil	Zug	91	18.25	18.33	Stadelhofen	Effretikon	Wetzen	20/24	23/24
18.07	SBZ	Burgdorf	Landquart	Ober	7	18.14	18.20	Schau	Glessenhögl	Sihlwald	1	18.26	18.34	Hardbrücke	Affoltern a/A	Zug	21/22	21/22
18.08	SBZ	Flughafen	Wohlen	Wohlen	12	18.15	18.21	Oerlikon	Wohlen	Winterthur	22	18.27	18.35	Sehneu	Tremli	Wetzen	Uetliberg	2
18.09	SBZ	Flughafen	Winterthur	Wohlen	9	18.16	18.22	Herdbrücke	Altstetten a/A	Winterthur	59	18.28	18.36	Oerlikon	Wohlen	Schaffhausen	Romanshorn	31/22
18.10	SBZ	Arlesheim	Feldbach	Wohlen	4	18.17	18.23	Stadelhofen	Wohlen	Plattikon	82	18.29	18.37	Stadelhofen	Winterthur	Ziegelbrücke	Langnau-G.	31/23
18.11	SBZ	Olten	Dielsdorf	Wohlen	50	18.18	18.24	Wetkiss Egg	Wohlen		23/34	18.30	18.38	Sehneu	Glessenhögl	Muri Rothreus	1	20/24
18.12	SBZ	Basel	Frankfurt	Wohlen	10	18.19	18.25	Wetkiss Egg	Wohlen			18.31	18.39	Wetkiss Egg	Thalwil			10
18.01	AFR	Doriken	Flughafen	Wohlen	4	18.20	18.26	Stadelhofen	Wohlen			18.32	18.40	Altstetten	Dielsikon			
18.01	AFR	Zug	Lucerne	Wohlen	5	18.21	18.27	Wetkiss Egg	Wohlen			18.33	18.41					

>120,000 commuters/day in 2013

>240,000 commuters/day by 2030*



Lausanne, Switzerland

* SBB CFF FFS report leman 2030

Track crowds in large spaces

Analytics

Flow in/out

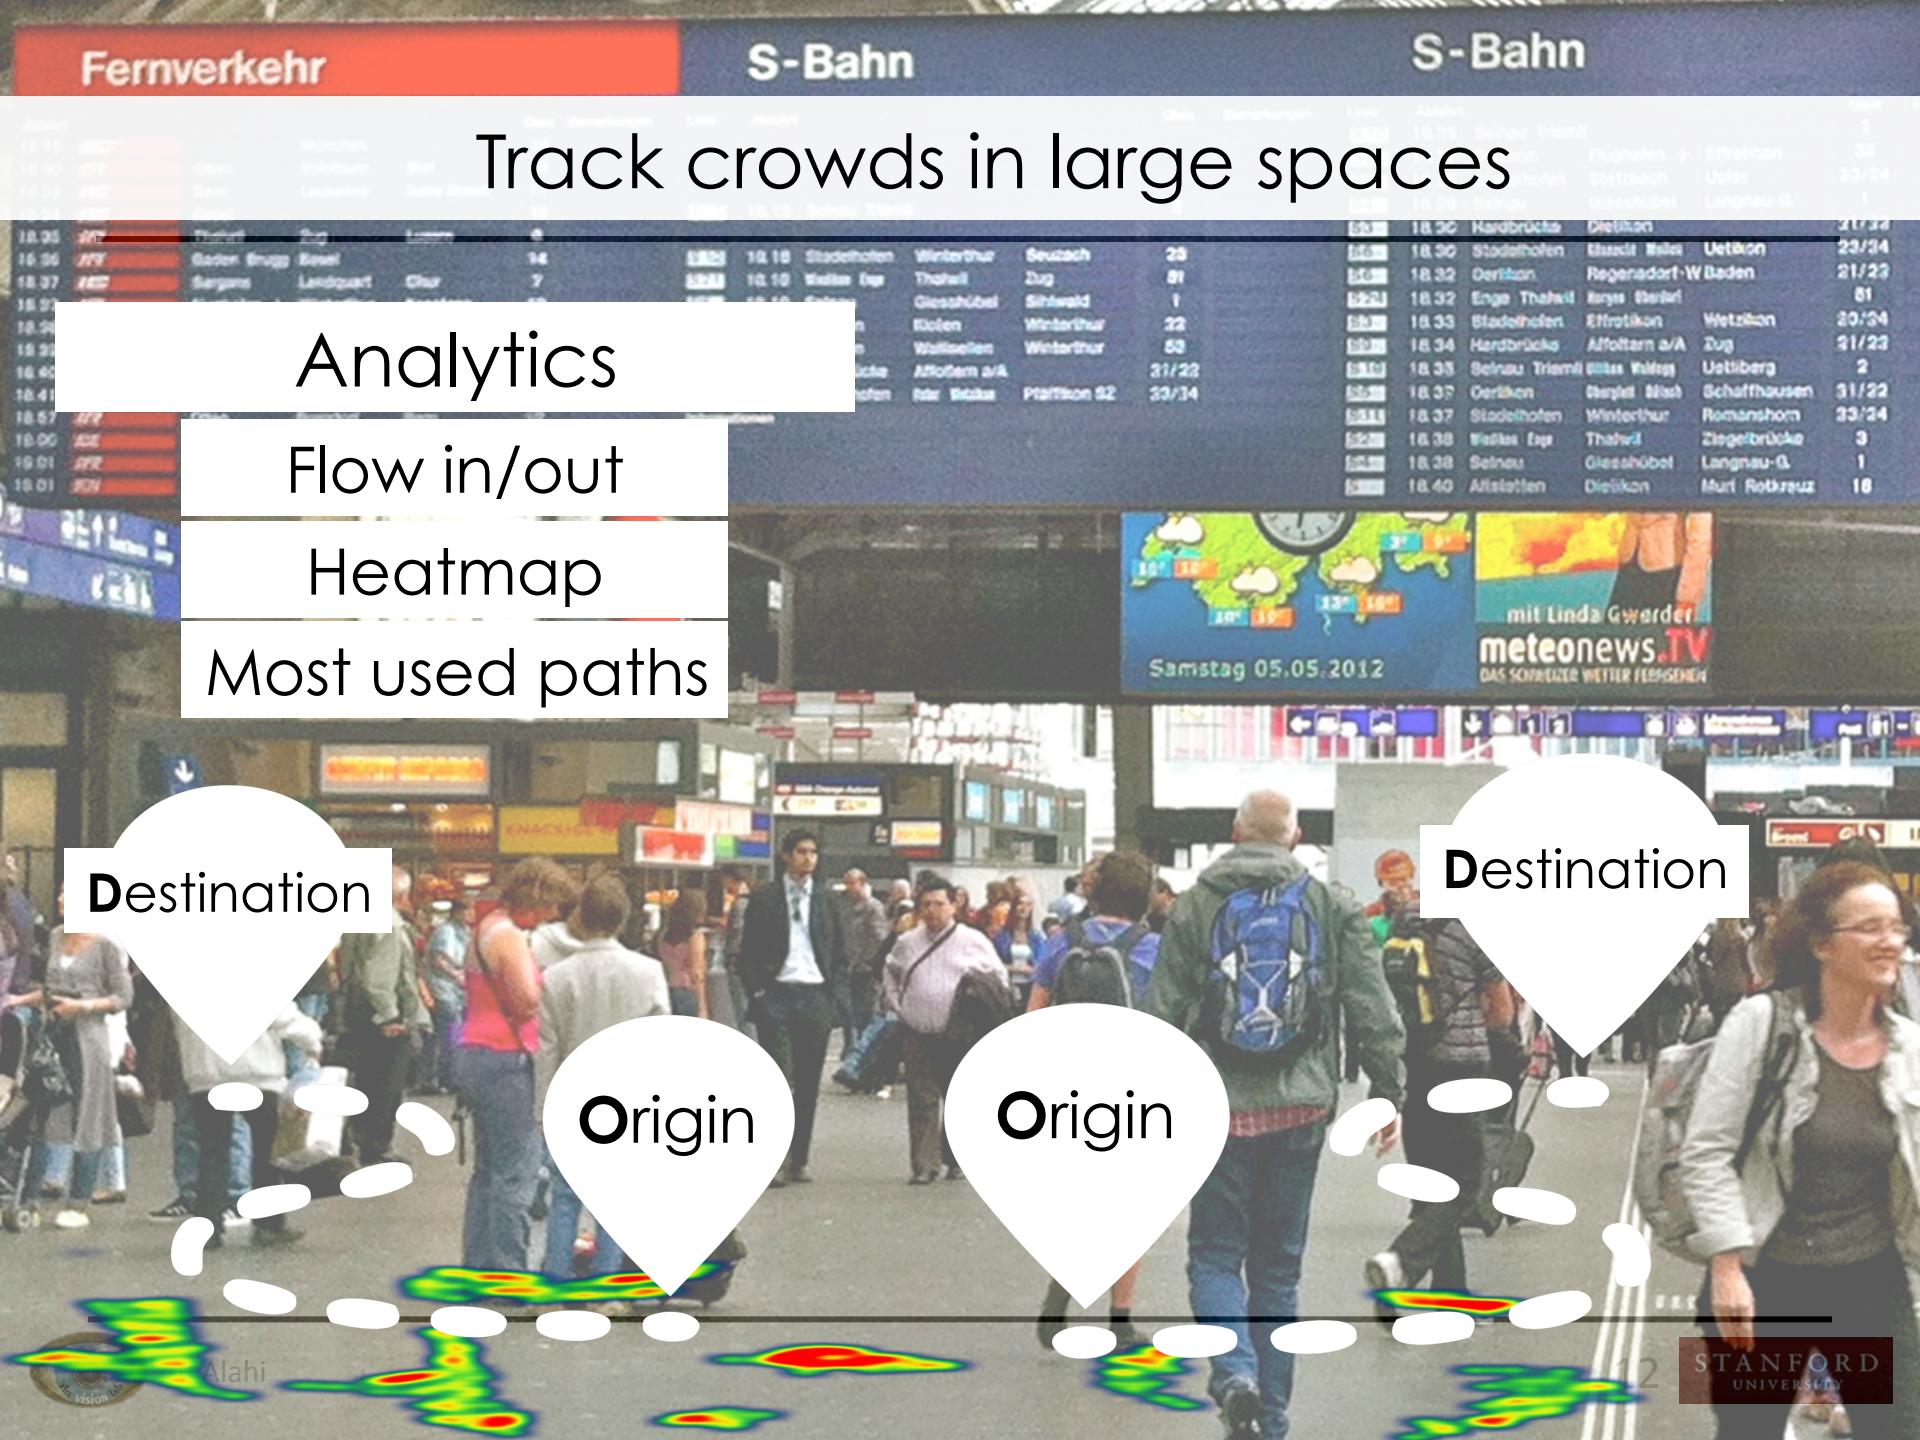
Heatmap

Most used paths

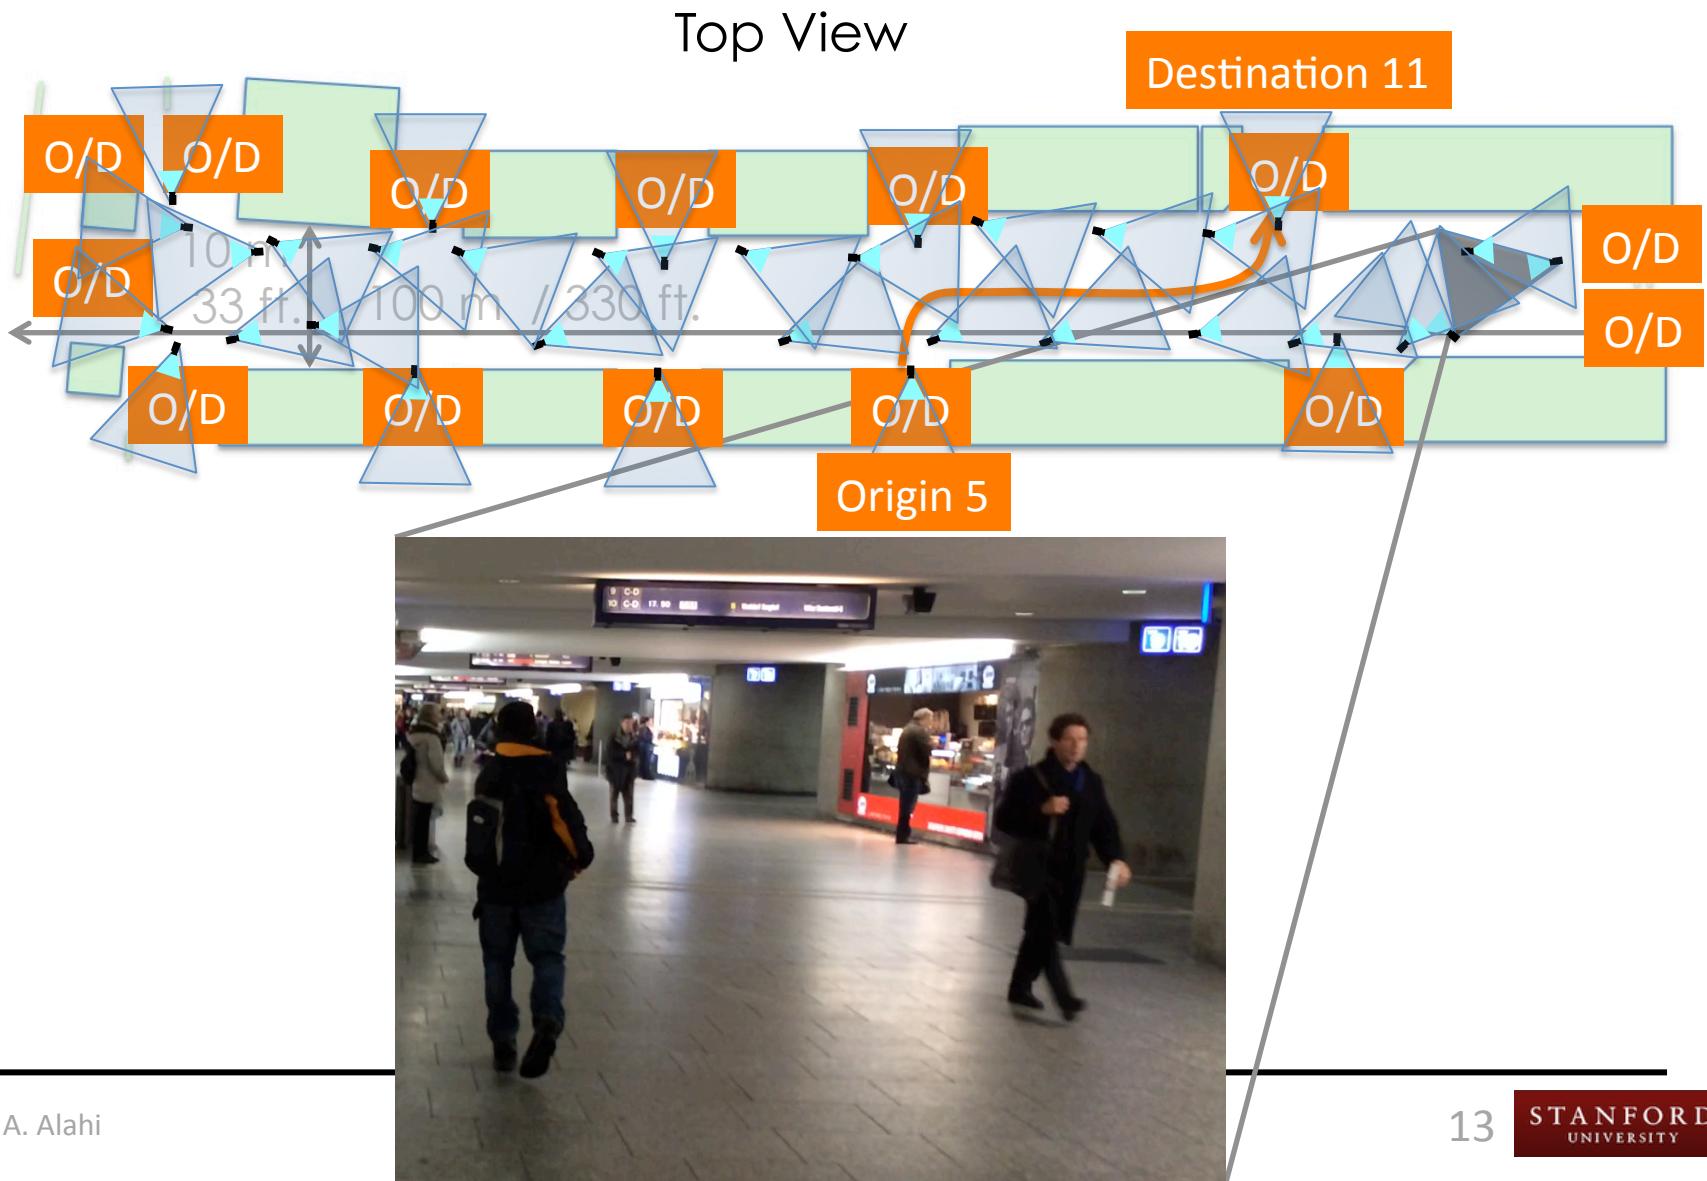
Destination

Origin

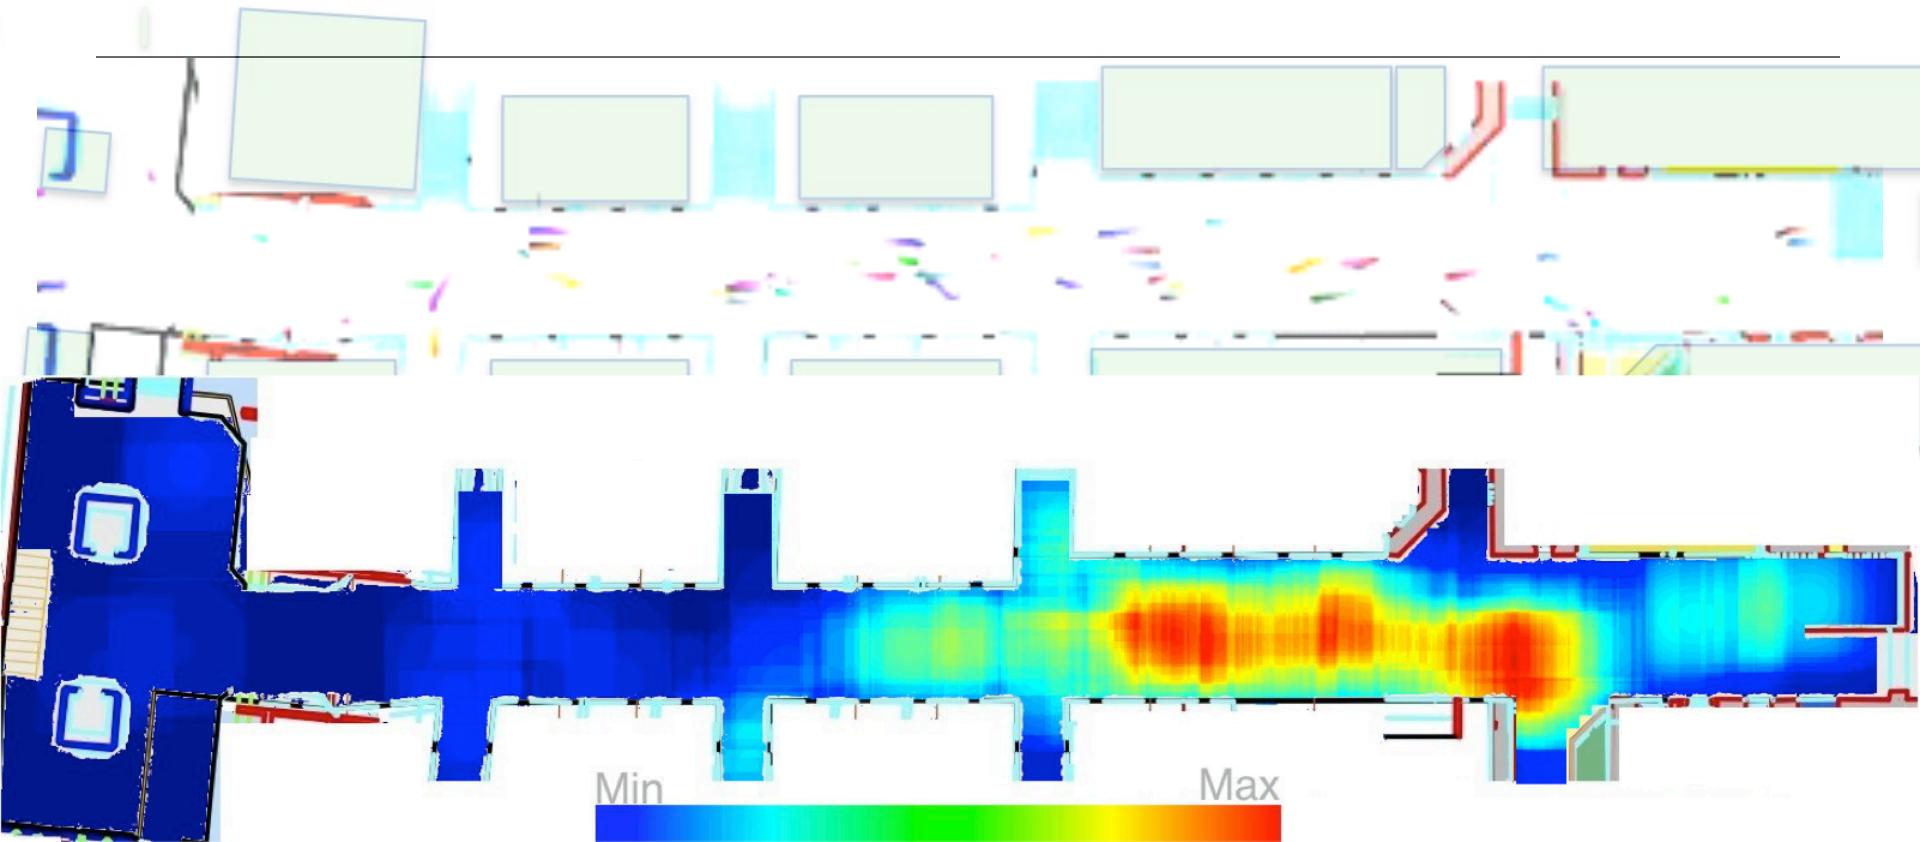
Destination



A corridor with 14 Origin/Destination (O/D)



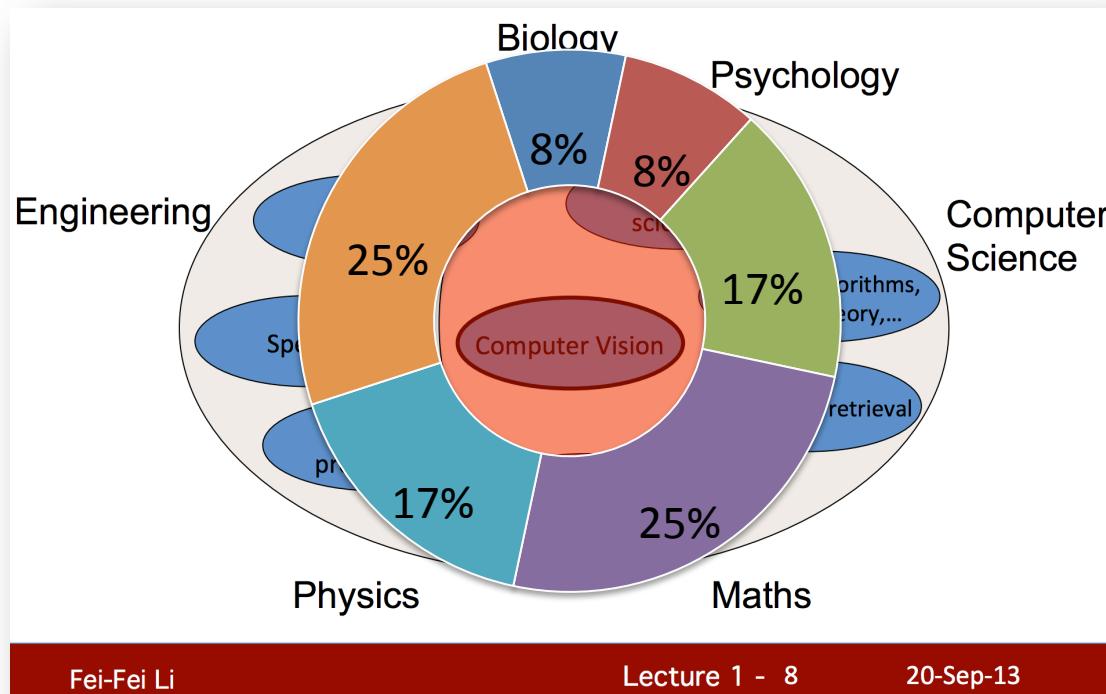
Collect long-term trajectories



# People	Av. duration	Av distance	Density (up to)	# Paths (O/D)
42 million	1 minute	100m	1 pedestrian/m ²	196

Tracking people

- Thousands scientific publications about tracking people
- What is related to tracking people?



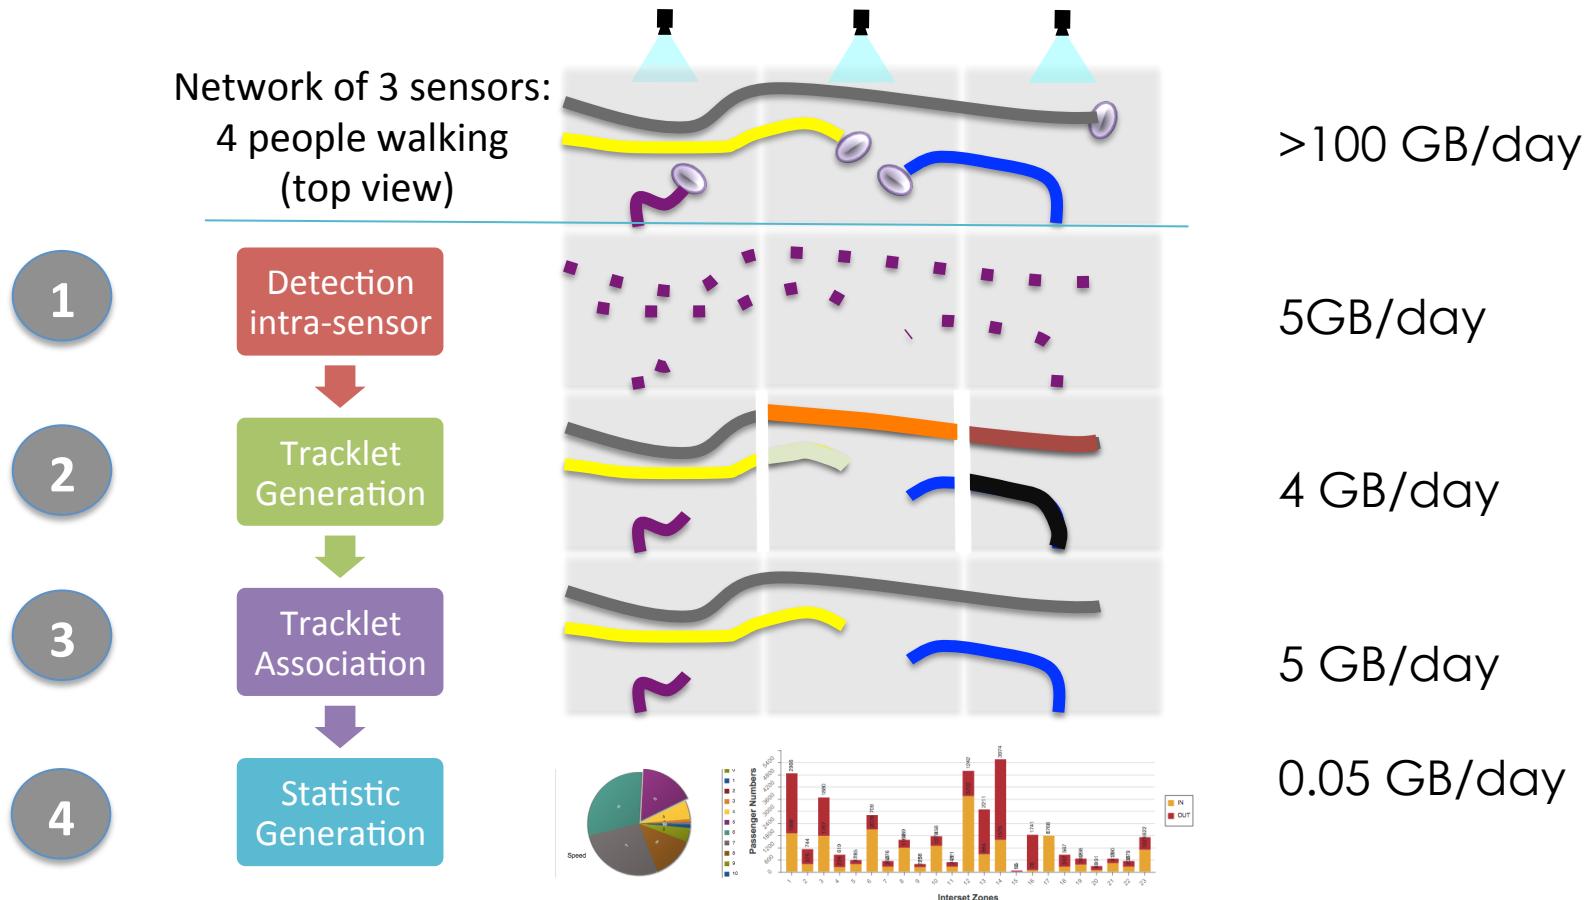
Tracking people

- Focus of today's lecture:
 - Tracking people with
 - Video streams (vs static images)
 - ⇒ Prior from the scene
 - ⇒ Calibration data
 - ⇒ Mapping from image plane to real-world
 - ⇒ Model temporal variation/changes



Outline:

From Foreground Extraction To Tracking 42 million Pedestrians



Outline:

From Foreground Extraction **To** Tracking 42 million Pedestrians

I. Detection

- I. Foreground extraction
- II. Pedestrian localization

II. Tracklet Generation

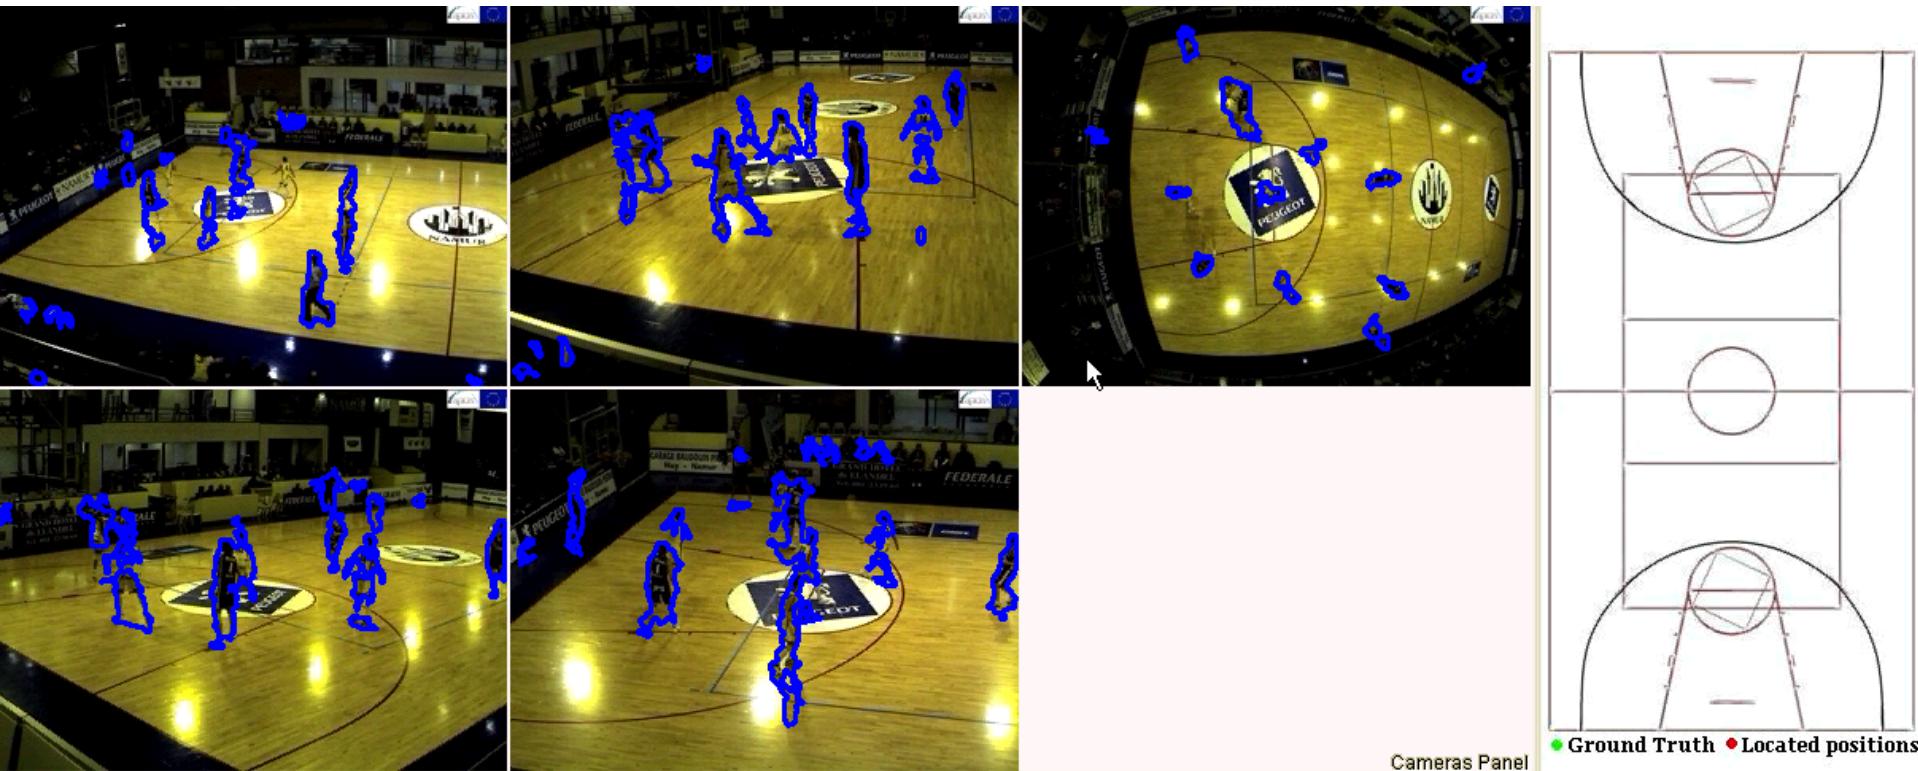
- I. Data association problem
- II. Matching appearance cues

III. Tracklet Association

- I. Modeling Social Affinities



I. Detection: Foreground extraction



- Severely degraded foreground silhouettes
- Spatially dense distribution
- Strong occlusions

I. Detection: Foreground extraction

$$F(x,y,t+1) = \begin{cases} 1 & \text{if } |I(x,y,t) - B(x,y)| > T \\ 0 & \text{otherwise} \end{cases}$$

- Frame differencing

$$B(x,y) = I(x,y,t-1)$$

- Mean filter

$$B(x,y) = 1/N \sum_{i=1 \dots N} I(x,y,t-i)$$

- Gaussian averaging

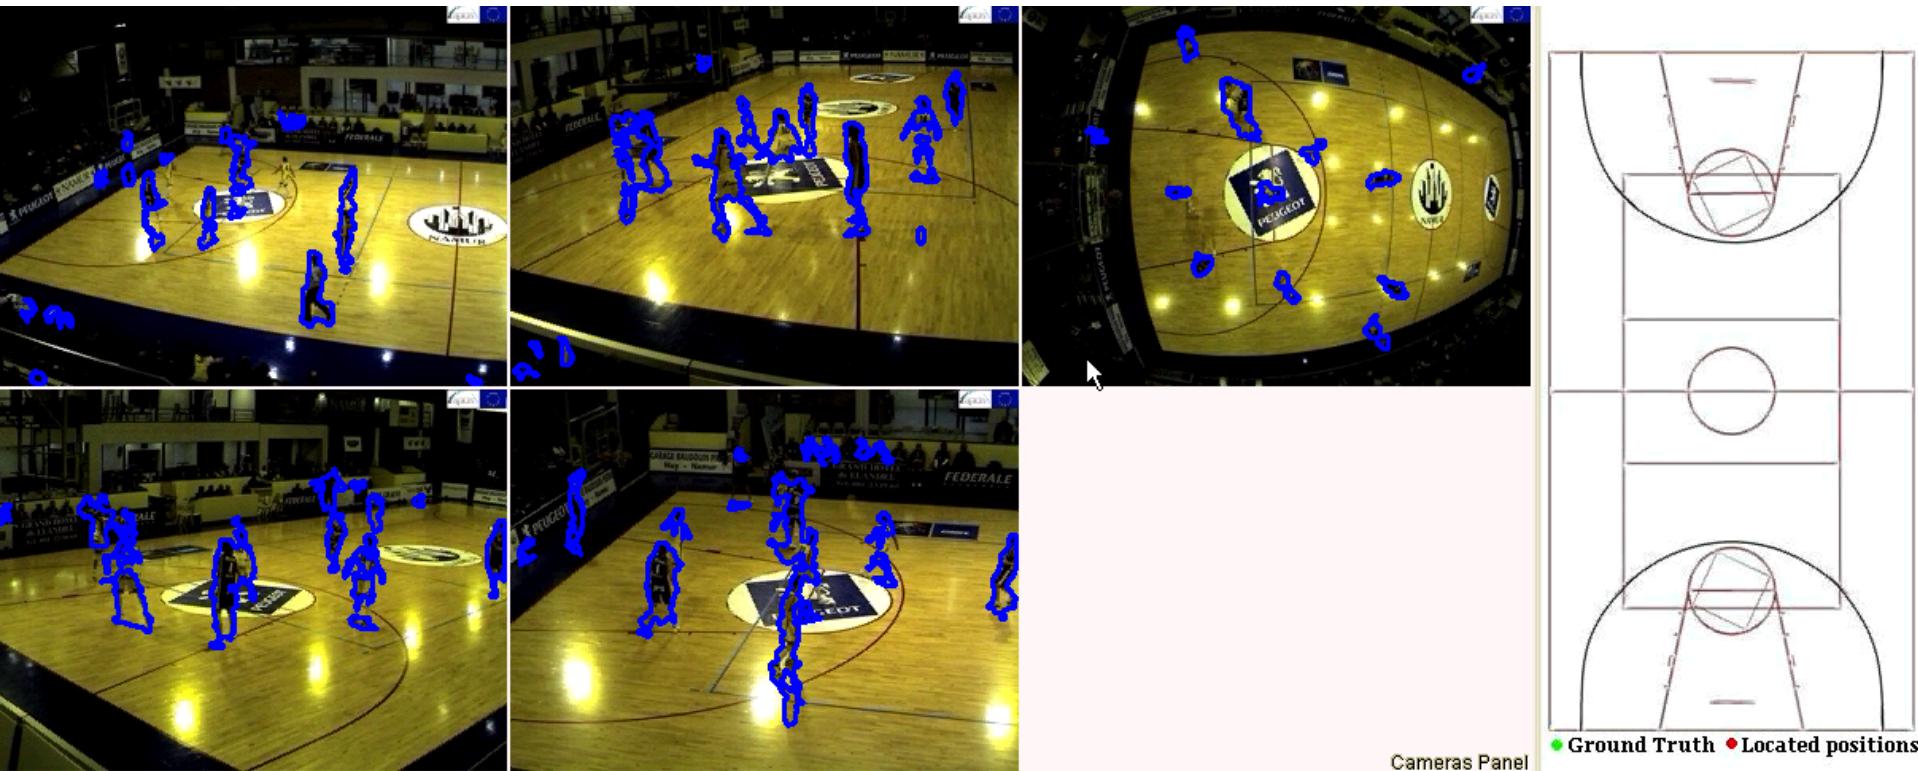
- GMM

- ...

=> Library of 32 algorithms (*BGS library*)

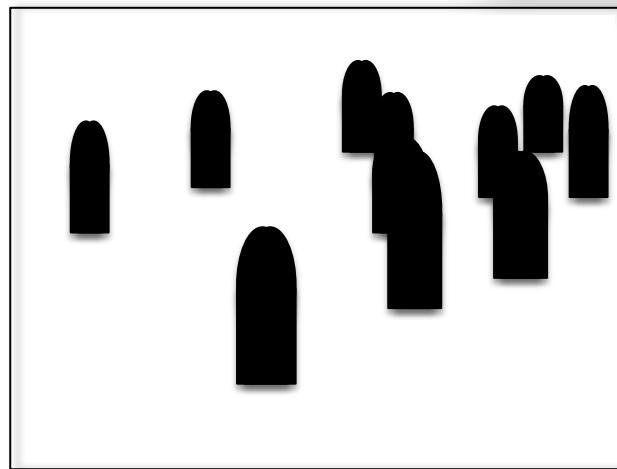


I. Detection: Foreground extraction



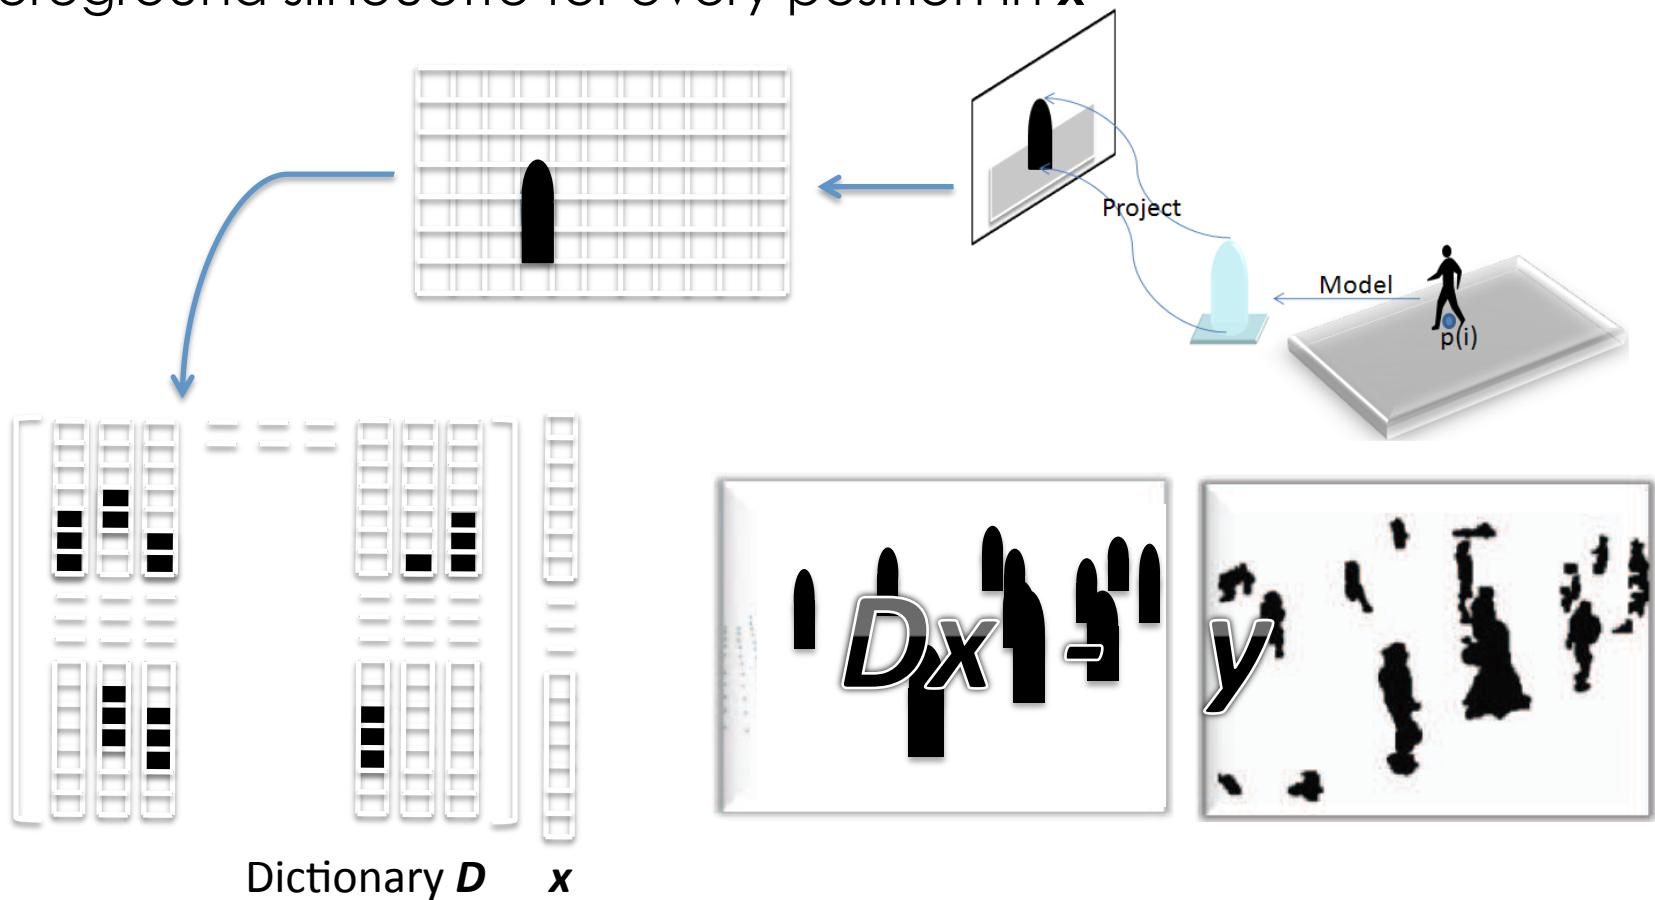
- Severely degraded foreground silhouettes
- Spatially dense distribution
- Strong occlusions

I. Detection: Pedestrian localization



I. Detection: Calibrated Camera

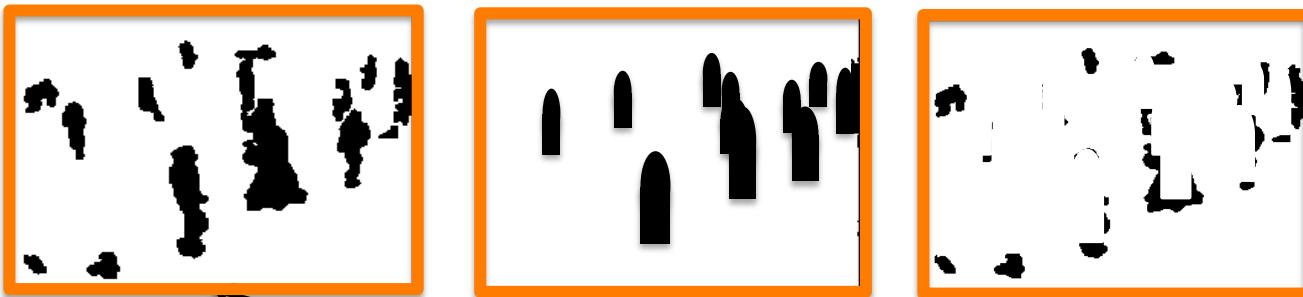
- Create a dictionary D of atoms approximating the ideal foreground silhouette for every position in x



I. Detection: Sparsity driven framework

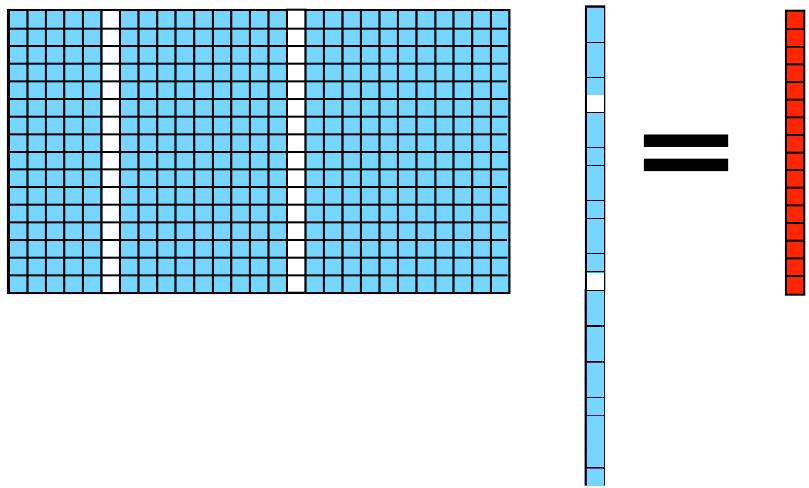
- Inverse problem:

$$y = Dx + n$$



I. Detection: Greedy approach

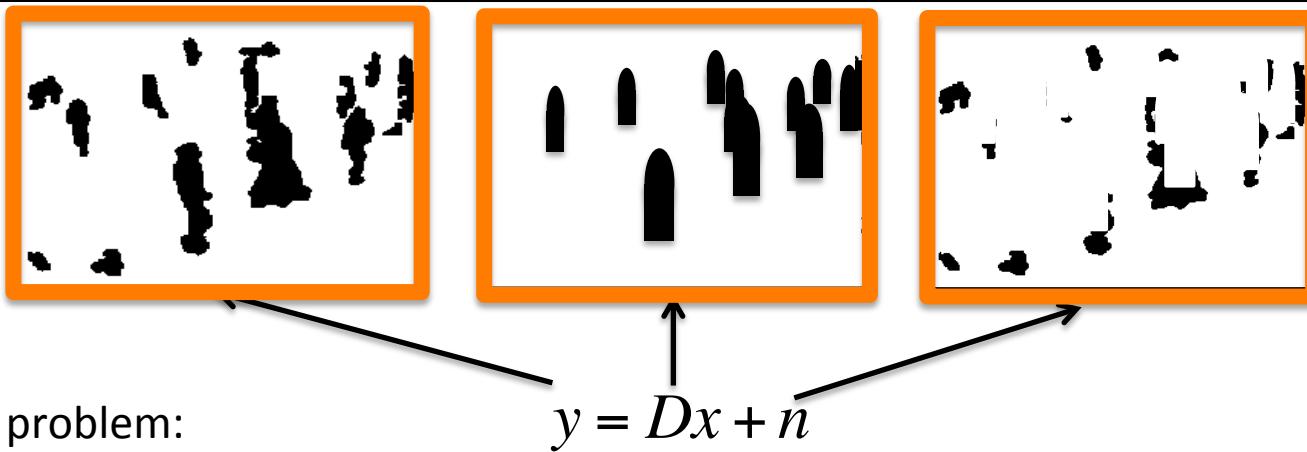
$$D \quad x = y$$



[1] S. Mallat and Z. Zhang, "Matching pursuits with time-frequency dictionaires," IEEE Transactions on signal processing, 1993.



I. Detection: Sparsity driven framework



- Inverse problem:

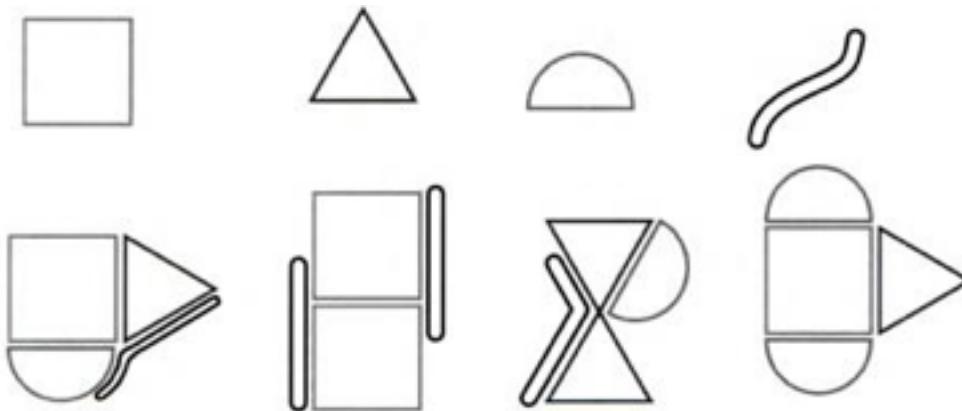
$$y = Dx + n$$

- Sparsity prior:

$$\min \|x\|_0 \quad \text{s. t.} \quad y = Dx + n$$

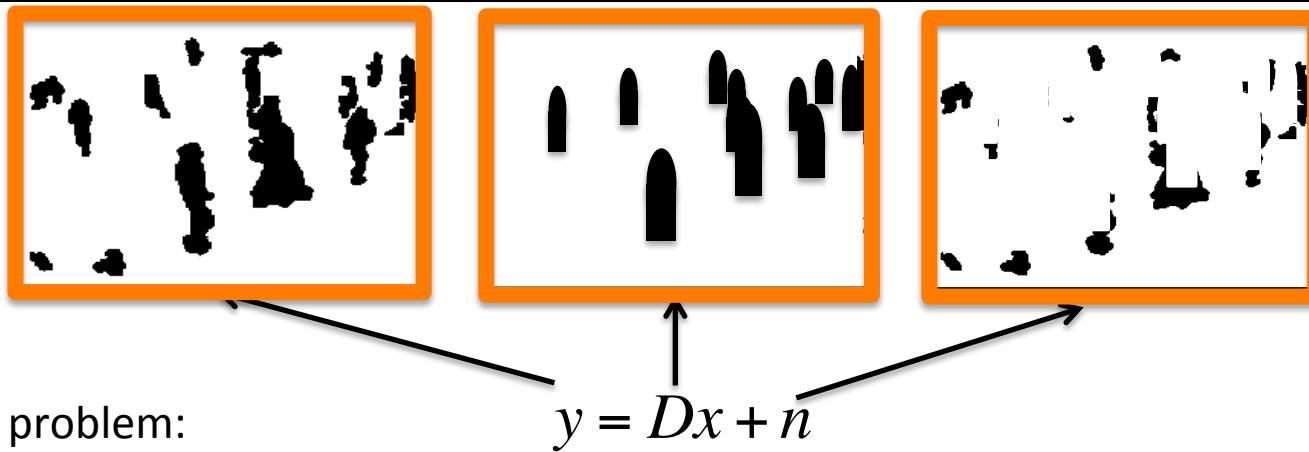
I. Detection: In praise of sparsity

“Creation is based on small number of primary, indivisible elements that combine with one another according to a few simple patterns.” [1]



[1] B. Elahi, “Spirituality is a science: Foundation of Natural Spirituality”

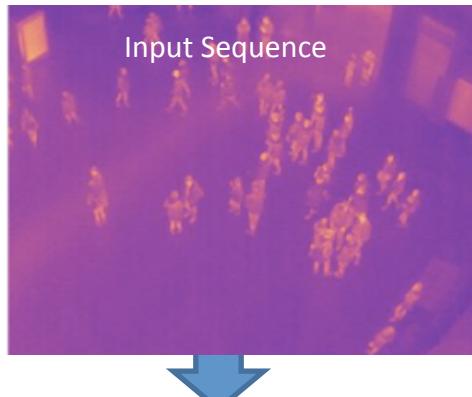
I. Detection: Sparsity driven framework



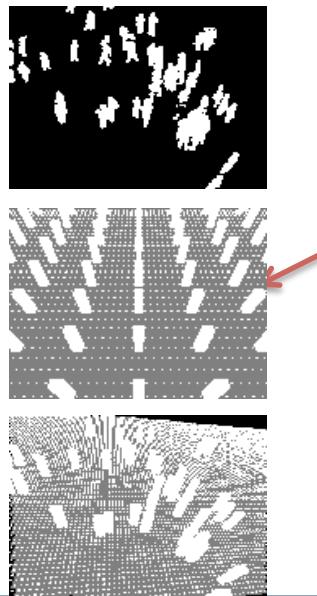
- Inverse problem:
- Sparsity prior: $\min \|x\|_0 \quad \text{s. t.} \quad y = Dx + n$
- Basis Pursuit [1]: $\min \|x\|_1 \quad \text{s. t.} \quad y = Dx$
- BPDN: $\min \|x\|_1 \quad \text{s. t.} \quad \|y - Dx\| \leq n$
- Lasso: $\min \|y - Dx\|_2 \quad \text{s. t.} \quad \|x\|_1 \leq \varepsilon$

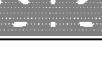


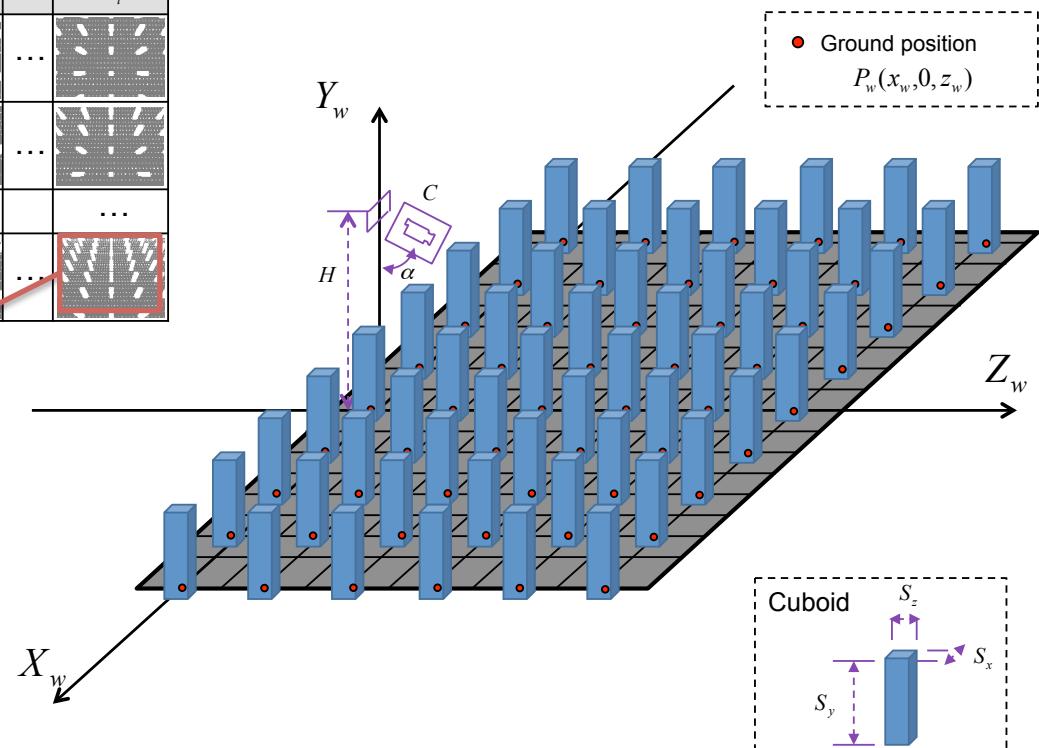
I. Detection: Pedestrian localization



- Foreground extraction
- Used Dictionary
- Localization

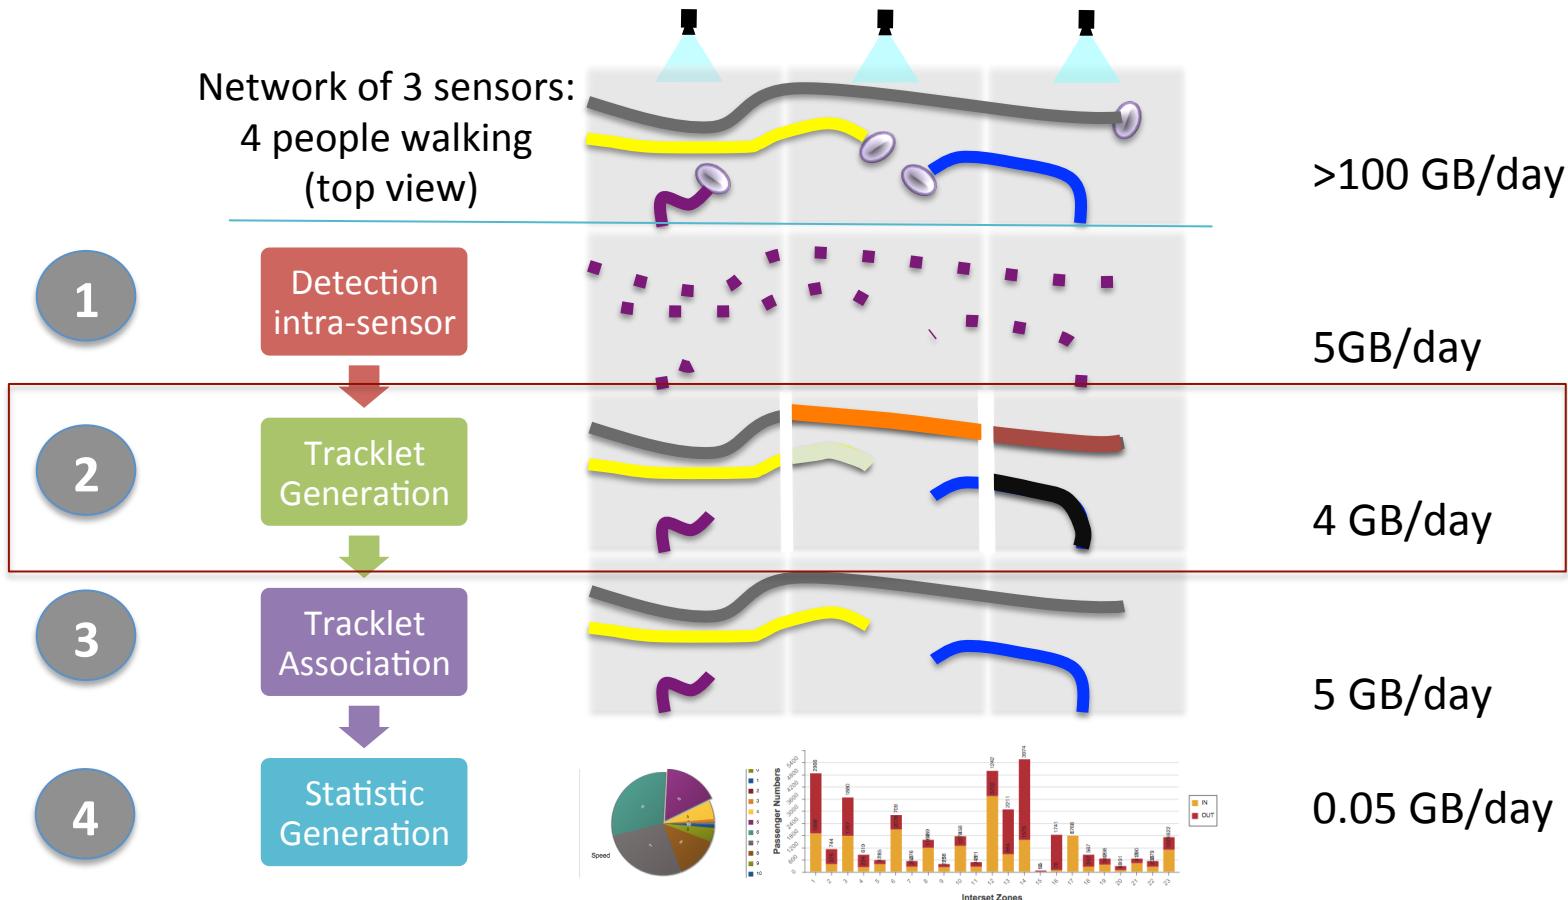


	H_1		H_l
α_1		...	
α_2		...	
α_k		...	



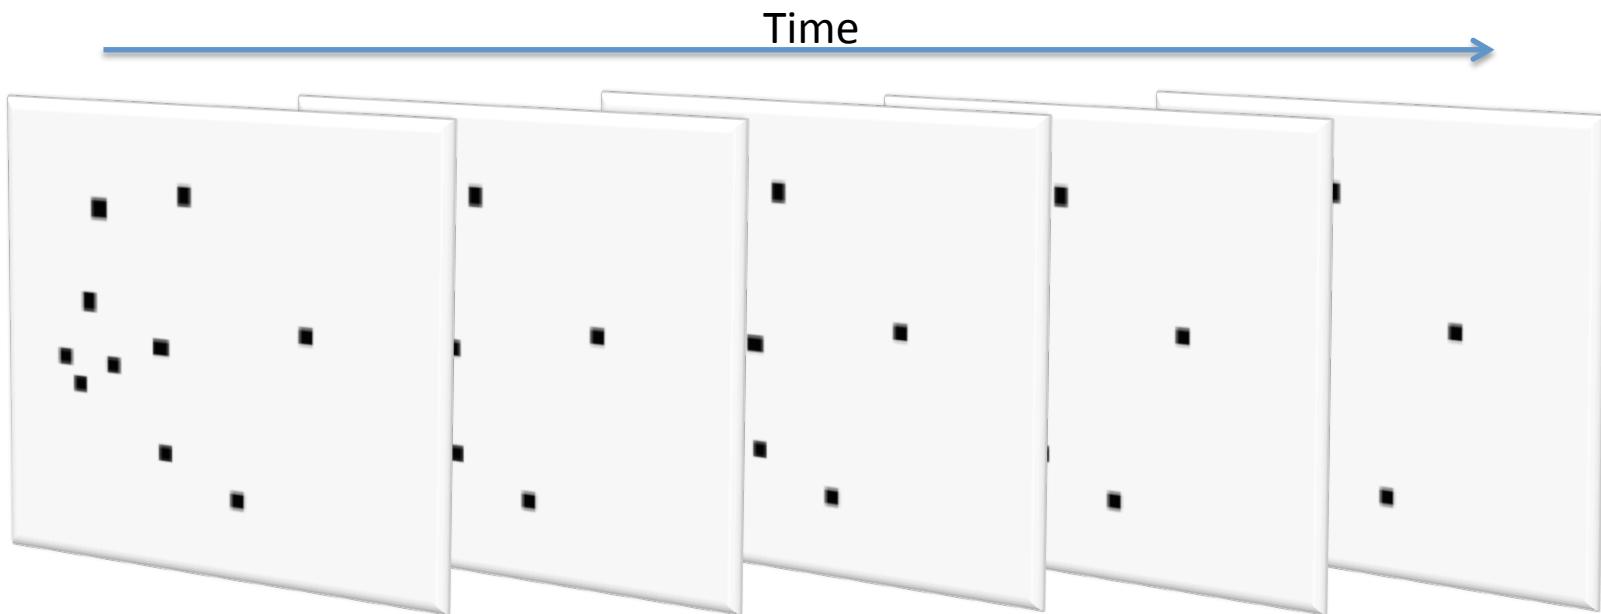
Outline:

From Foreground Extraction To Tracking 42 million Pedestrians

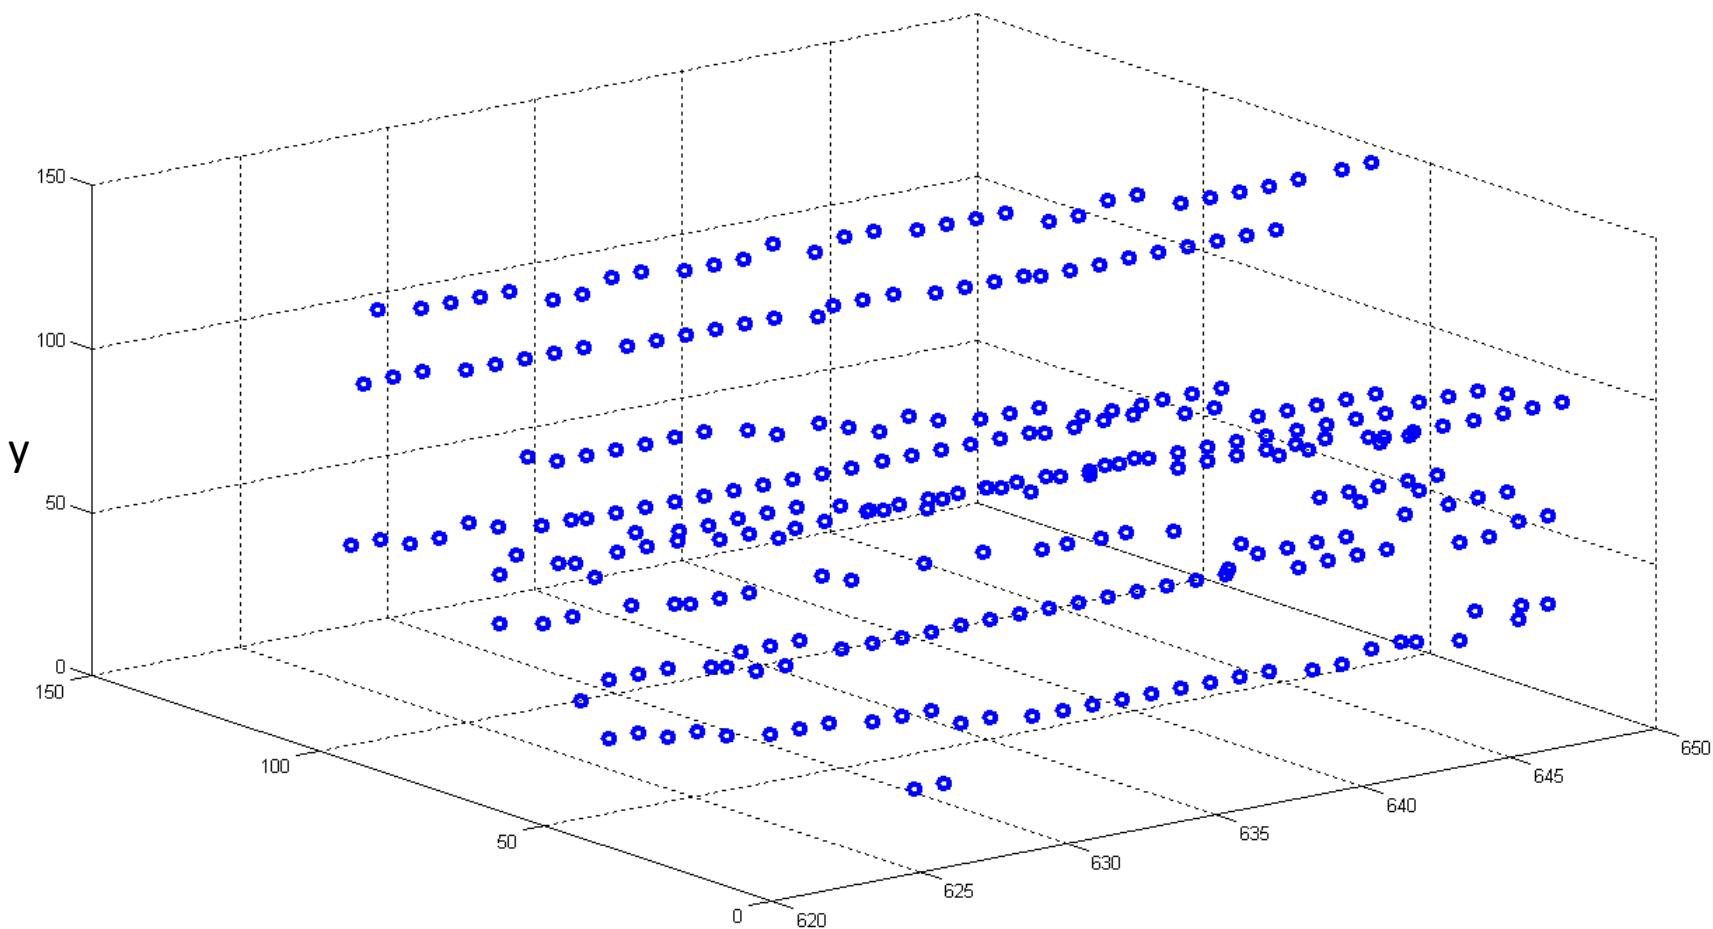


II. Tracklet generation: Data Association Problem

- Create a Directed Acyclic Graph $G = (N, E)$ where
 - N = The detected ground plane points across time
 - E = The connectivity cost between the detections (based on motion/appearance model)



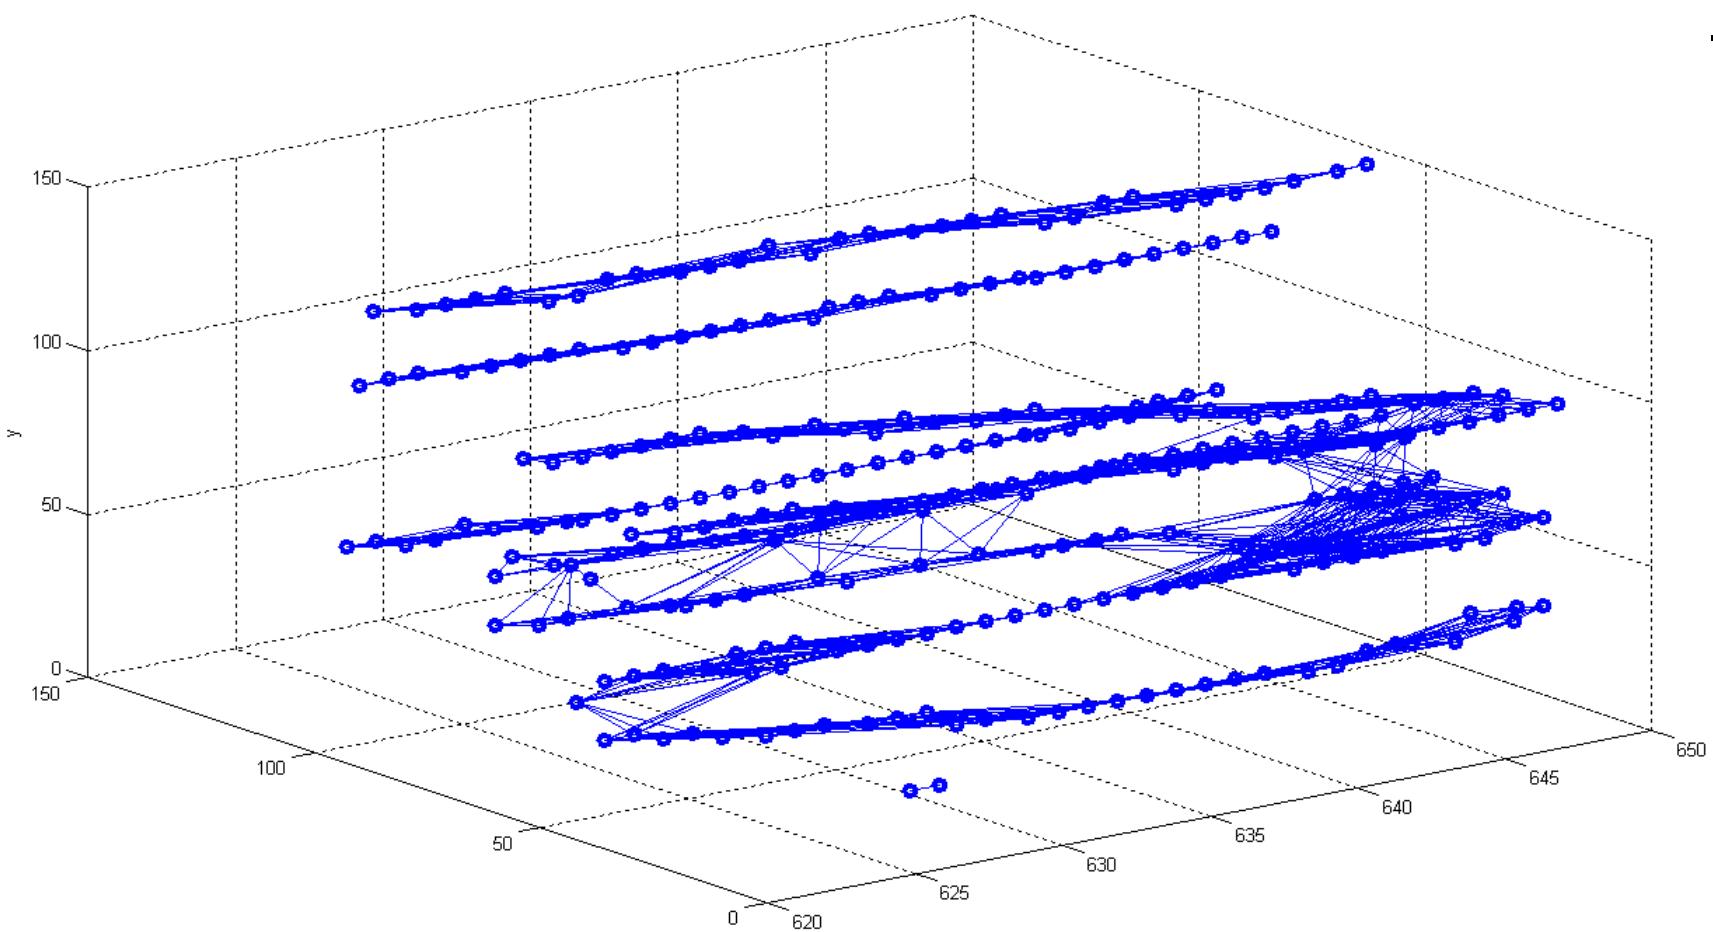
II. Tracklet generation: Data Association Problem



- [1] A.Alahi, L.Jacques, Y.Boursier, and P.Vandergheynst, "Sparsity-Driven People Localization with a Heterogeneous Network of Cameras," Journal of Mathematical Imaging and Vision, 2011
- [2] F. Fleuret, J. Berclaz, R. Lengagne, and P. Fua, "Multi-Camera People Tracking With a Probabilistic Occupancy Map", PAMI 2008



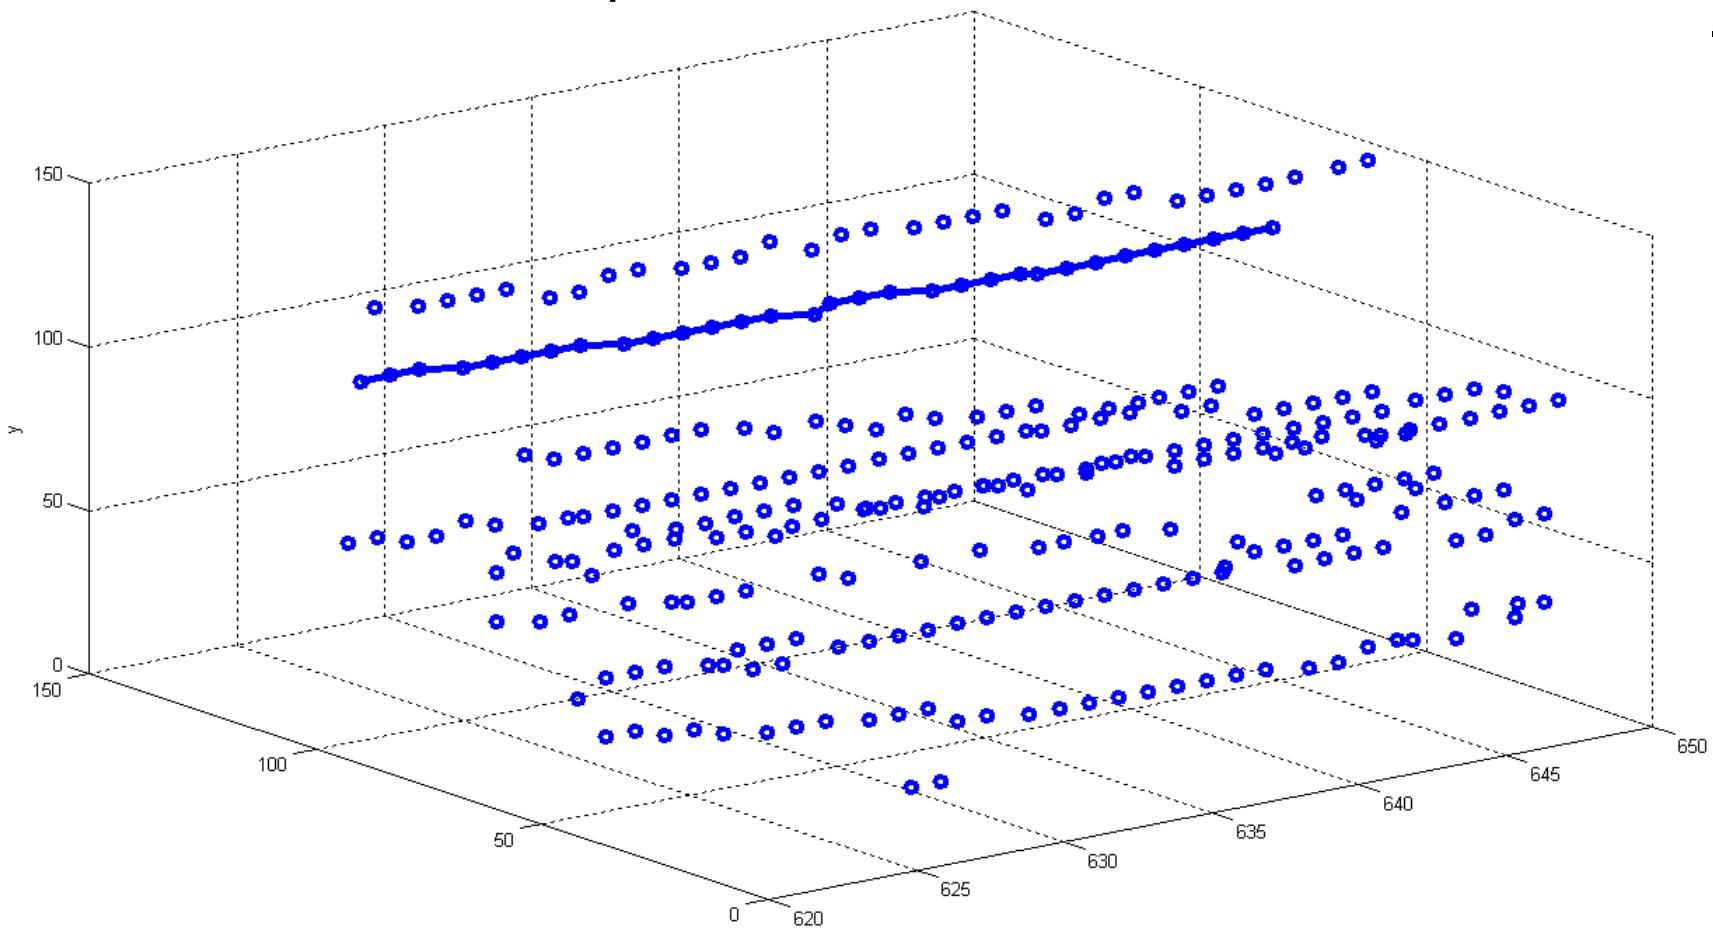
II. Tracklet generation: Create a DAG



- [1] A.Alahi, L.Jacques, Y.Boursier, and P.Vandergheynst, "Sparsity-Driven People Localization with a Heterogeneous Network of Cameras," Journal of Mathematical Imaging and Vision, 2011
- [2] F. Fleuret, J. Berclaz, R. Lengagne, and P. Fua, "Multi-Camera People Tracking With a Probabilistic Occupancy Map", PAMI 2008



II. Tracklet generation: Select longest shortest path with smallest cost

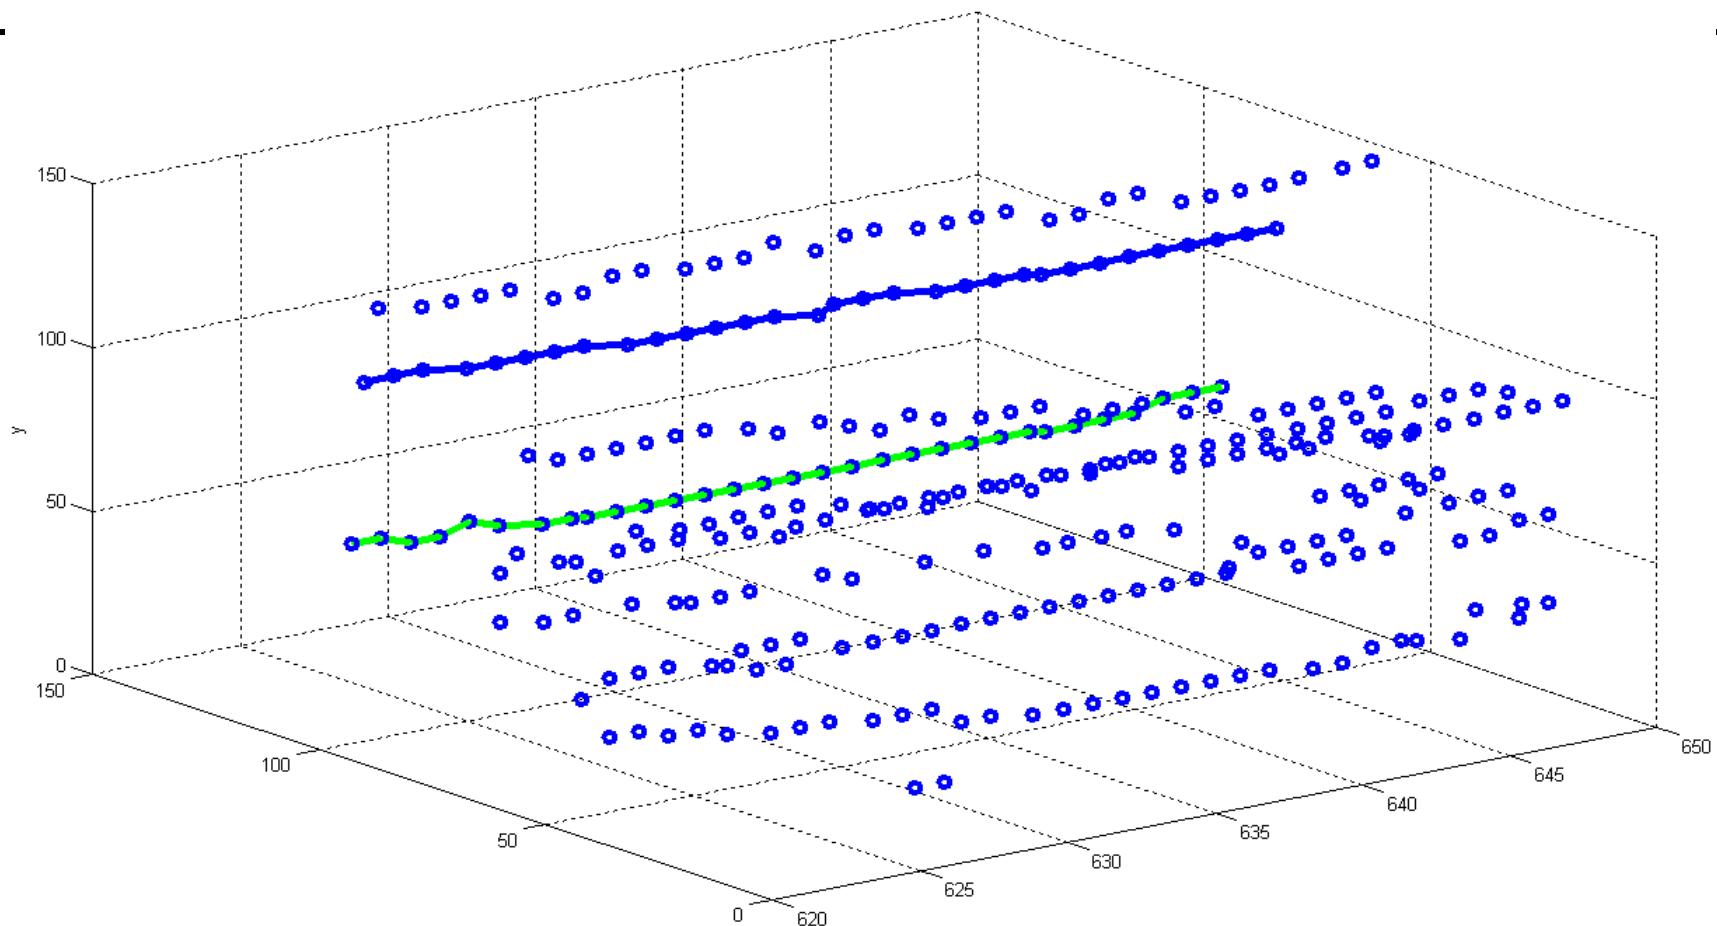


[1] A.Alahi, L.Jacques, Y.Boursier, and P.Vandergheynst, "Sparsity-Driven People Localization with a Heterogeneous Network of Cameras," Journal of Mathematical Imaging and Vision, 2011

[2] F. Fleuret, J. Berclaz, R. Lengagne, and P. Fua, "Multi-Camera People Tracking With a Probabilistic Occupancy Map", PAMI 2008



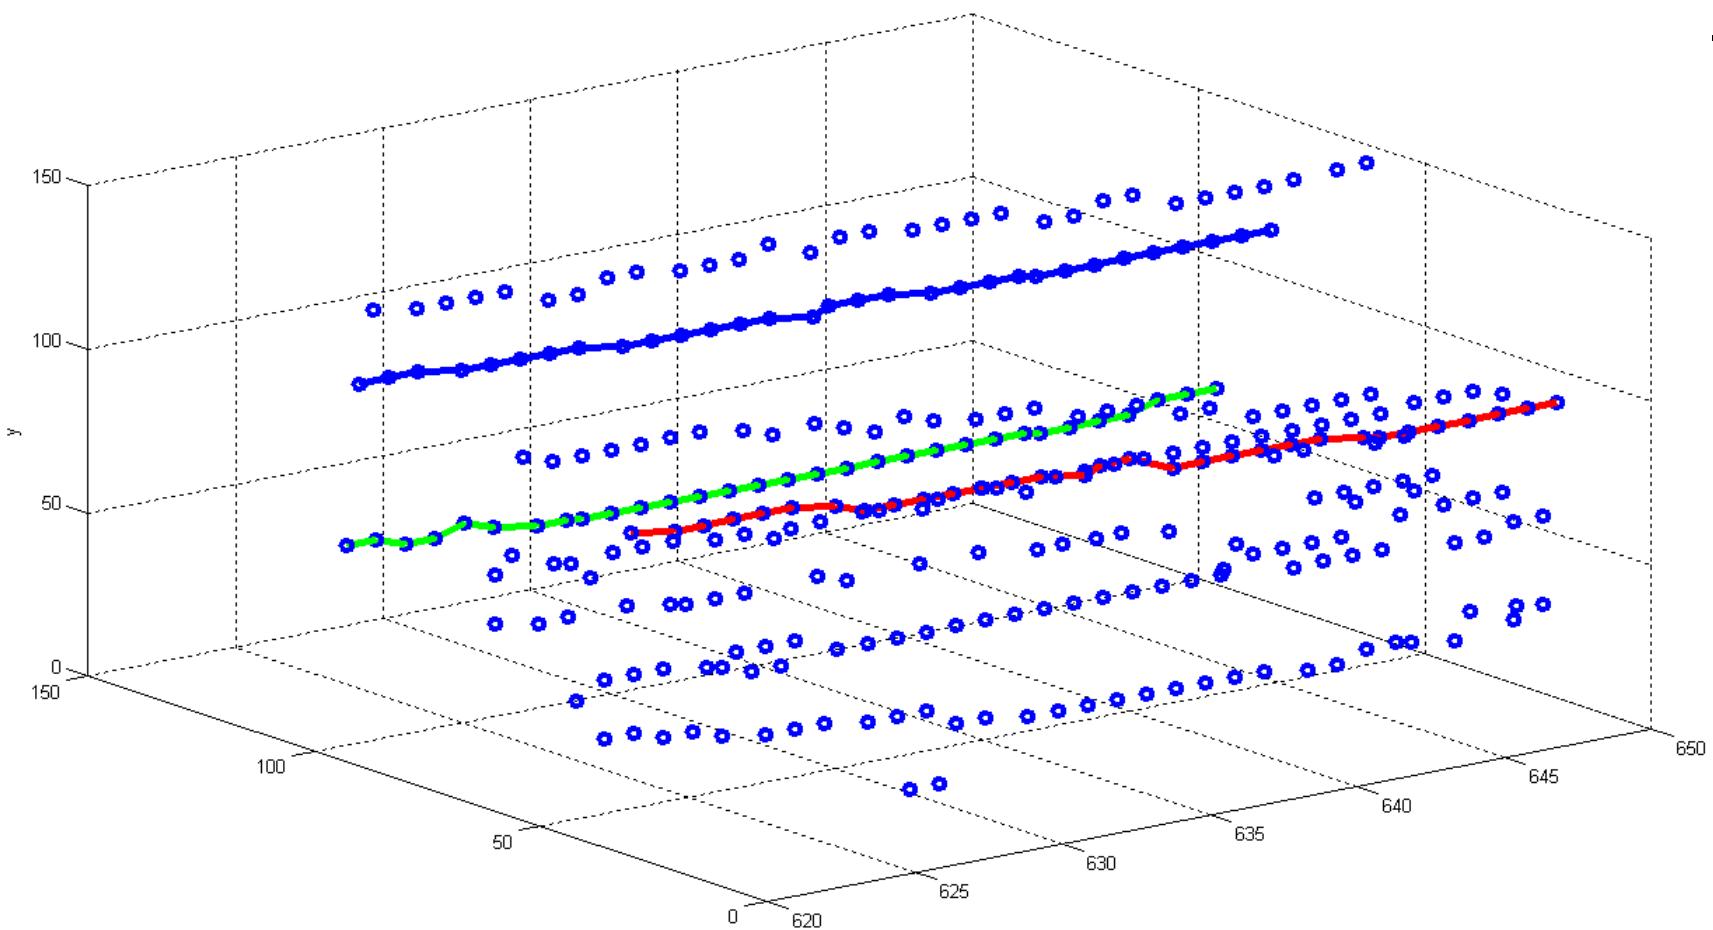
II. Tracklet generation: Iterate



- [1] A.Alahi, L.Jacques, Y.Boursier, and P.Vandergheynst, "Sparsity-Driven People Localization with a Heterogeneous Network of Cameras," Journal of Mathematical Imaging and Vision, 2011
- [2] F. Fleuret, J. Berclaz, R. Lengagne, and P. Fua, "Multi-Camera People Tracking With a Probabilistic Occupancy Map", PAMI 2008



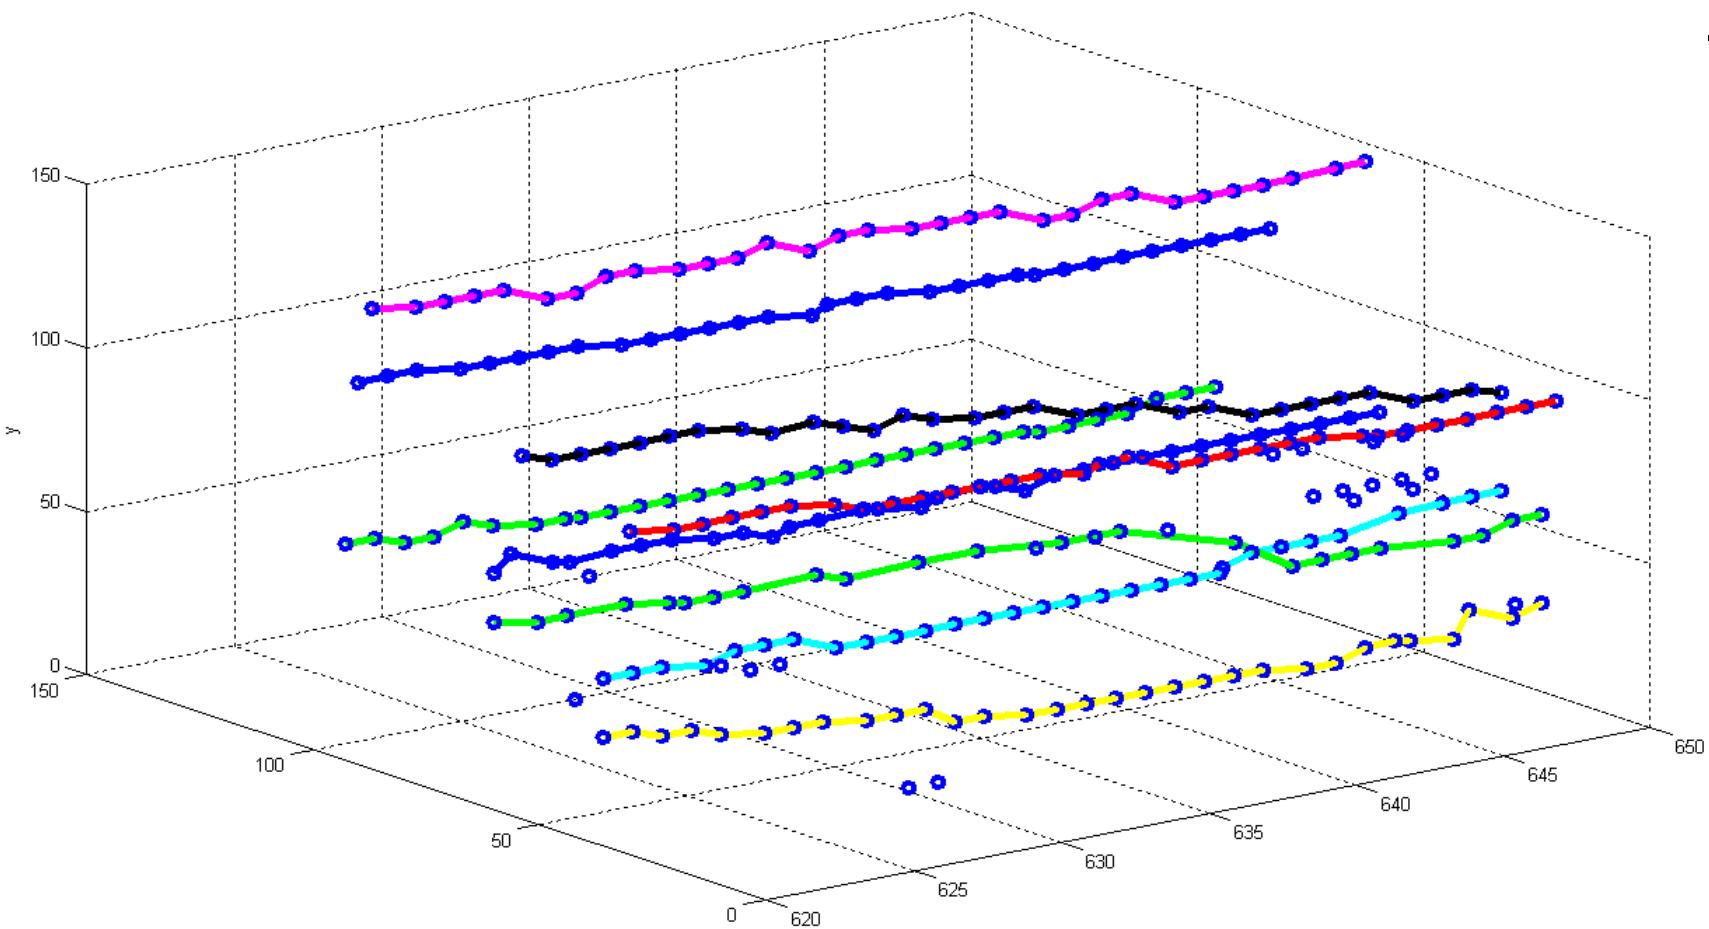
II. Tracklet generation: Iterate



- [1] A.Alahi, L.Jacques, Y.Boursier, and P.Vandergheynst, "Sparsity-Driven People Localization with a Heterogeneous Network of Cameras," Journal of Mathematical Imaging and Vision, 2011
- [2] F. Fleuret, J. Berclaz, R. Lengagne, and P. Fua, "Multi-Camera People Tracking With a Probabilistic Occupancy Map", PAMI 2008



II. Tracklet generation: Till no more paths



[1] A.Alahi, L.Jacques, Y.Boursier, and P.Vandergheynst, "Sparsity-Driven People Localization with a Heterogeneous Network of Cameras," Journal of Mathematical Imaging and Vision, 2011

[2] F. Fleuret, J. Berclaz, R. Lengagne, and P. Fua, "Multi-Camera People Tracking With a Probabilistic Occupancy Map", PAMI 2008



II. Tracklet Generation: Edge cost

- Motion model with social interactions
- Appearance model



II. Tracklet generation: Modeling social interactions



$$\mathbf{F}_i = \mathbf{F}_i^{Goal} + \boxed{\mathbf{F}_i^{Avoidance}} + \mathbf{F}_i^{Attraction} + \mathbf{F}_i^{Scene}$$



$$\mathbf{F}_i^{Avoidance} = \sum_{j \in P \setminus i} \mathbf{f}_{j \rightarrow i}^{Avoidance},$$

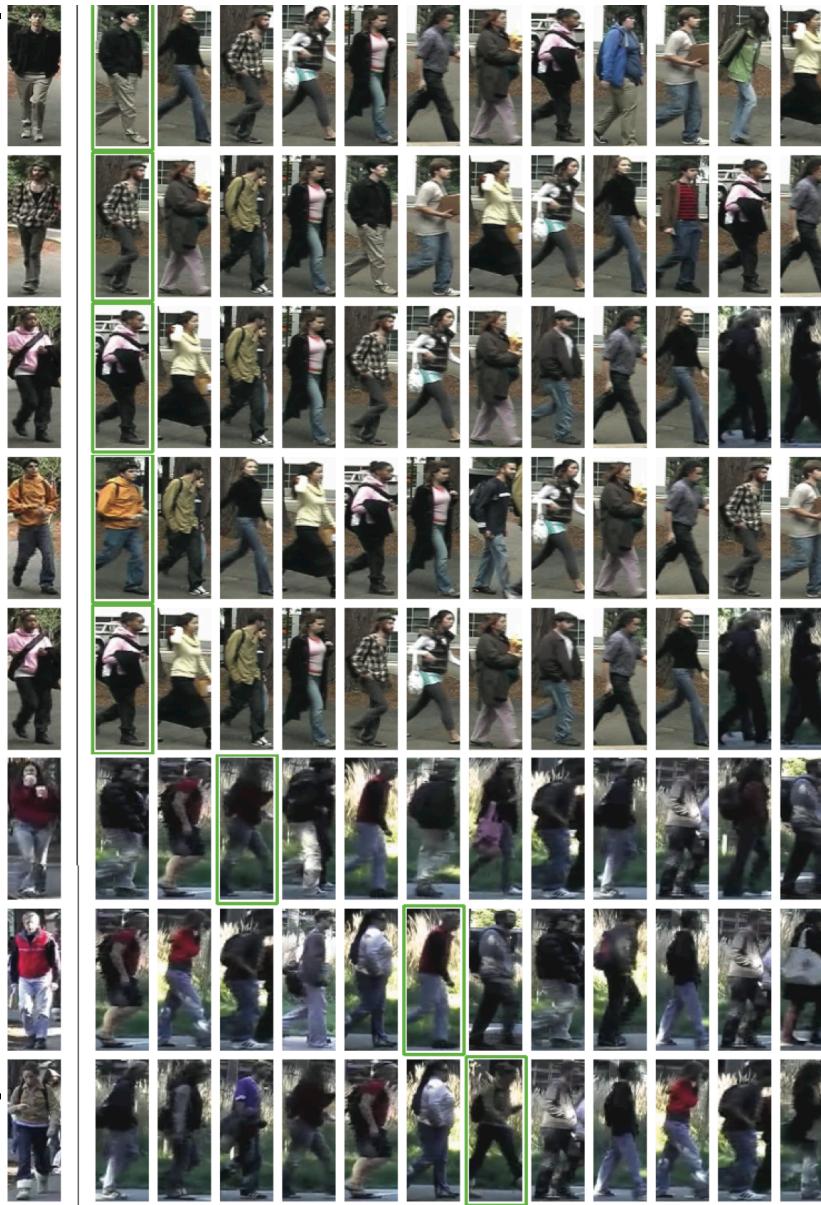
where

$$\mathbf{f}_{j \rightarrow i}^{Avoidance} = \alpha e^{\frac{d_p - d_{ij}}{\beta}} \mathbf{n}_{j \rightarrow i}$$

$$\frac{d}{dt} \mathbf{v} = \frac{\mathbf{F}_i}{m},$$



II. Tracklet Generation: Modeling appearance cues



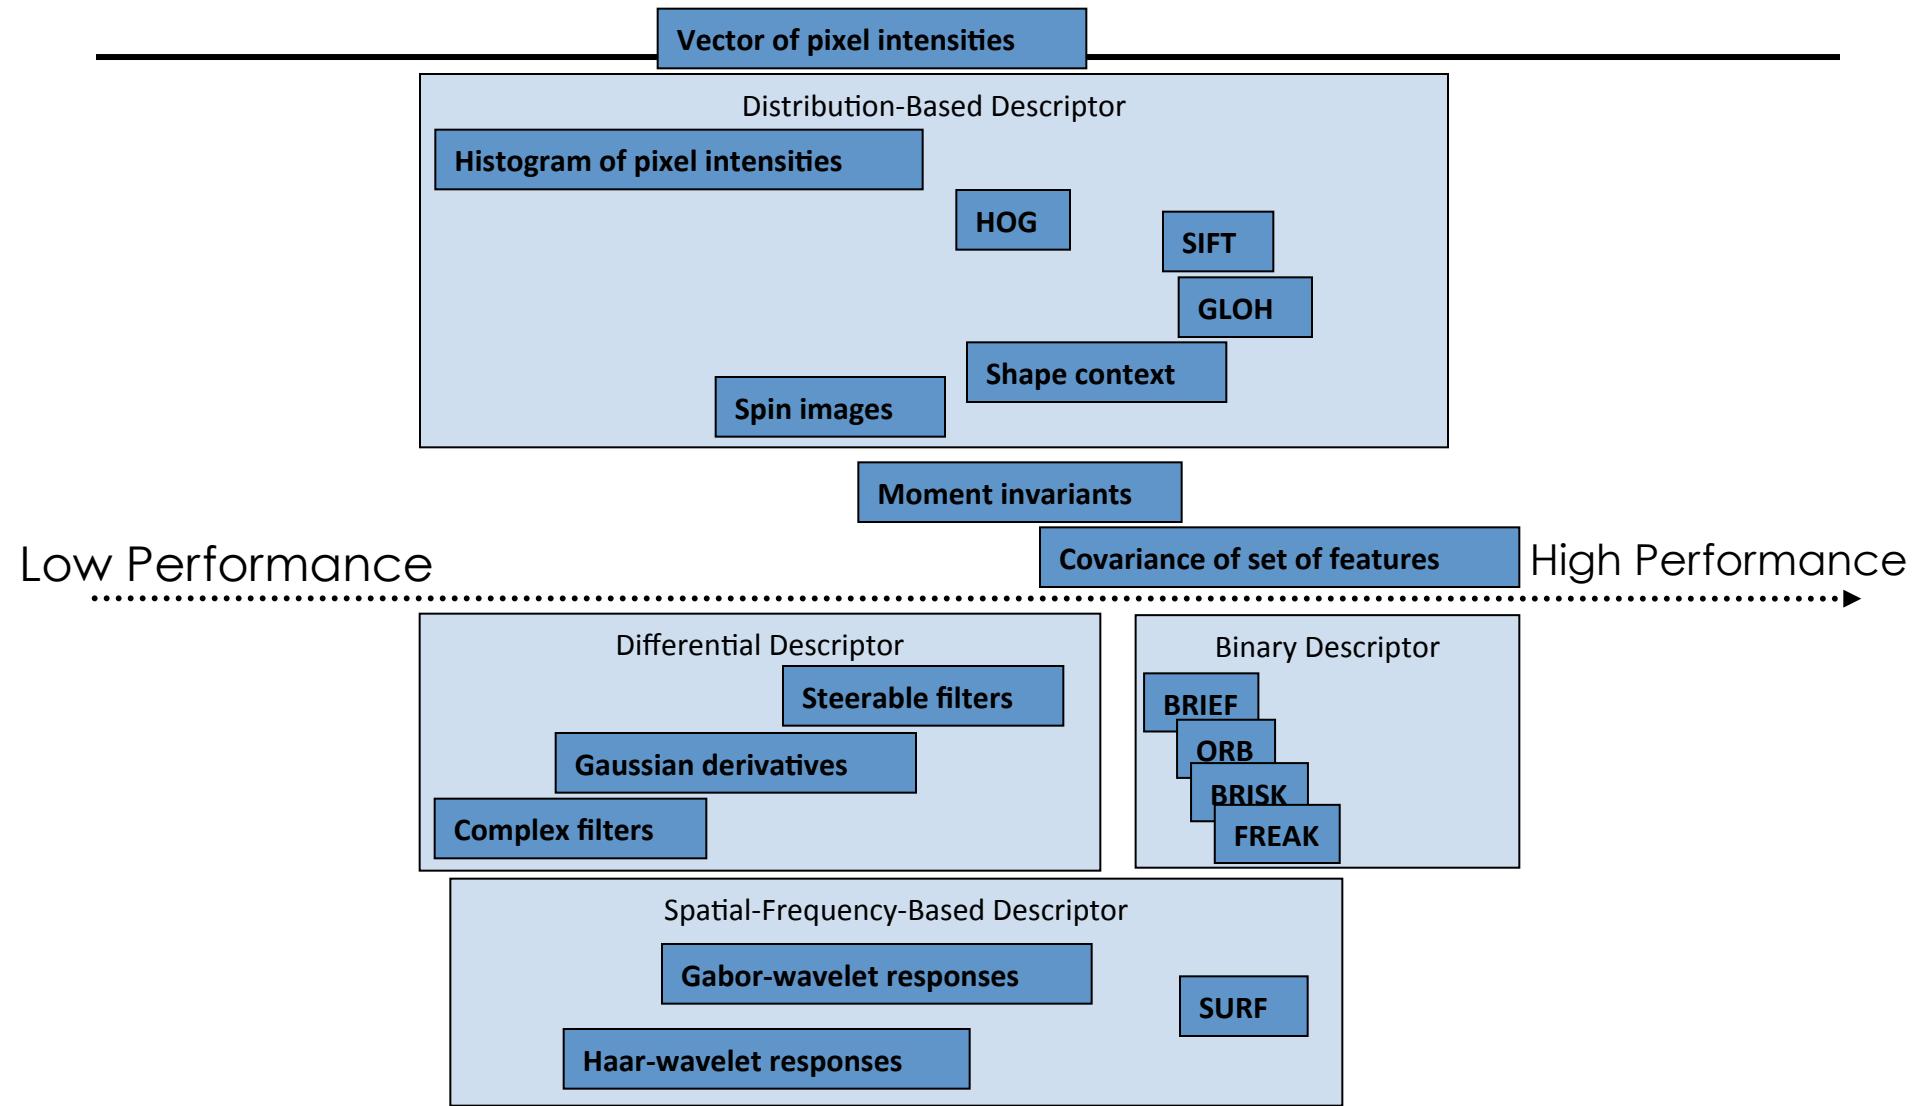
A. Alahi



40

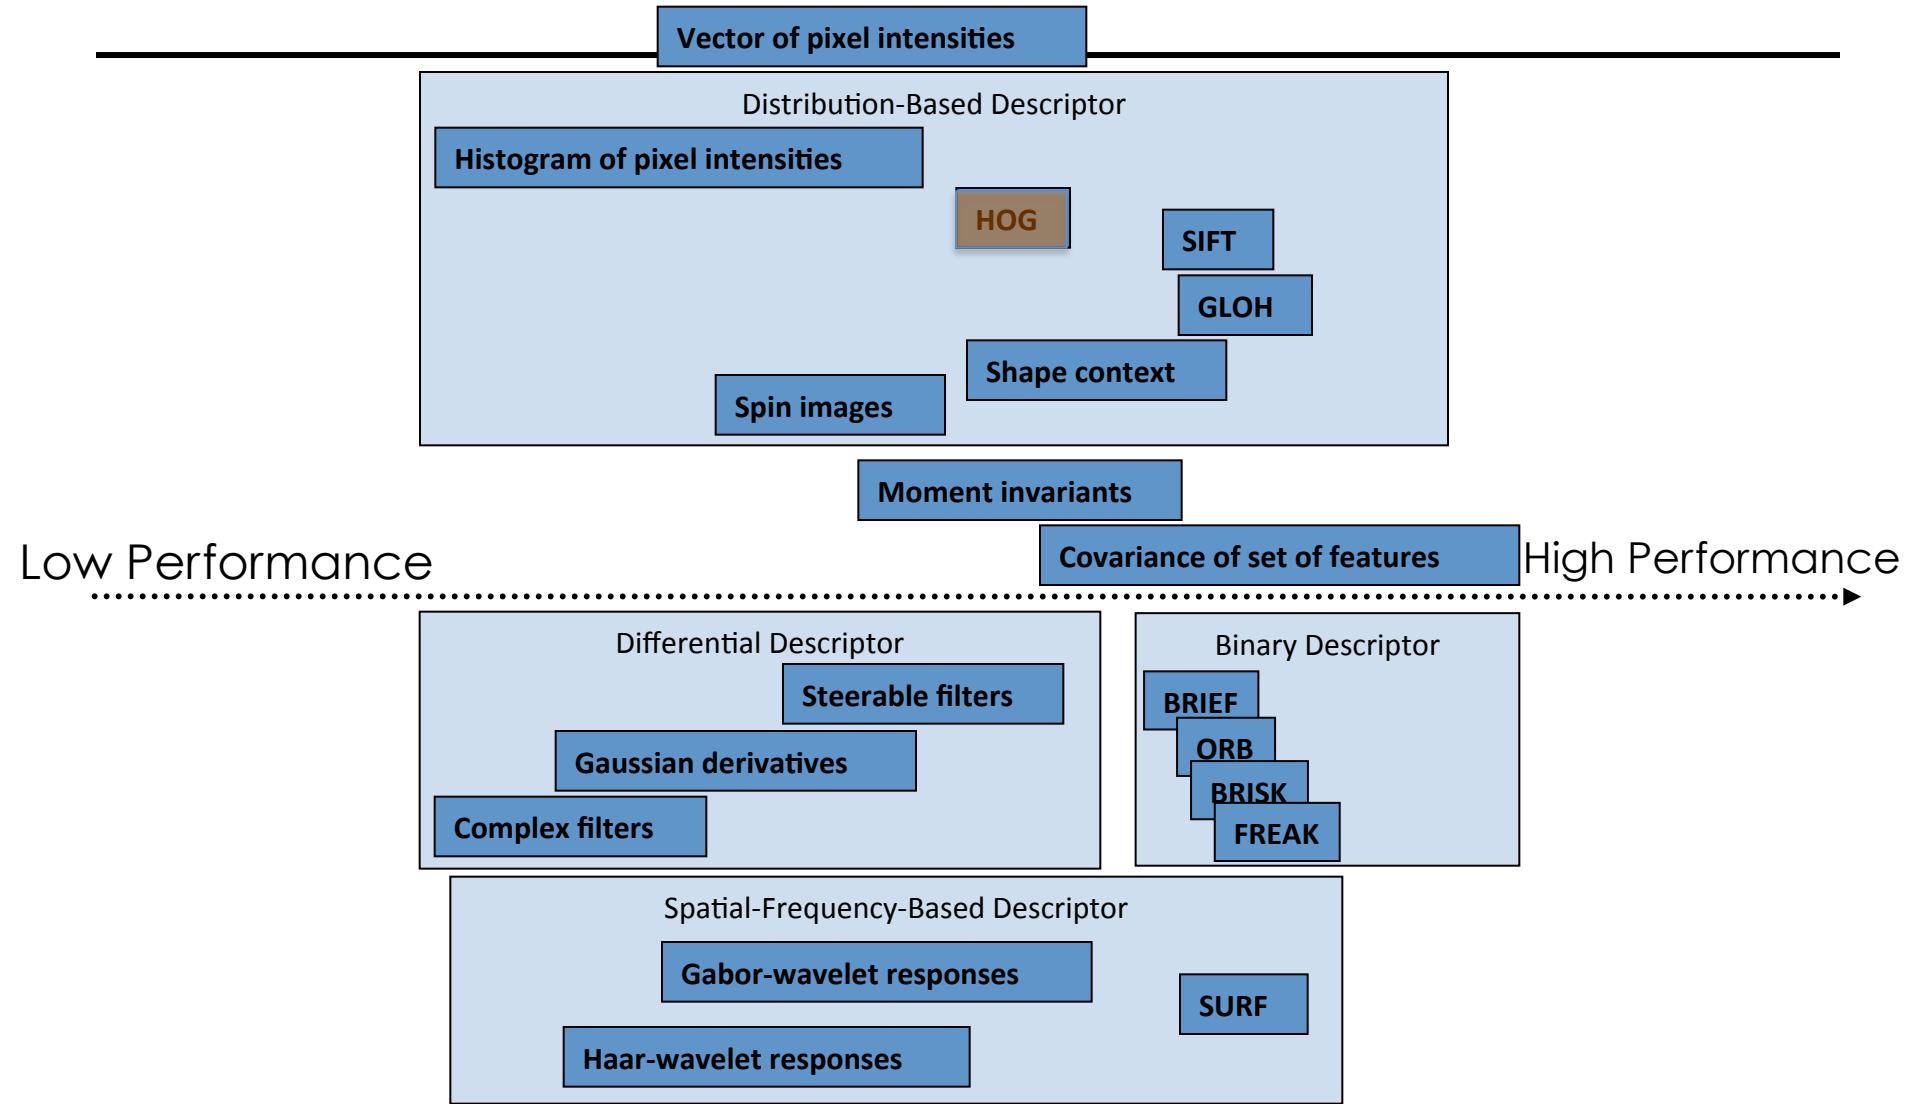
STANFORD
UNIVERSITY

II. Tracklet Generation: An arm-race of image descriptors



- [1] Gabriel, P., Hayet, J., Piater, J., Verly, J.: Object tracking using color interest points
[2] N. Dalal and B. Triggs, “Histograms of oriented gradients for human detection,”,
[3] O. Tuzel, F. Porikli, and P. Meer, “Region covariance: A fast descriptor for detection and classification,”.

II. Tracklet Generation: An arm-race of image descriptors



[1] Gabriel, P., Hayet, J., Piater, J., Verly, J.: Object tracking using color interest points

[2] N. Dalal and B. Triggs, "Histograms of oriented gradients for human detection,"

[3] O. Tuzel, F. Porikli, and P. Meer, "Region covariance: A fast descriptor for detection and classification,".

II. Tracklet Generation: HOG

Image gradient

- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

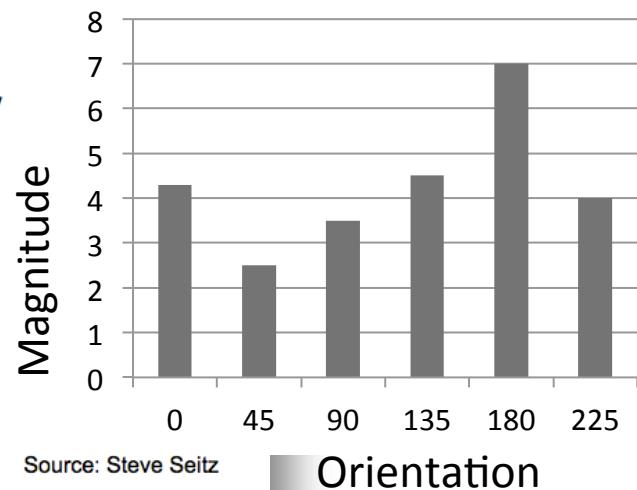
The gradient points in the direction of most rapid increase in intensity

The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

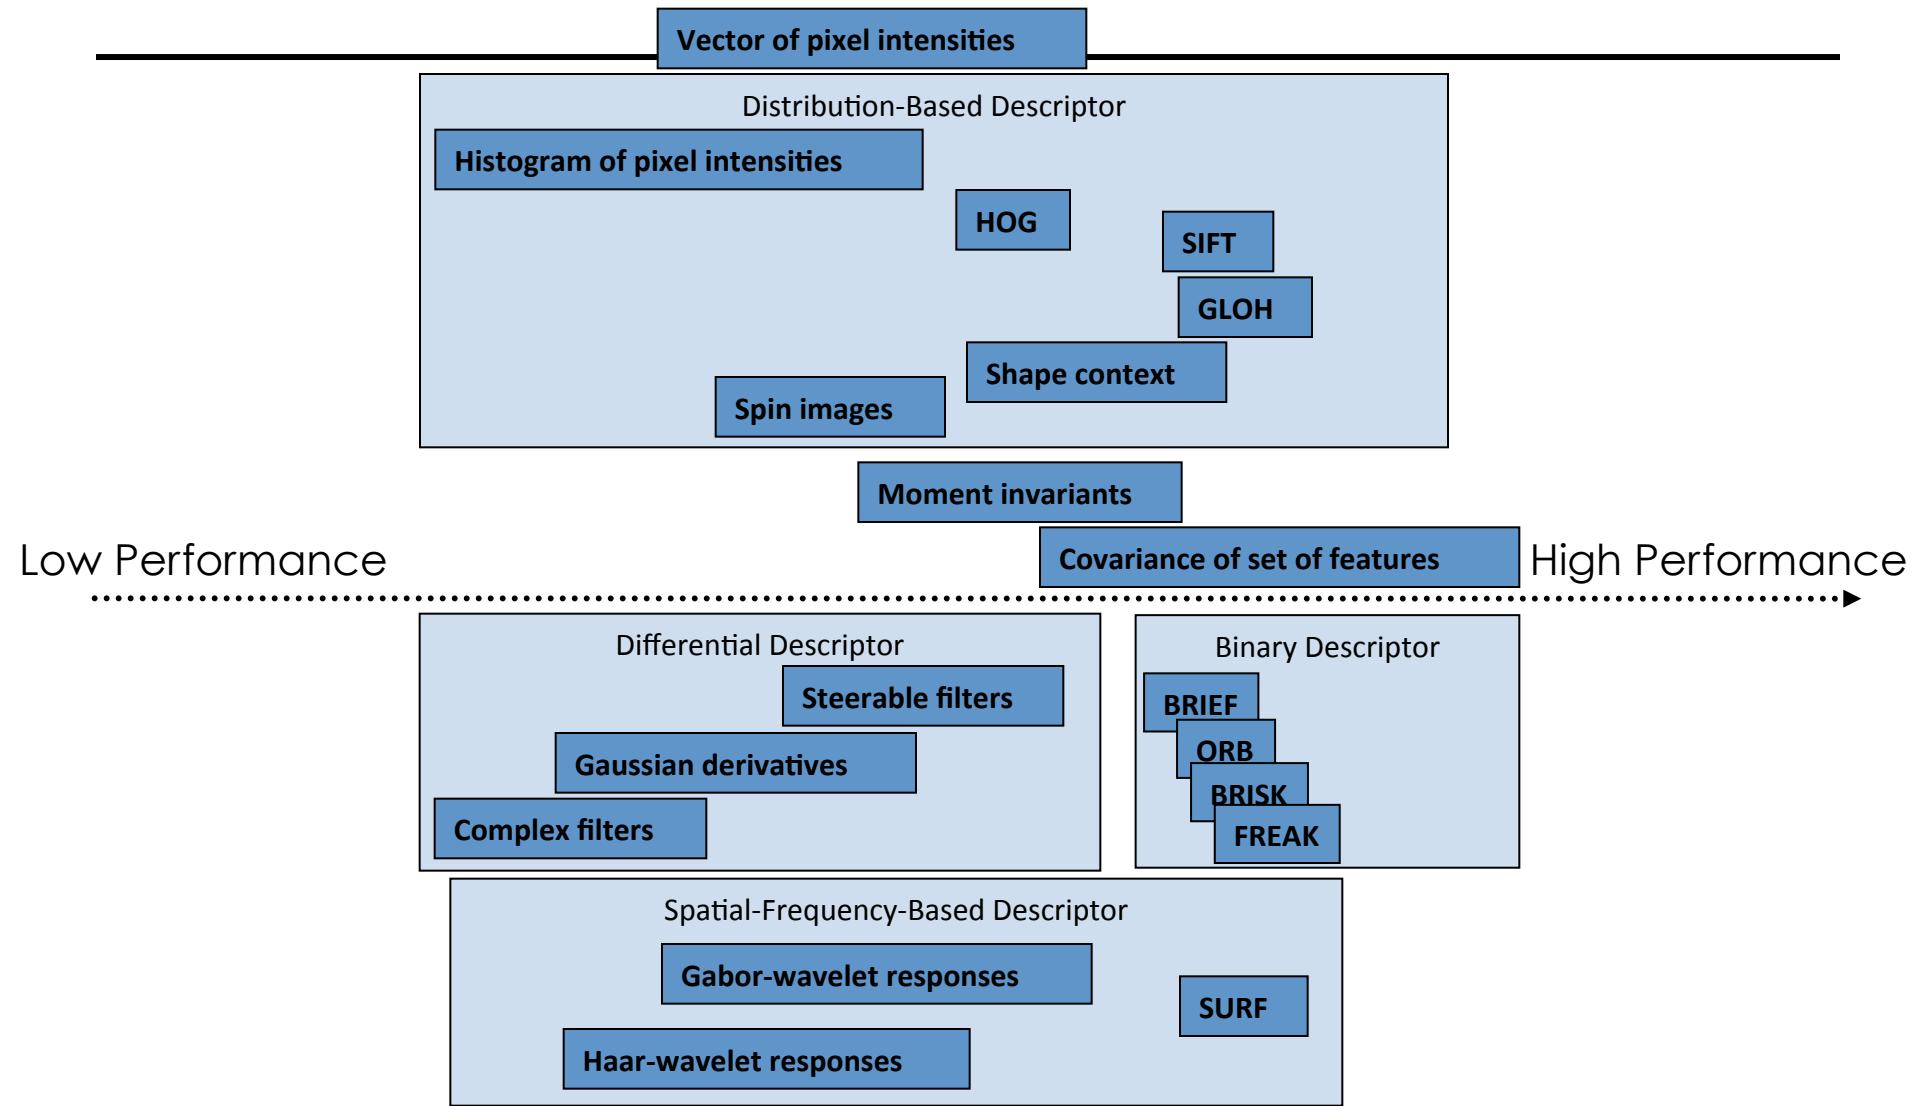
- how does this relate to the direction of the edge?

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

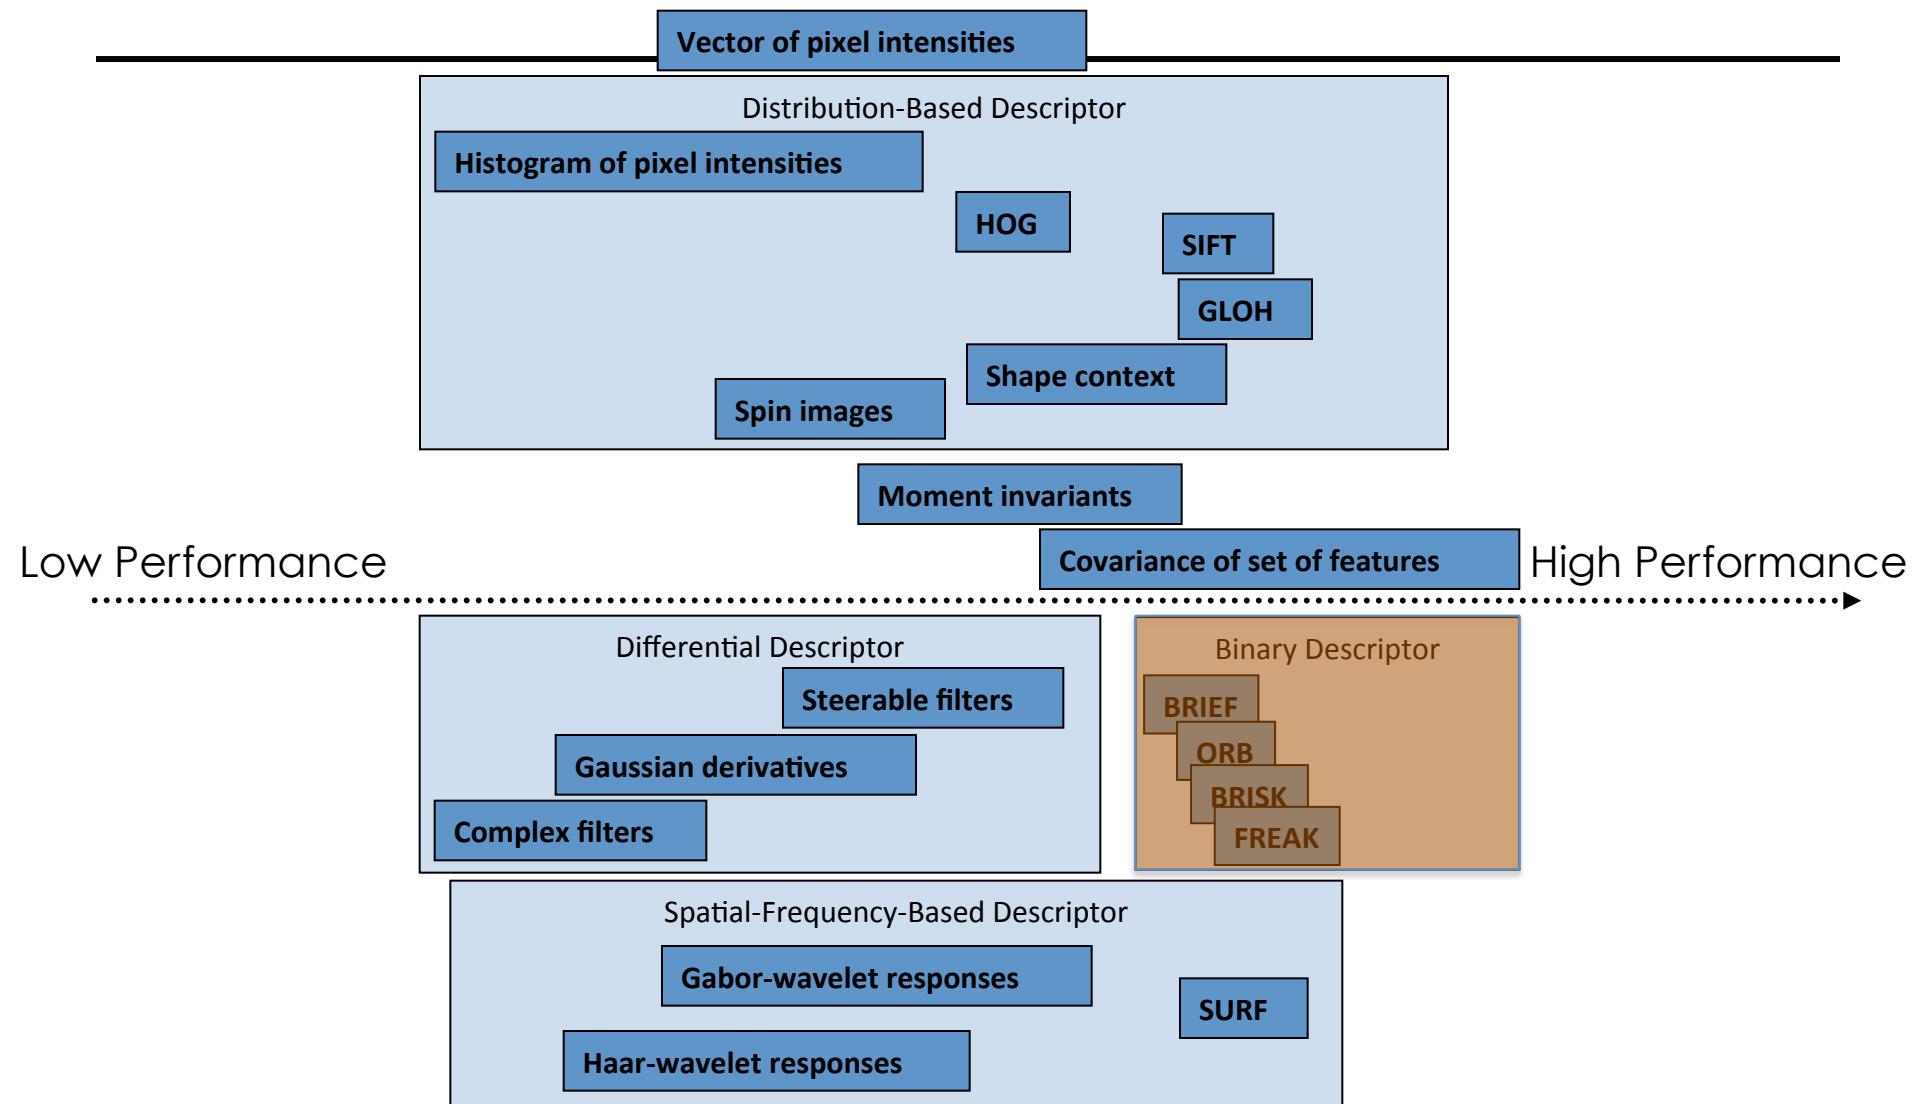


II. Tracklet Generation: An arm-race of image descriptors



- [1] Gabriel, P., Hayet, J., Piater, J., Verly, J.: Object tracking using color interest points
[2] N. Dalal and B. Triggs, "Histograms of oriented gradients for human detection,"
[3] O. Tuzel, F. Porikli, and P. Meer, "Region covariance: A fast descriptor for detection and classification,".

II. Tracklet Generation: An arm-race of image descriptors



[1] Gabriel, P., Hayet, J., Platé, J., Verly, J.. Object tracking using color interest points

[2] N. Dalal and B. Triggs, "Histograms of oriented gradients for human detection,"

[3] O. Tuzel, F. Porikli, and P. Meer, "Region covariance: A fast descriptor for detection and classification,".

II. Tracklet Generation: Binary descriptors



II. Tracklet Generation: BRIEF[1] / ORB[2] / BRISK[3]



[1] Calonder, Michael, et al. "Brief: Binary robust independent elementary features." *ECCV 2010*.

[2] Rublee, Ethan, et al. "ORB: an efficient alternative to SIFT or SURF." *ICCV 2011*

[3] Leutenegger,.., et al. "BRISK: Binary robust invariant scalable keypoints." *ICCV 2011*

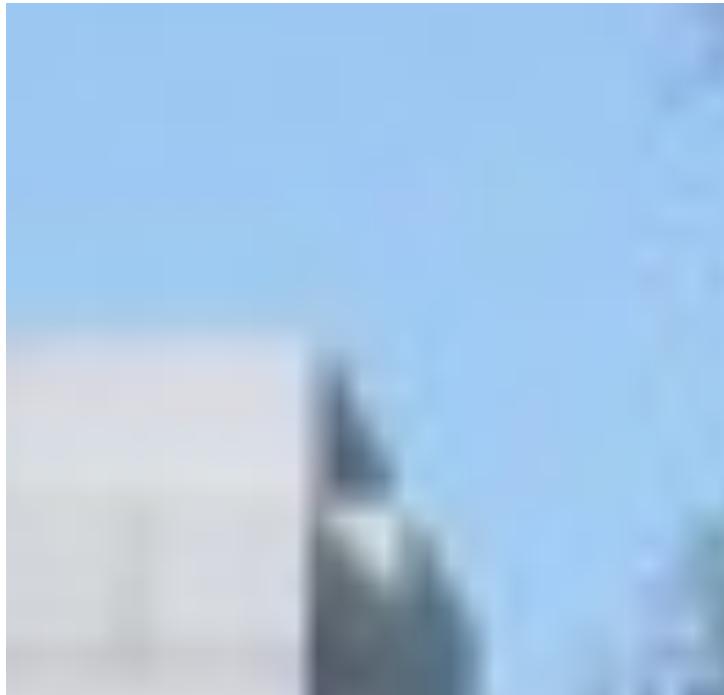
II. Tracklet Generation: BRIEF[1] / ORB[2] / BRISK[3]



-
- A. Alahi
- [1] Calonder, Michael, et al. "Brief: Binary robust independent elementary features." *ECCV 2010*.
 - [2] Rublee, Ethan, et al. "ORB: an efficient alternative to SIFT or SURF." *ICCV 2011*
 - [3] Leutenegger,.., et al. "BRISK: Binary robust invariant scalable keypoints." *ICCV 2011*



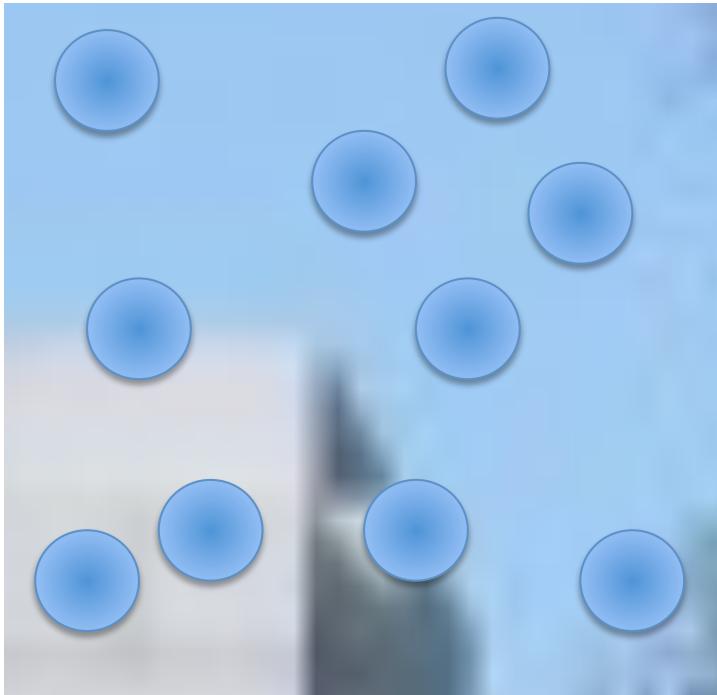
II. Tracklet Generation: BRIEF[1] / ORB[2] / BRISK[3]



-
- [1] Calonder, Michael, et al. "Brief: Binary robust independent elementary features." *ECCV 2010*.
 - [2] Rublee, Ethan, et al. "ORB: an efficient alternative to SIFT or SURF." *ICCV 2011*
 - [3] Leutenegger,.., et al. "BRISK: Binary robust invariant scalable keypoints." *ICCV 2011*

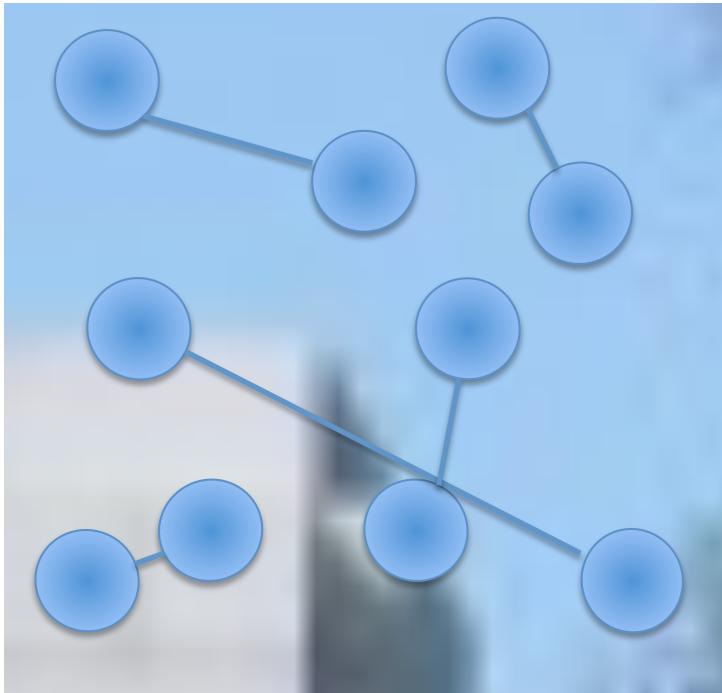


II. Tracklet Generation: **BRIEF[1]** / **ORB[2]** / **BRISK[3]**



-
- [1] Calonder, Michael, et al. "Brief: Binary robust independent elementary features." *ECCV 2010*.
 - [2] Rublee, Ethan, et al. "ORB: an efficient alternative to SIFT or SURF." *ICCV 2011*
 - [3] Leutenegger,.., et al. "BRISK: Binary robust invariant scalable keypoints." *ICCV 2011*

II. Tracklet Generation: **BRIEF**[1] / **ORB**[2] / **BRISK**[3]



A sequence of 1-bit DoG

[1] Calonder, Michael, et al. "Brief: Binary robust independent elementary features." *ECCV 2010*.

[2] Rublee, Ethan, et al. "ORB: an efficient alternative to SIFT or SURF." *ICCV 2011*

[3] Leutenegger,.., et al. "BRISK: Binary robust invariant scalable keypoints." *ICCV 2011*

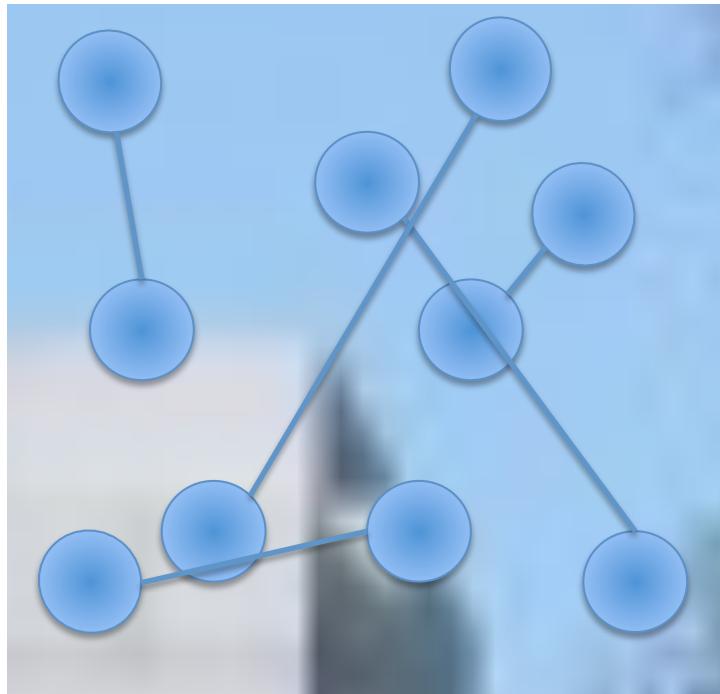
II. Tracklet Generation: BRIEF[1] / ORB[2] / BRISK[3]

- Select

1) Most discriminant

AND

2) Less correlated

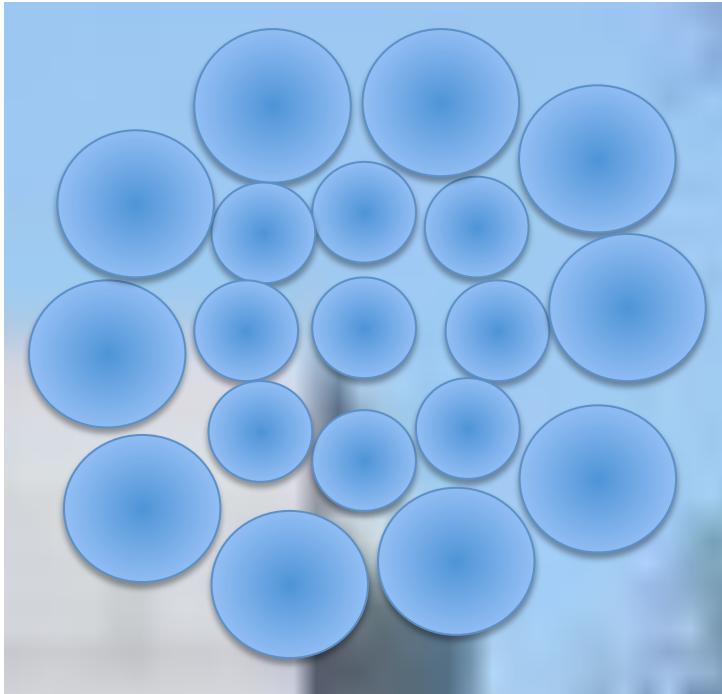


[1] Calonder, Michael, et al. "Brief: Binary robust independent elementary features." *ECCV 2010*.

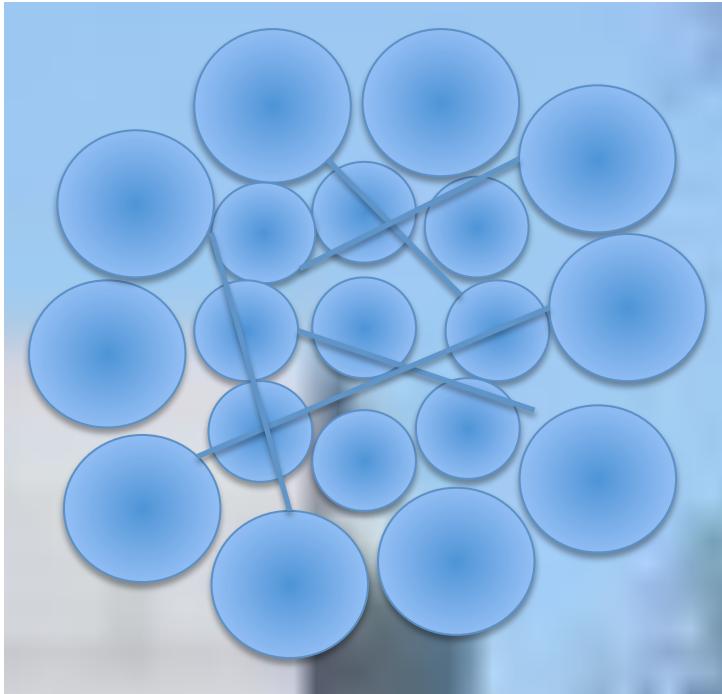
[2] Rublee, Ethan, et al. "ORB: an efficient alternative to SIFT or SURF." *ICCV 2011*

[3] Leutenegger,.., et al. "BRISK: Binary robust invariant scalable keypoints." *ICCV 2011*

II. Tracklet Generation: BRIEF[1] / ORB[2] / **BRISK[3]**



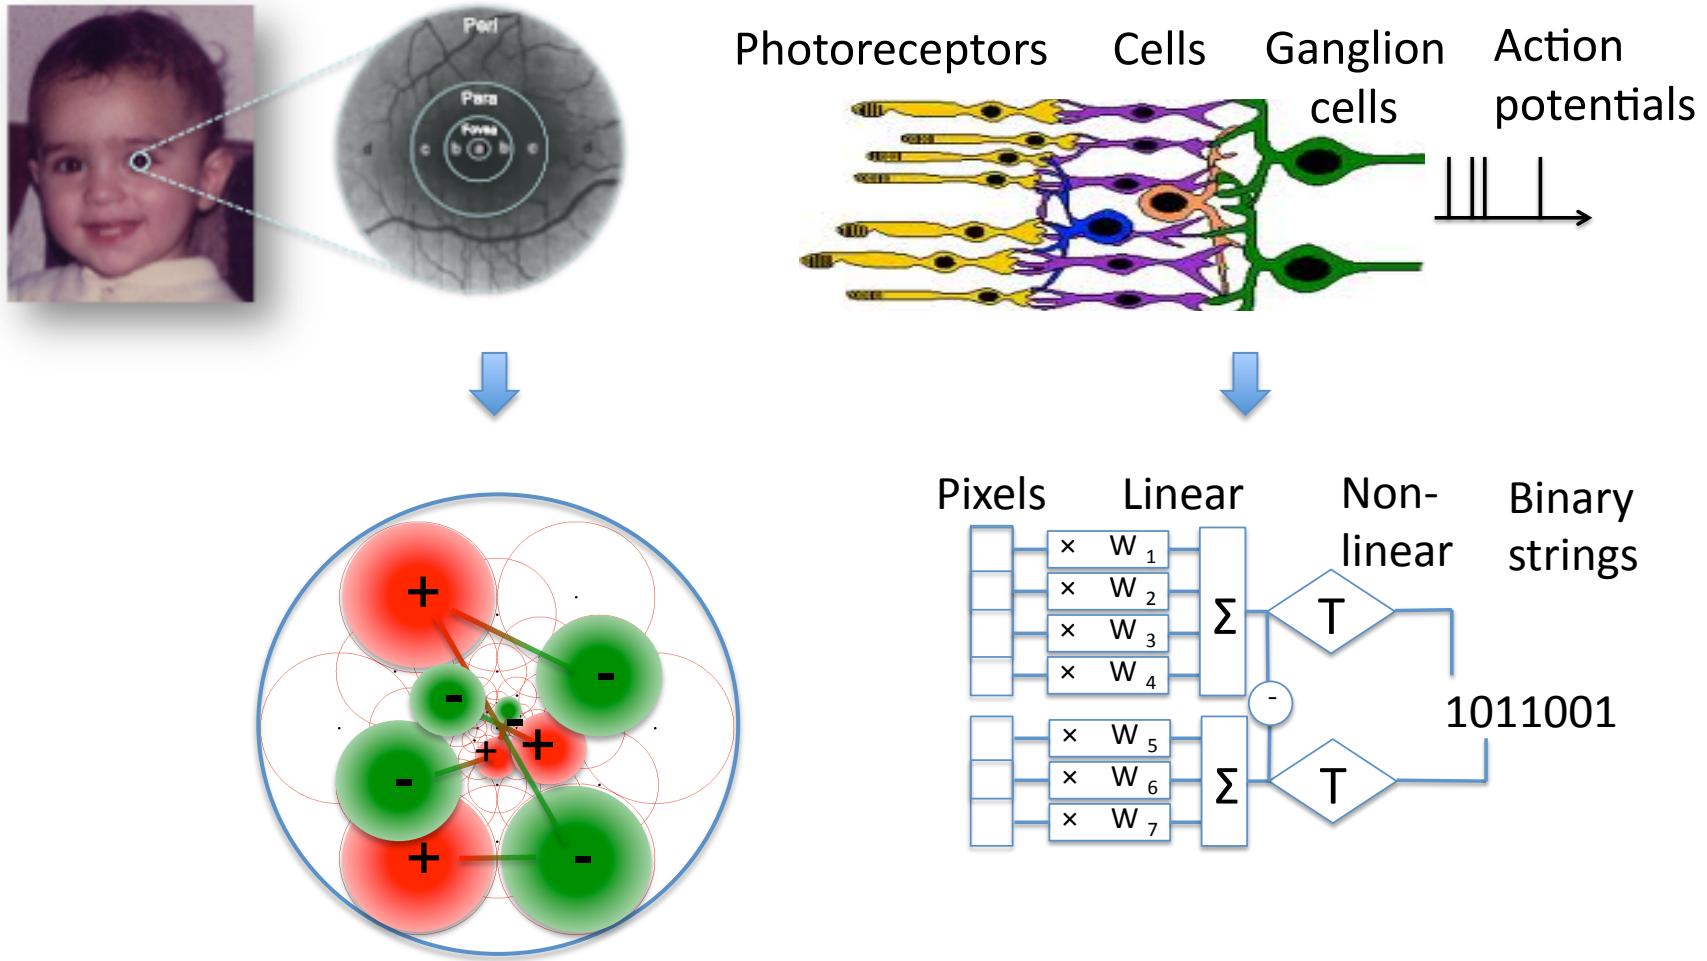
II. Tracklet Generation: BRIEF[1] / ORB[2] / **BRISK[3]**



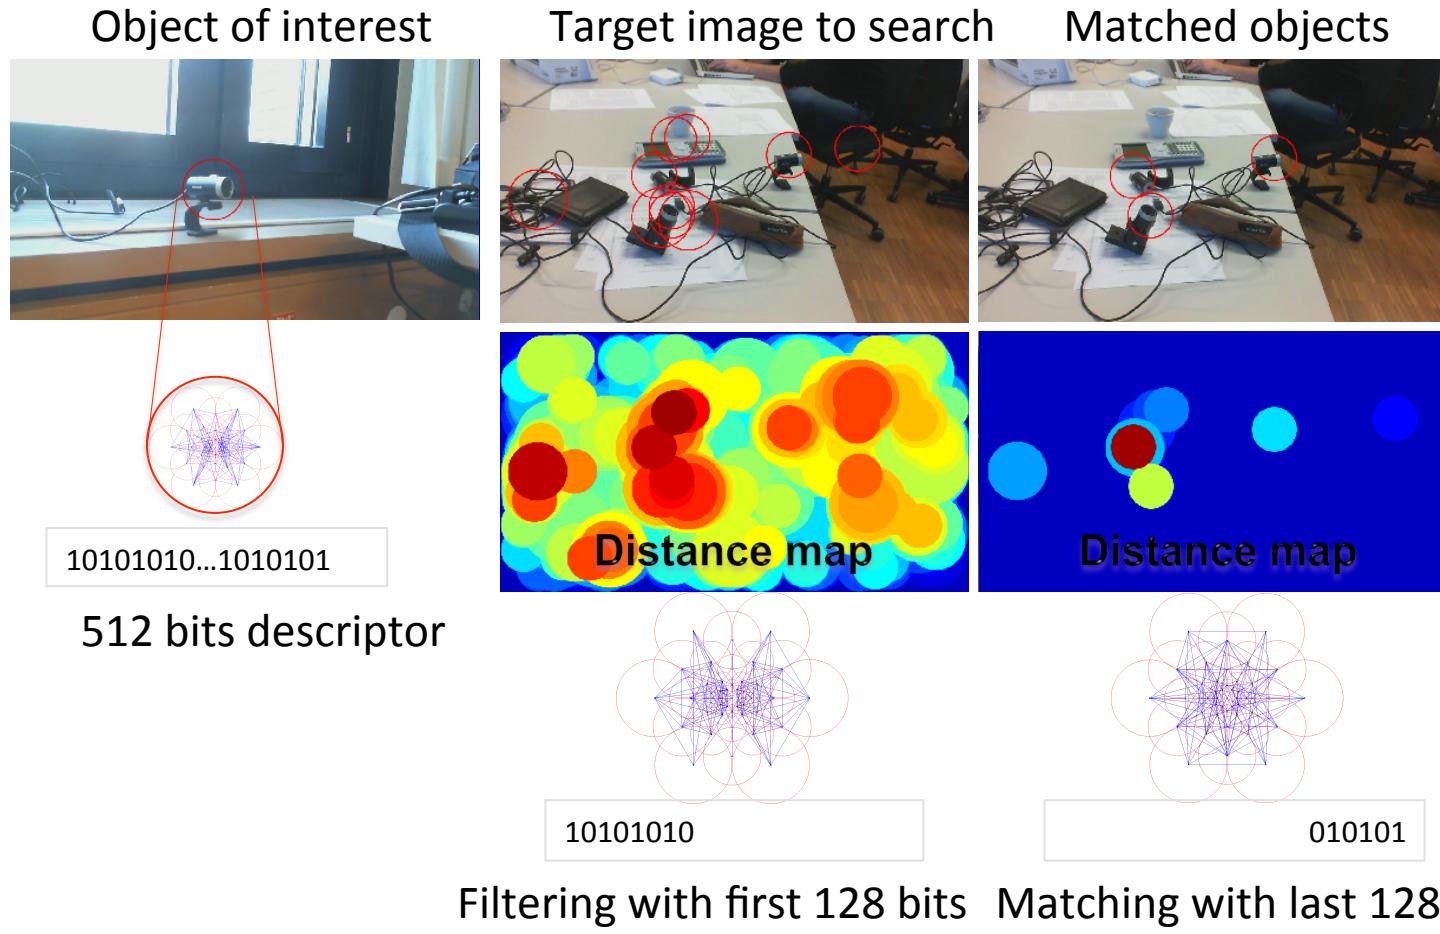
-
- [1] Calonder, Michael, et al. "Brief: Binary robust independent elementary features." *ECCV 2010*.
 - [2] Rublee, Ethan, et al. "ORB: an efficient alternative to SIFT or SURF." *ICCV 2011*
 - [3] Leutenegger,.., et al. "BRISK: Binary robust invariant scalable keypoints." *ICCV 2011*



II. Tracklet Generation: Retina-inspired [1]



II. Tracklet Generation: Saccadic search

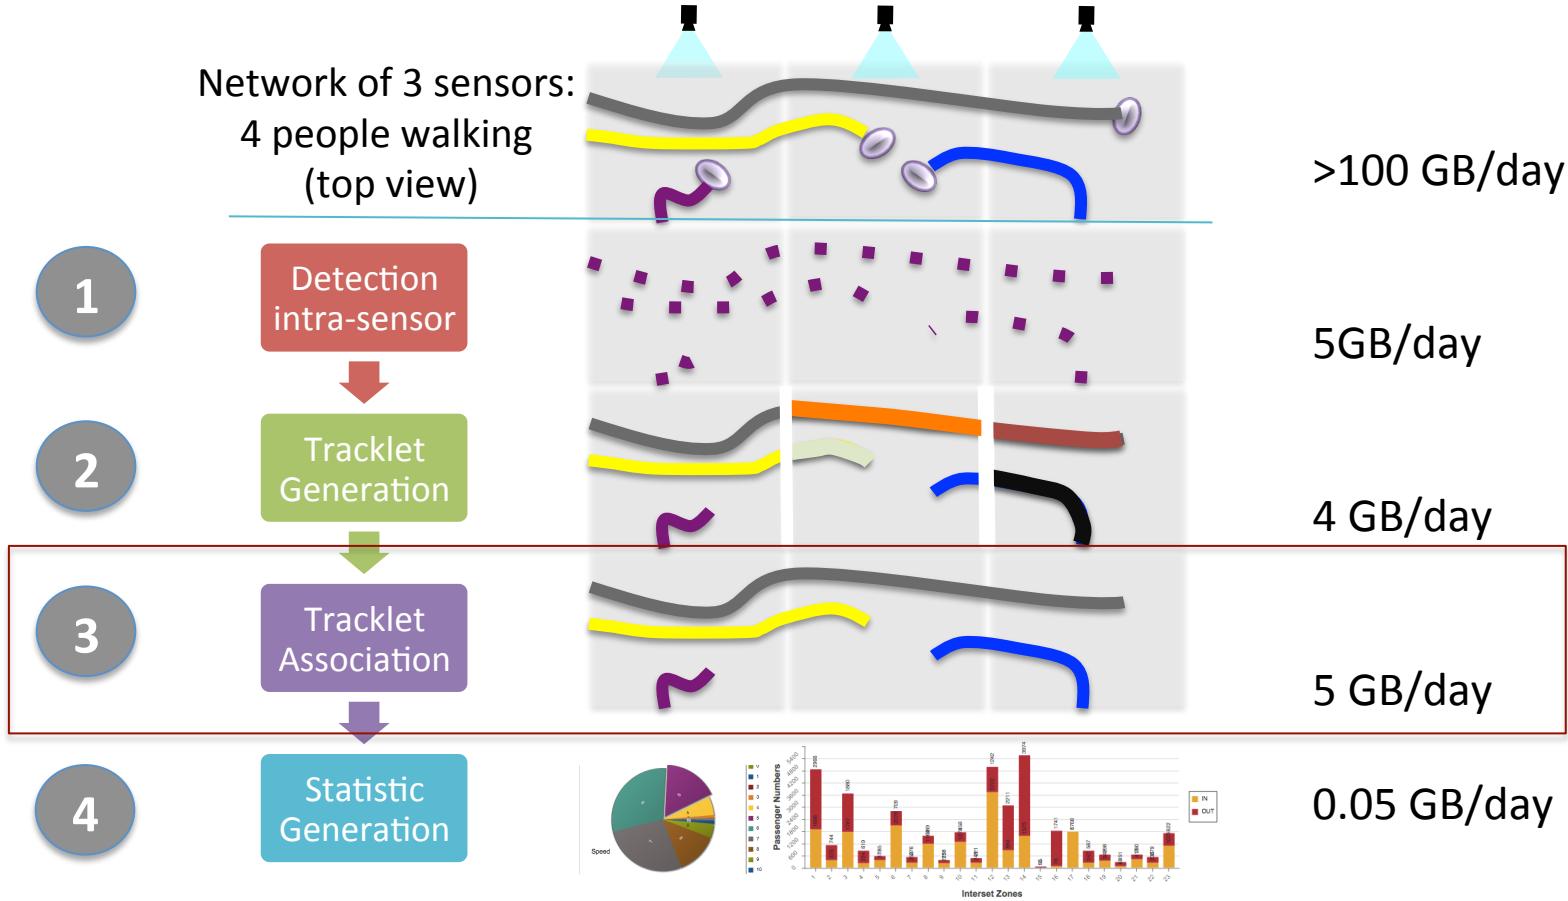


Tracking Example

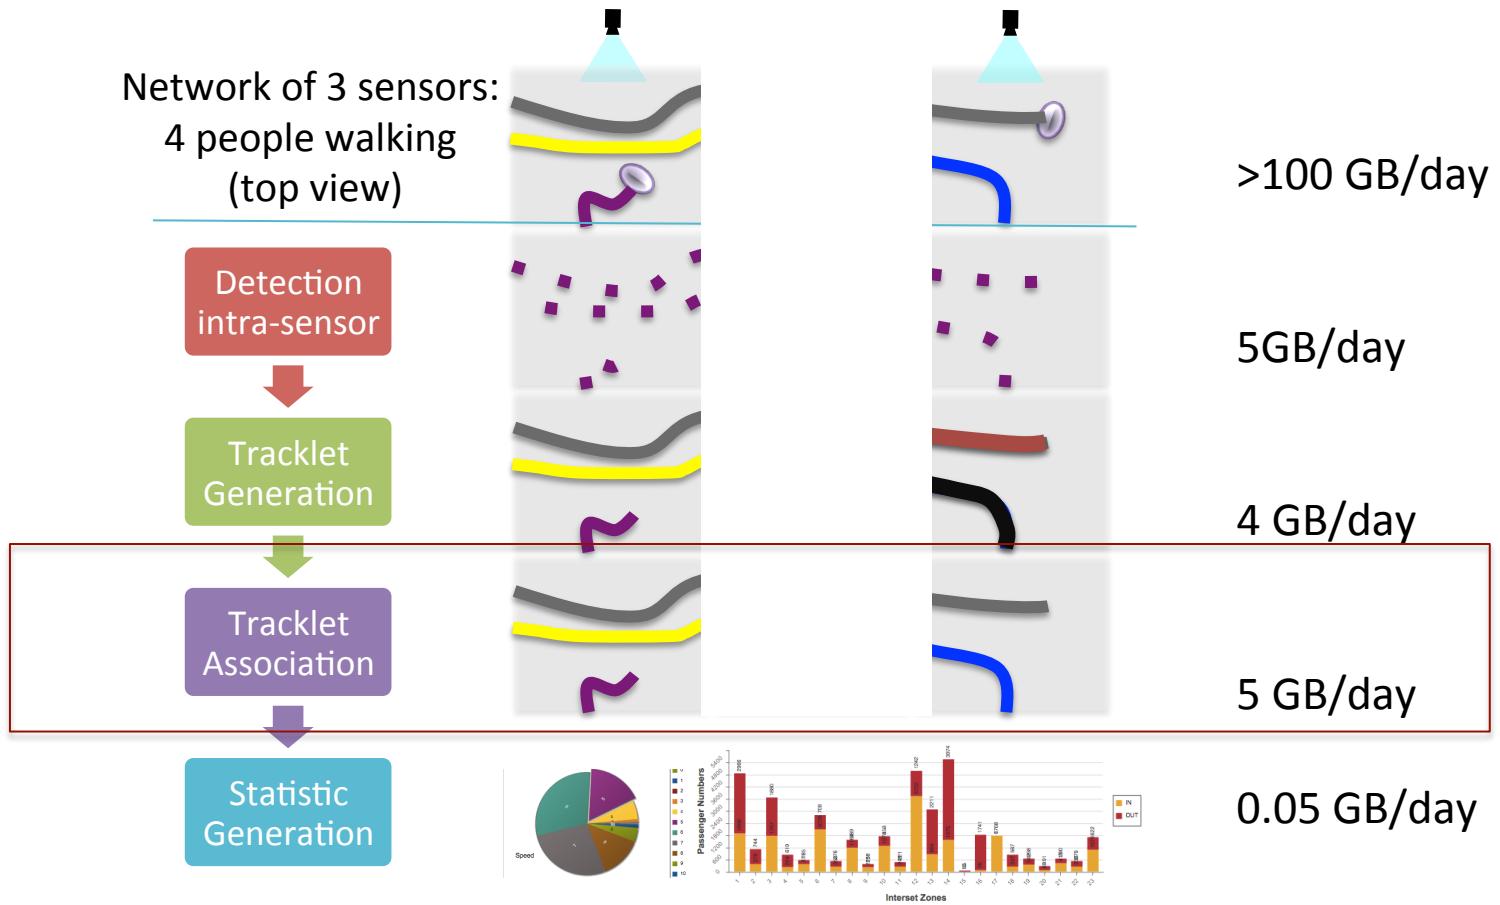


Outline:

From Foreground Extraction To Tracking 42 million Pedestrians



Tracklet association in scattered network



III- Tracklet association: Problem formulation

Let:

- \mathbf{T} : all long term trajectories
- \mathbf{t} : tracklets (tracklets capture within each camera)
- Problem: Maximizing the a posterior probability (MAP) of \mathbf{T} :

$$\mathbf{T}^* = \arg \max_{\mathbf{T}} P(\mathbf{T} | \mathbf{t}) \quad (1)$$

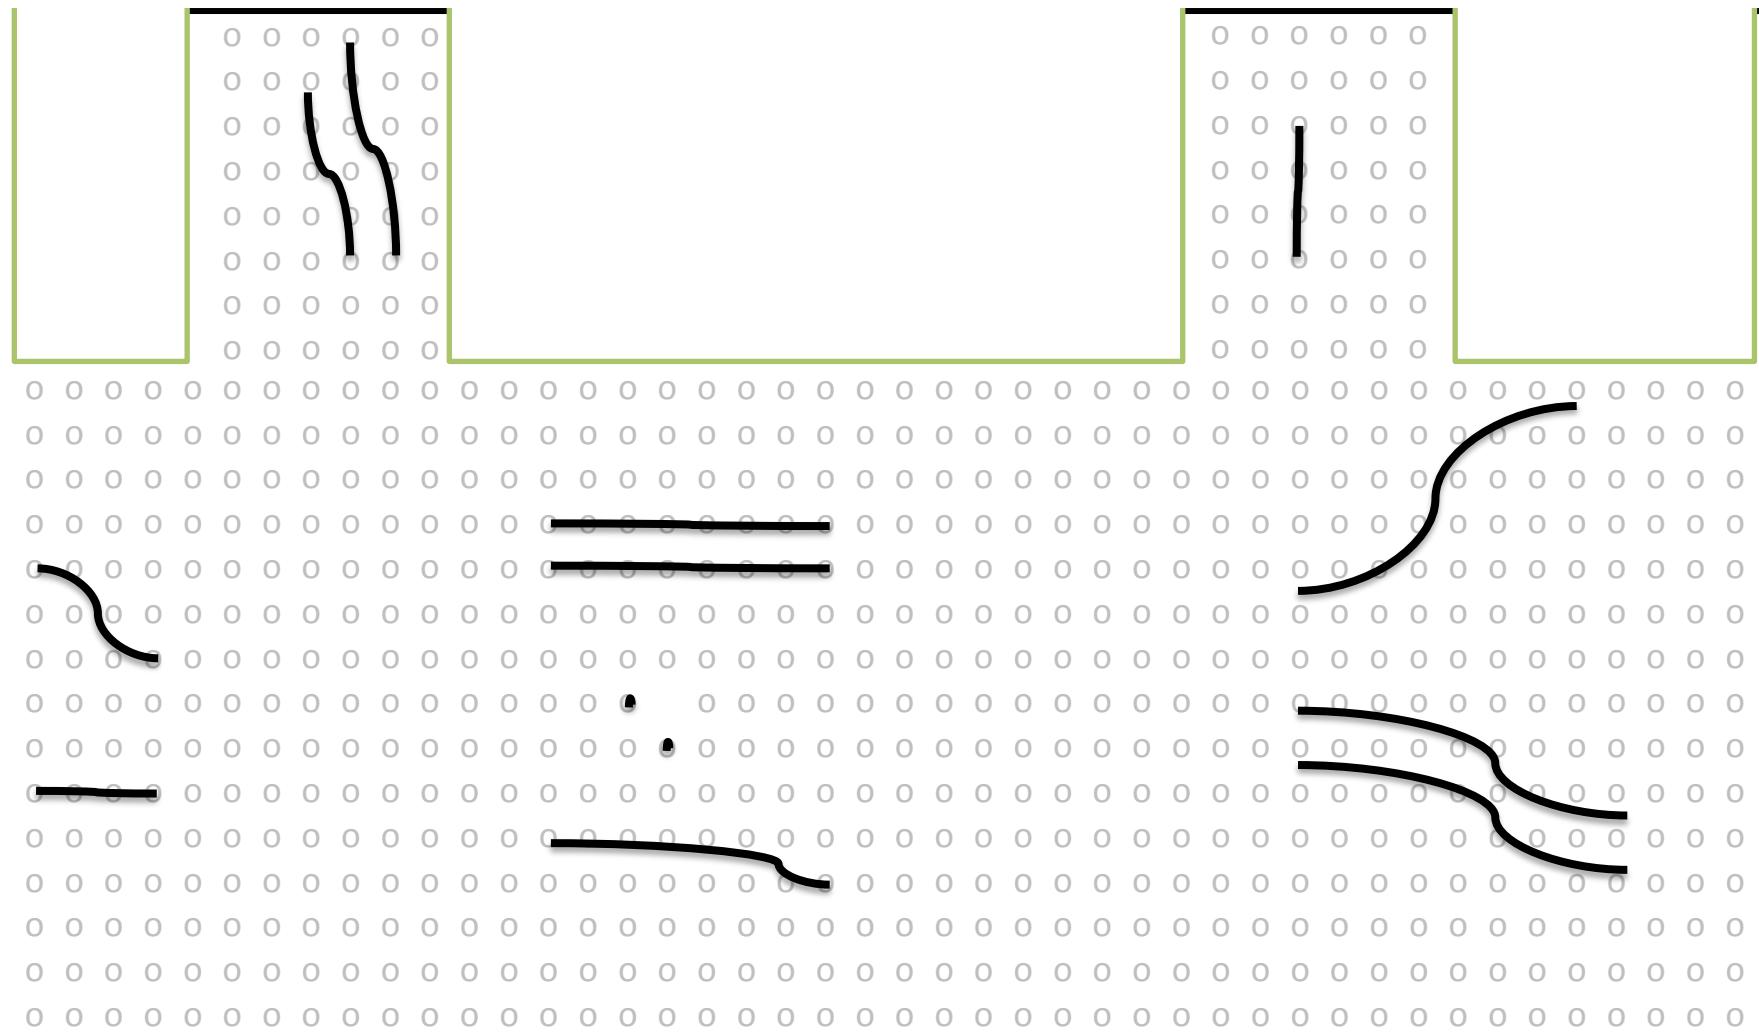
$$= \arg \max \prod_i P(t_i | \mathbf{T})P(\mathbf{T}) \quad (2),$$

where $P(\mathbf{T}) = \prod_k P(\mathbf{T}_k)$ (since trajectories do not overlap)

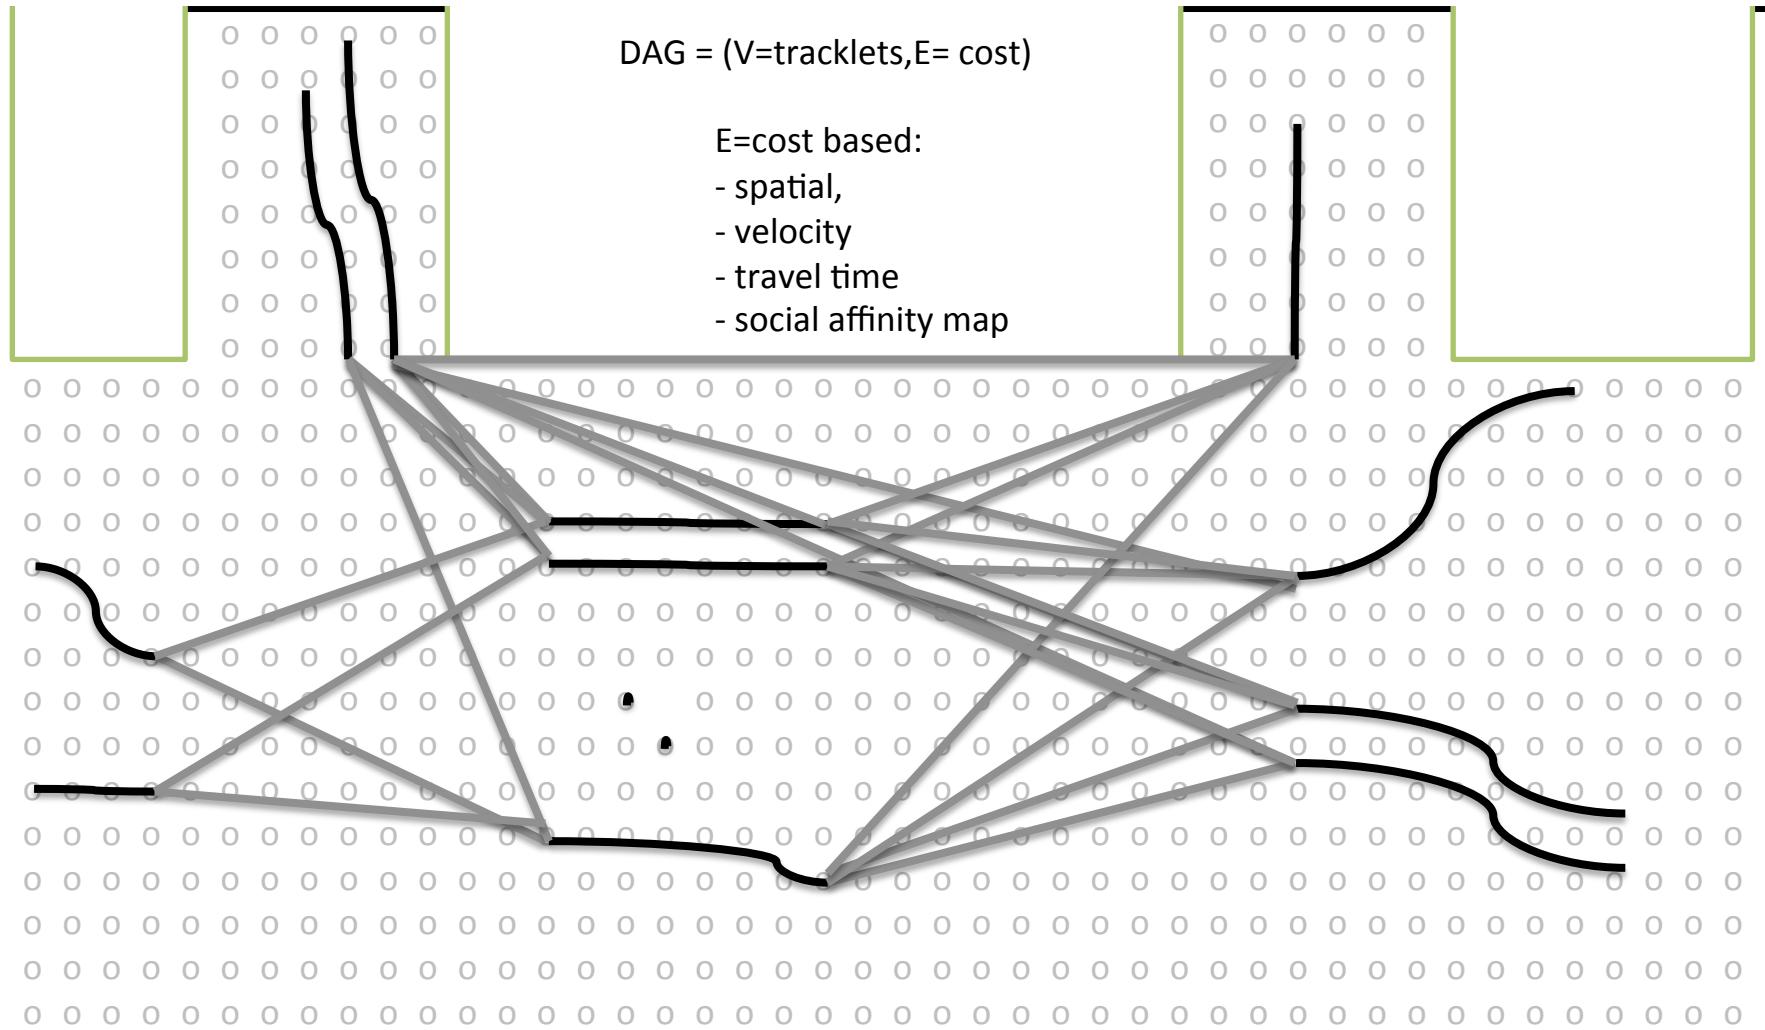
$$P(\mathbf{T}_k) = P(t_k^s) \dots P(t_k^t | t_{k-1}^t) P(t_k^e) \quad (\text{markov chain})$$



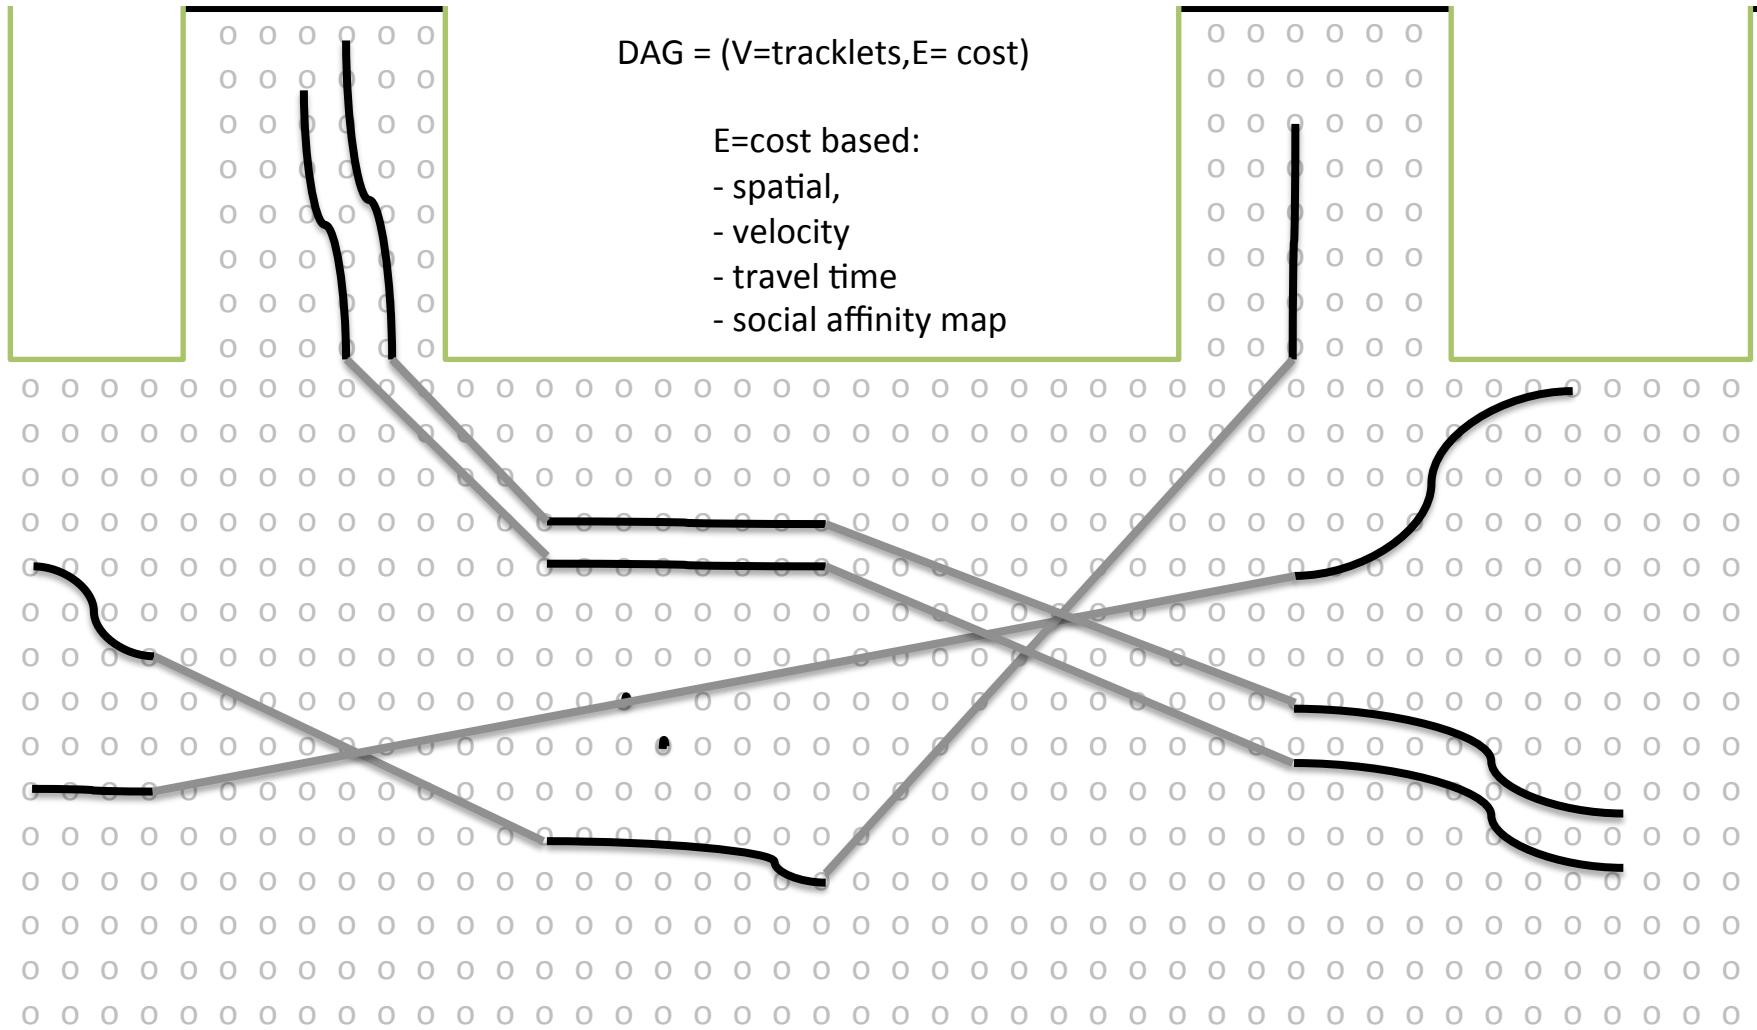
III. Tracklet association (Top view)



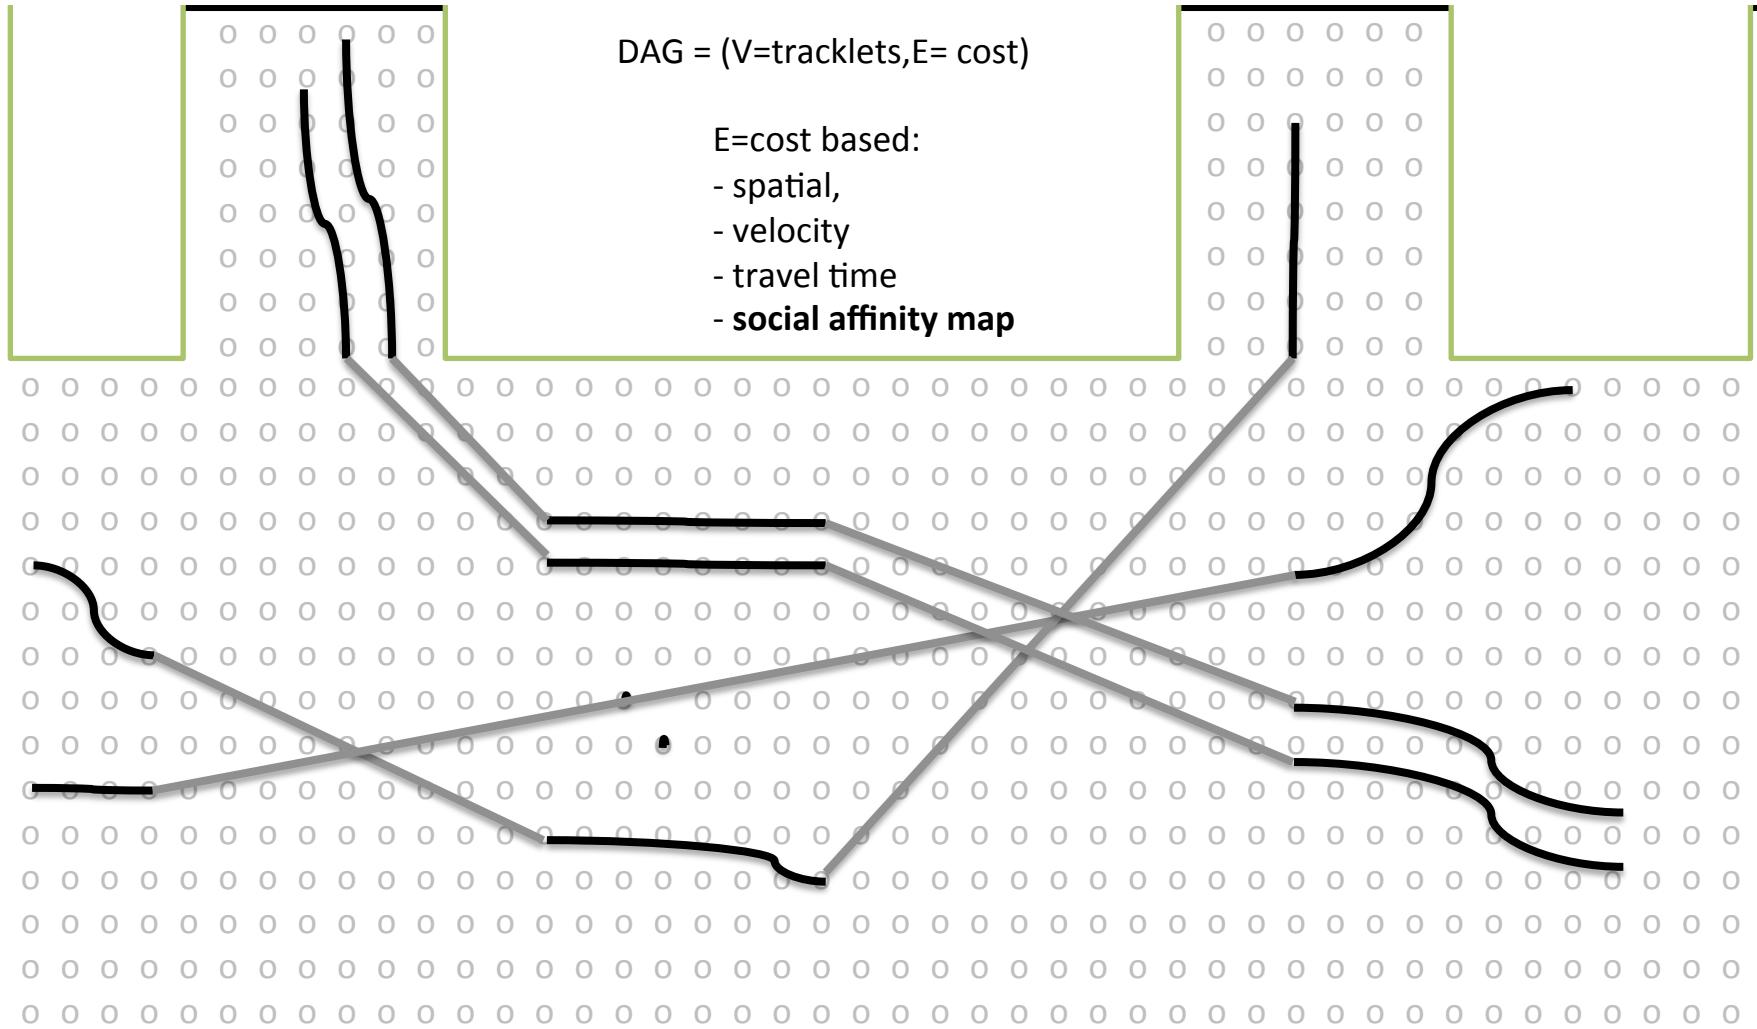
III. Tracklet association (Top view)



III. Tracklet association (Top view)



III. Tracklet association (Top view)

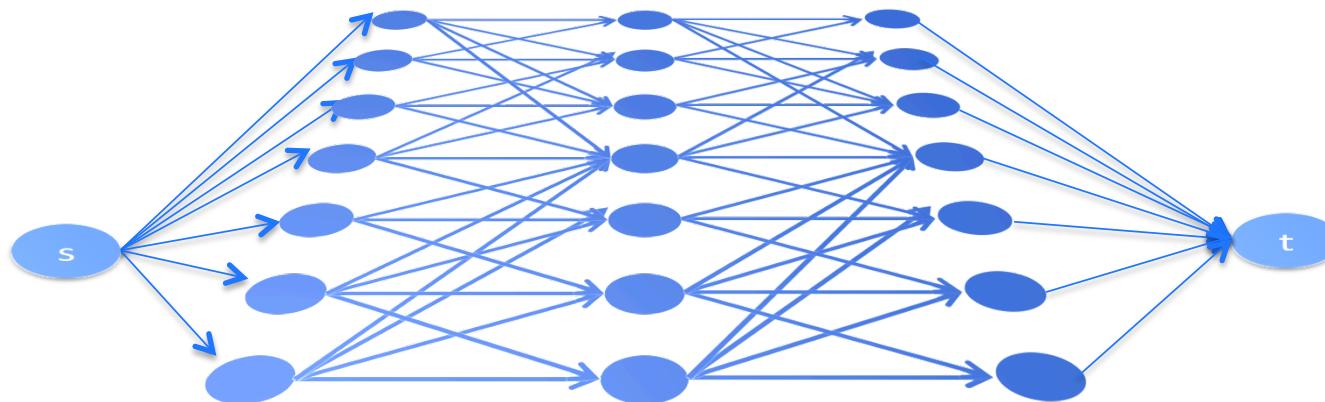


Network flow optimization

Objective: minimum cost maximum flow

$$\arg \min_f c(f)$$

$$c(f) = \sum \alpha_i f_i + \sum \beta_{ij} f_{ij}$$



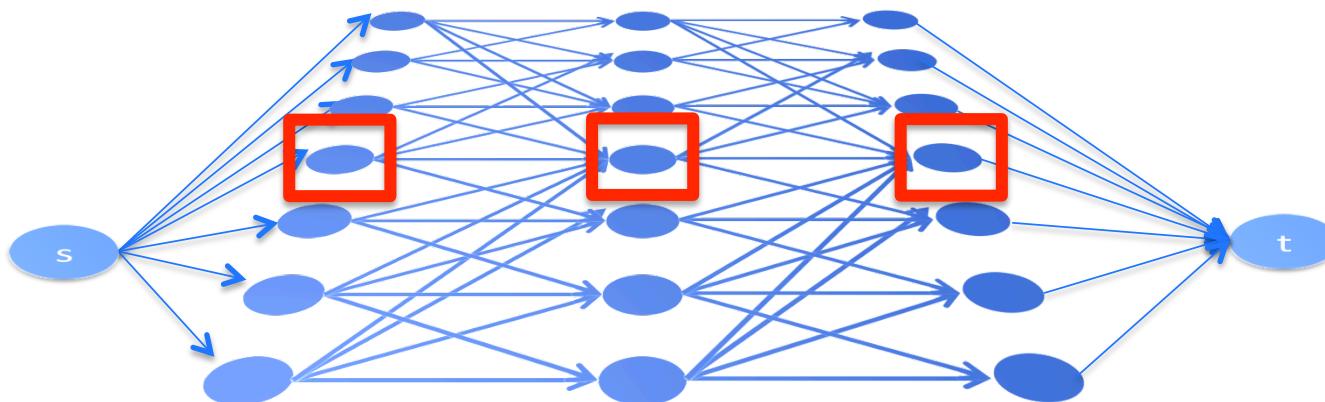
Where $\alpha_i, \beta_{ij}, \gamma_{OD}$ are the costs,
and f_i the flows



Network flow optimization

Objective: minimum cost maximum flow

$$\arg \min_f c(f)$$
$$c(f) = \sum \alpha_i f_i + \sum \beta_{ij} f_{ij}$$



Cost α_i based:
- Detection likelihood

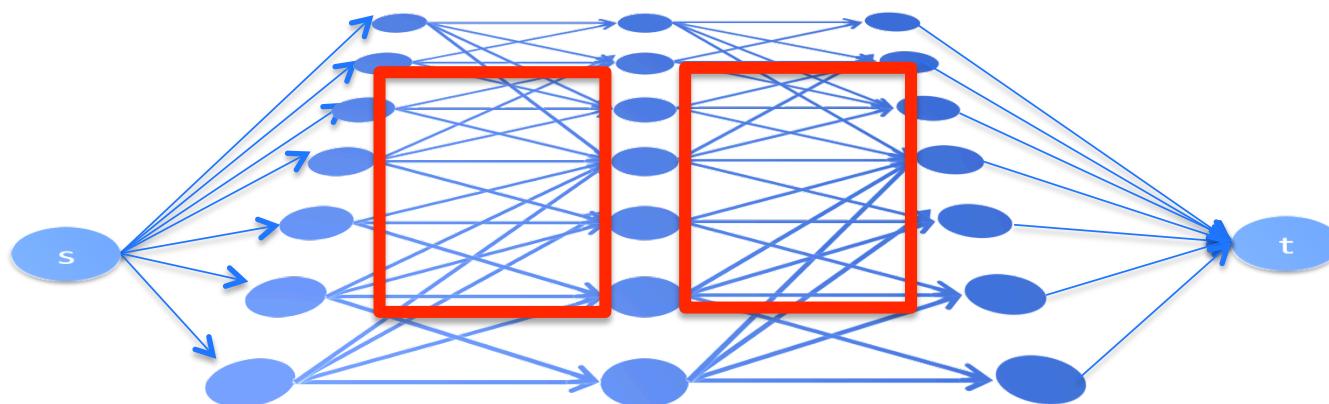


Network flow optimization

Objective: minimum cost maximum flow

$$\arg \min_f c(f)$$

$$c(f) = \sum \alpha_i f_i - \boxed{\sum \beta_{ij} f_{ij}}$$

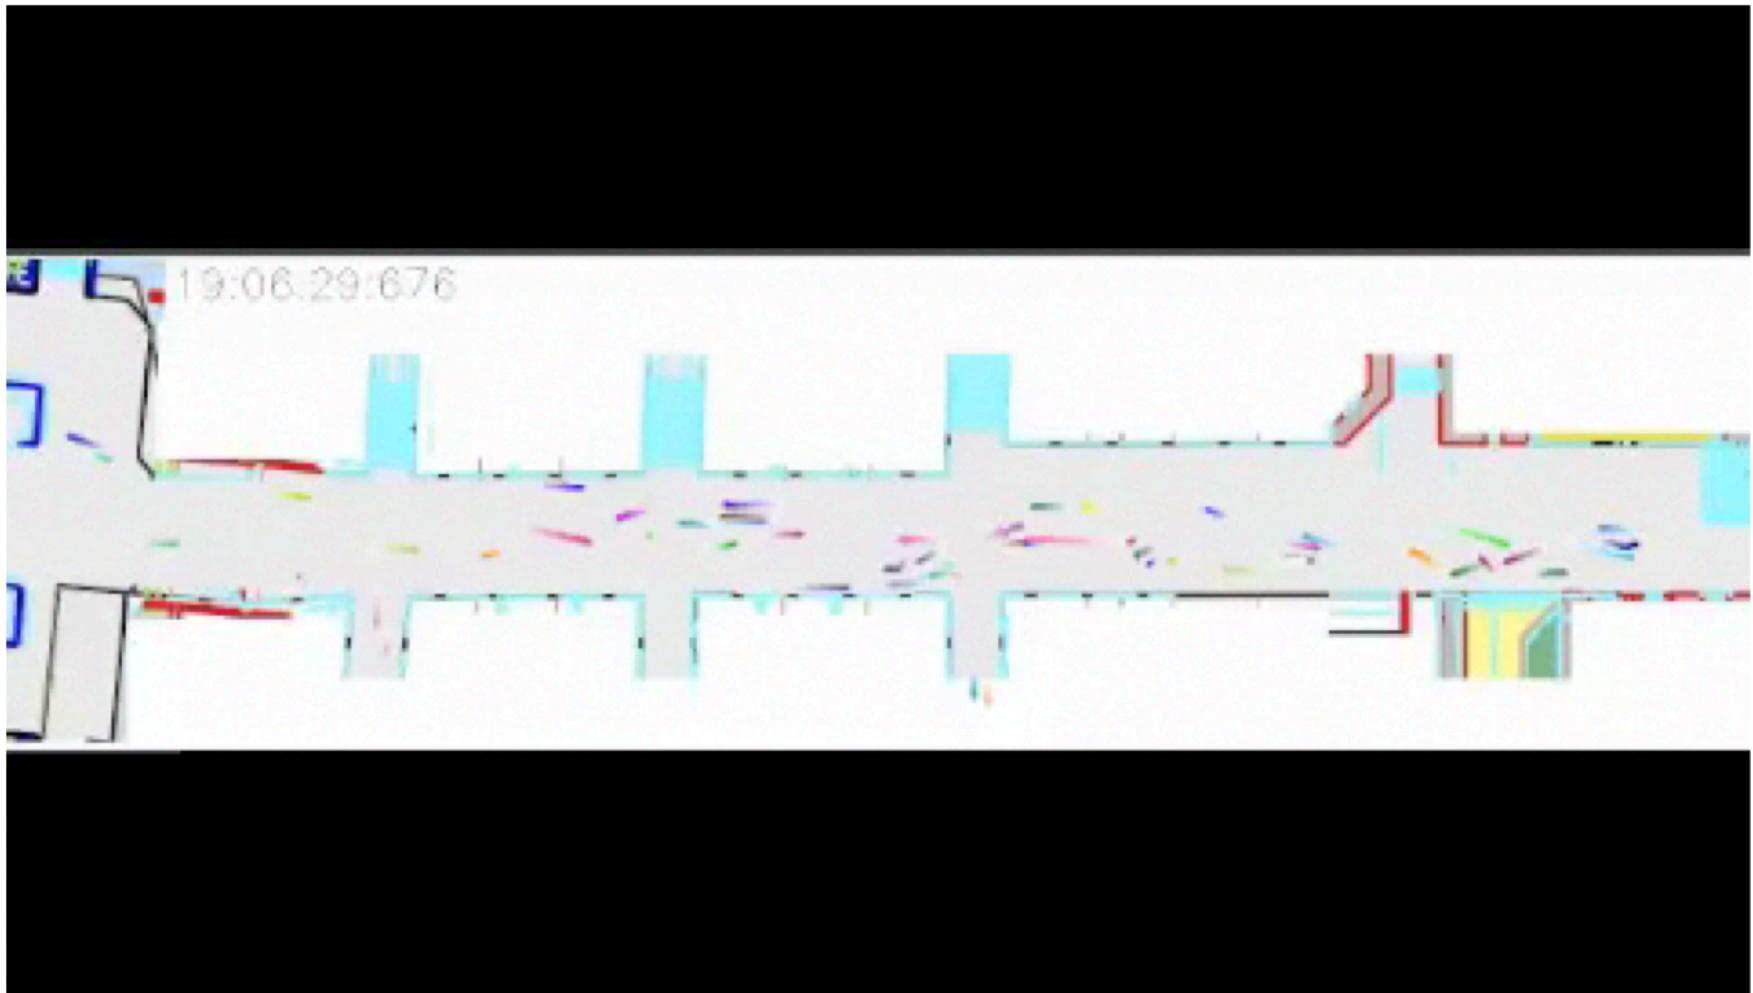


Cost β_{ij} based:

- spatial
- velocity
- **Social Affinity Map**



Tracklet association With Social Affinity Map



Conclusion

A new dimension to “Google Analytics”:
Analyzing people outside of website

