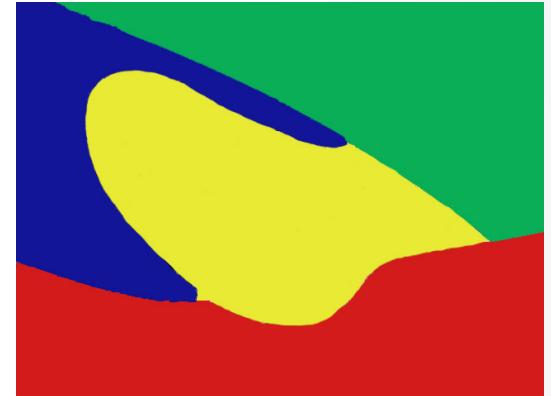
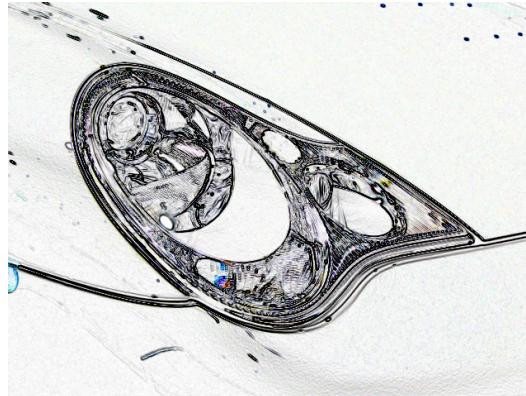


CSE578: Computer Vision

Spring 2017:

MRF for Segmentation



Anoop M. Namboodiri

Center for Visual Information Technology

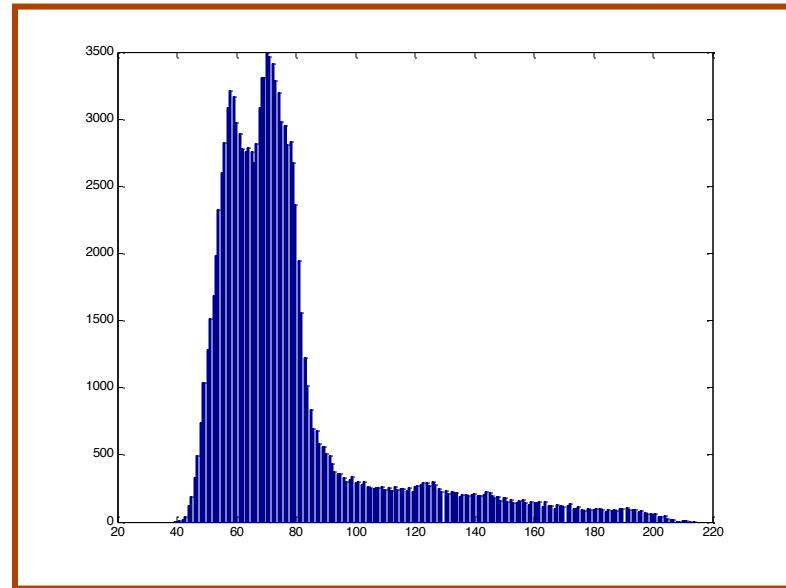
IIIT Hyderabad, INDIA

Optimal Thresholding

- The graylevel histogram is approximated using a mixture of two gaussian distributions and set the threshold to minimize the segmentation error



Grayscale Image



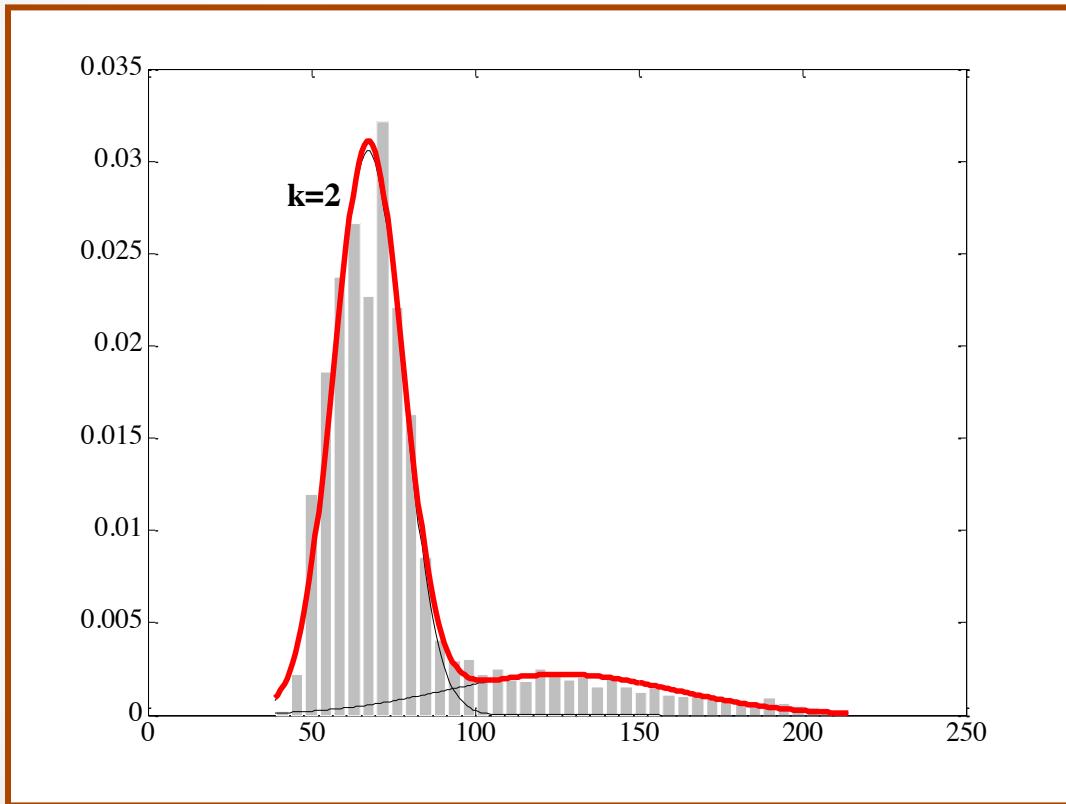
Histogram

Gaussian Mixture Estimation by EM

- Obj: =
$$N(\mu_1, \sigma_1) = \frac{1}{\sigma_1 \sqrt{2\pi}} e^{\frac{(x-\mu_1)^2}{2\sigma_1^2}}$$
- Bkg: =
$$N(\mu_2, \sigma_2) = \frac{1}{\sigma_2 \sqrt{2\pi}} e^{\frac{(x-\mu_2)^2}{2\sigma_2^2}}$$

- Initialize μ_1 , σ_1 , μ_2 , and σ_2 .
- E-Step: Computed the expected pixel label assignments. This could be either hard or soft assignment.
- M-Step: Computed Maximum-(Log)Likelihood estimates of the parameters: $\mu_1, \sigma_1, \mu_2, \sigma_2$
- Repeat the E and M steps until convergence

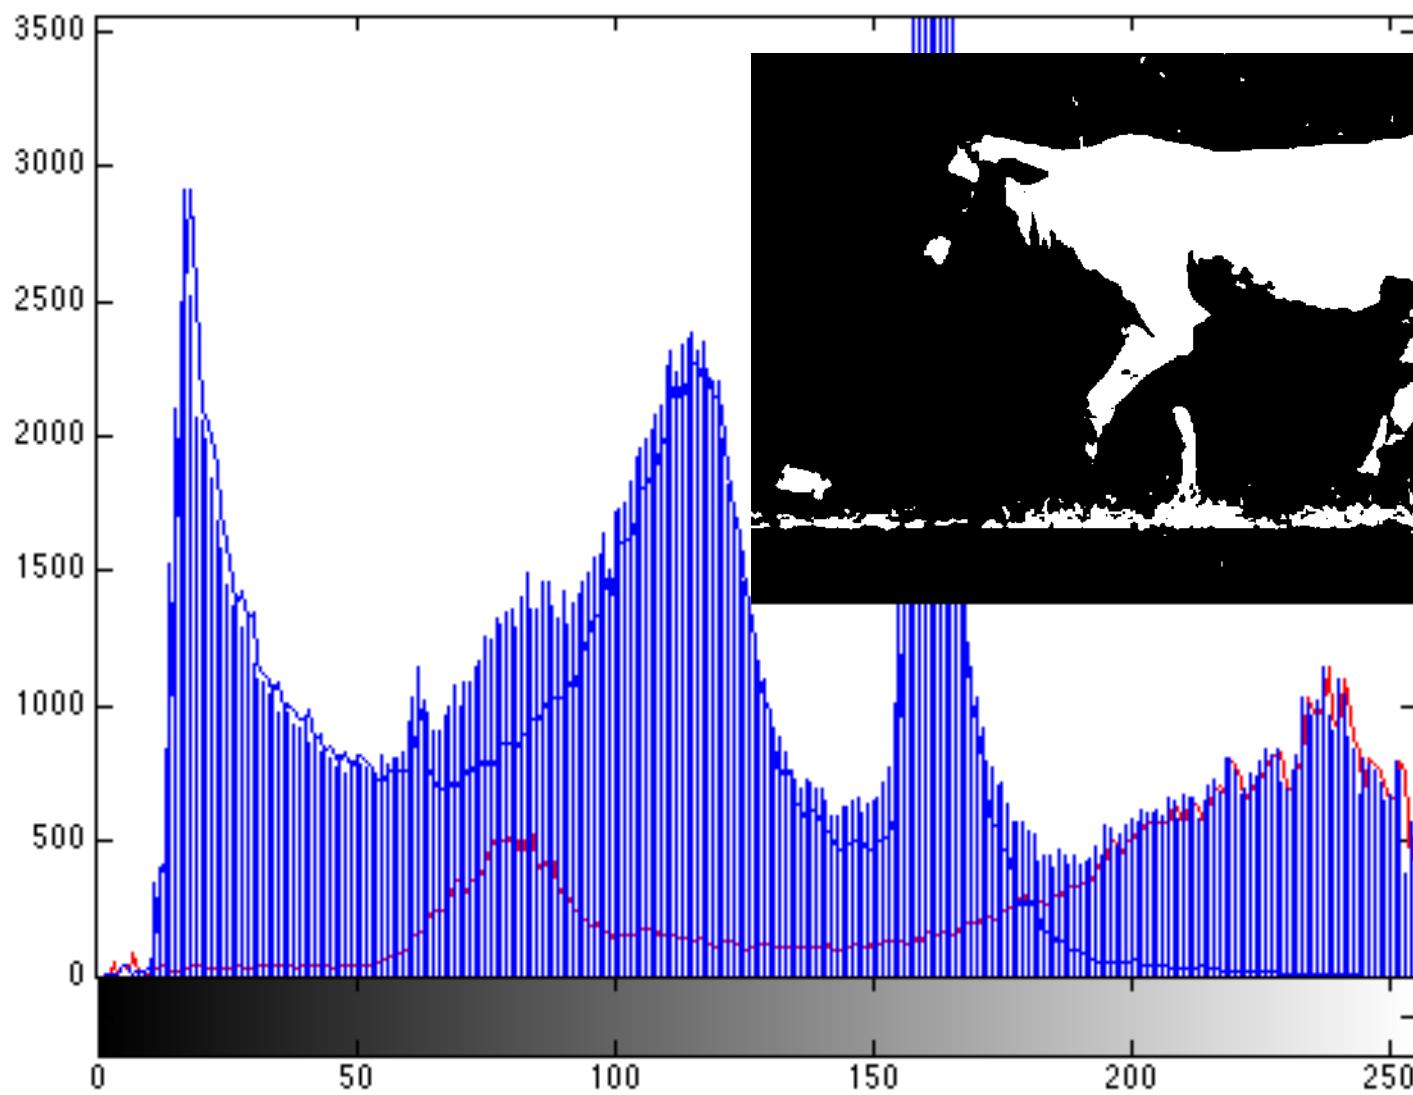
Optimal Thresholding



Histogram with bimodal fit



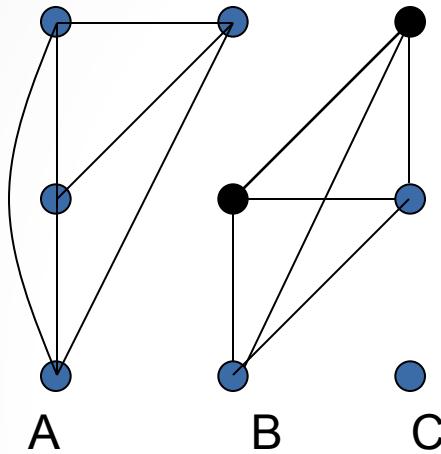
Thresholded ($T=94$)



Segmentation as Optimal Labeling

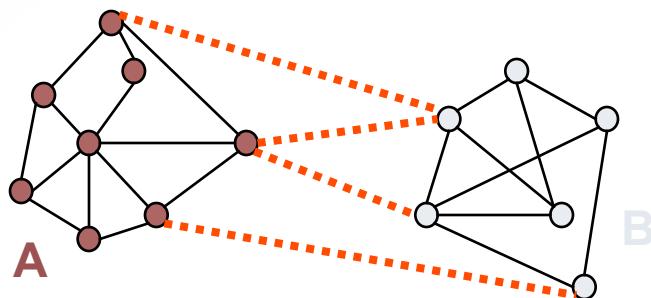
- Model knowledge about the world
- Classify each pixel as belonging to a specific object
 - Independent classification does not work
 - Need to incorporate neighborhood information
- Consider a graph over the image
 - Each node in the graph need to be labeled
 - Edges in the graph represent neighborhood constraints
- Define a cost function, $Q(f)$, using the above
- Compute the optimal labeling wrt $Q(f)$.

Segmentation by Graph Cuts



- Break Graph into Segments
 - Delete links that cross between segments
 - Easiest to break links that have low cost (low similarity)
 - similar pixels should be in the same segments
 - dissimilar pixels should be in different segments

Cuts in a graph



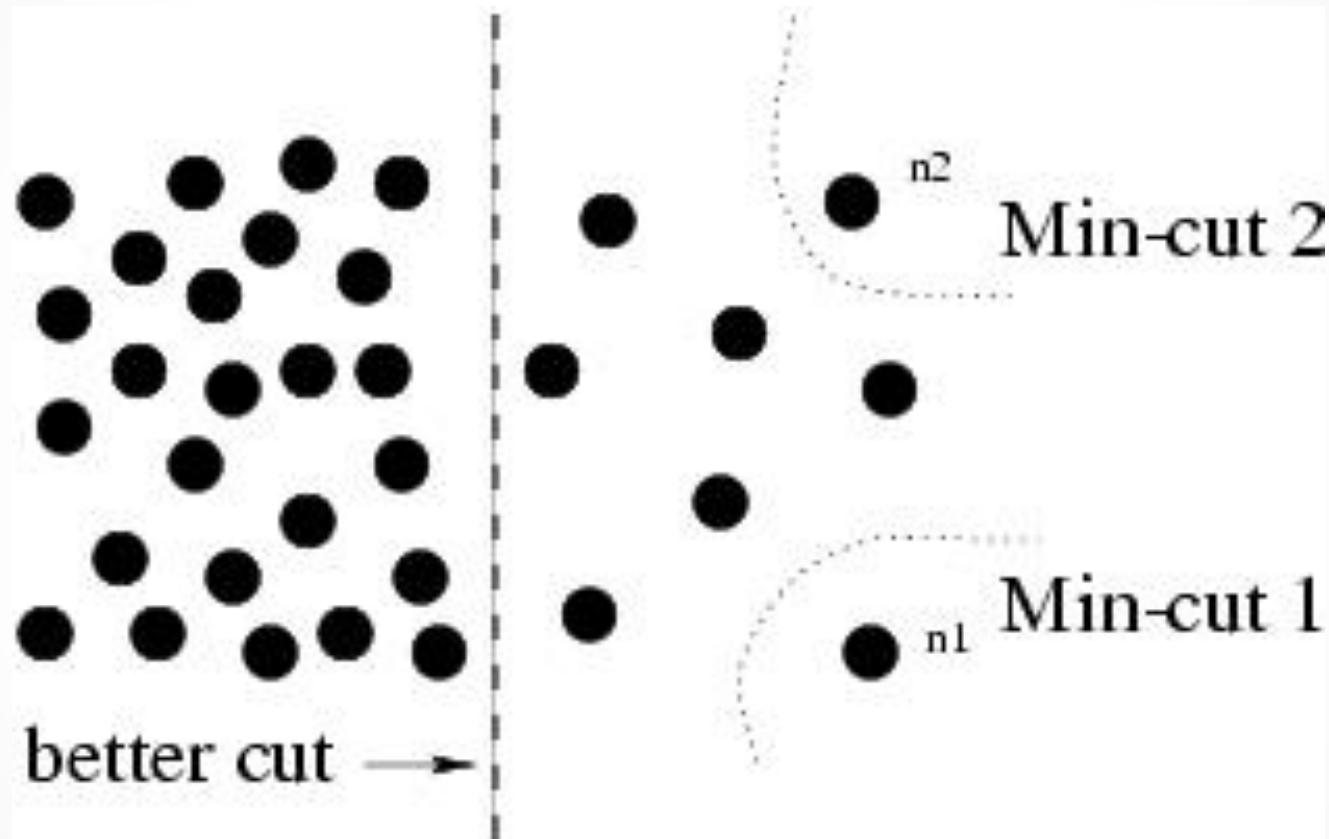
- Link Cut
 - set of links whose removal makes a graph disconnected
 - cost of a cut:

$$cut(A, B) = \sum_{u \in A, v \in B} w(u, v).$$

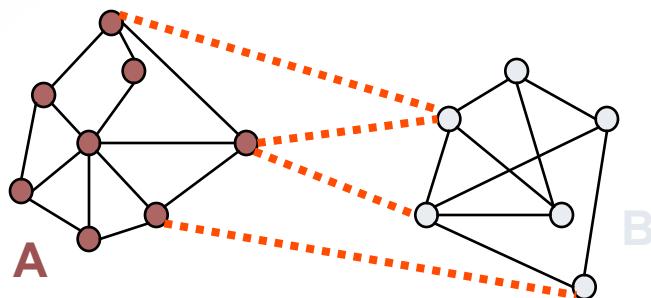
One idea: Find minimum cut

- gives you a segmentation
- fast algorithms exist for doing this

Min cut is not always the best



Cuts in a graph



Normalized Cut

- a cut penalizes large segments
- fix by normalizing for size of segments

$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)},$$

where $assoc(A, V) = \sum_{u \in A, t \in V} w(u, t)$

Recursive normalized cuts

1. Given an image or image sequence, set up a weighted graph: $G=(V, E)$
 - Vertex for each pixel
 - Edge weight for nearby pairs of pixels
 2. Solve for eigenvectors with the smallest eigenvalues:
 $(\mathbf{D} - \mathbf{W})\mathbf{y} = \lambda \mathbf{D}\mathbf{y}$, where
 - Use the eigenvector with the second smallest eigenvalue to bipartition the graph
 - Note: this is an approximation
 3. Recursively repartition the segmented parts if necessary
- Details: <http://www.cs.berkeley.edu/~malik/papers/SM-ncut.pdf>

Normalized cuts results



Graphcut for Image Segmentation



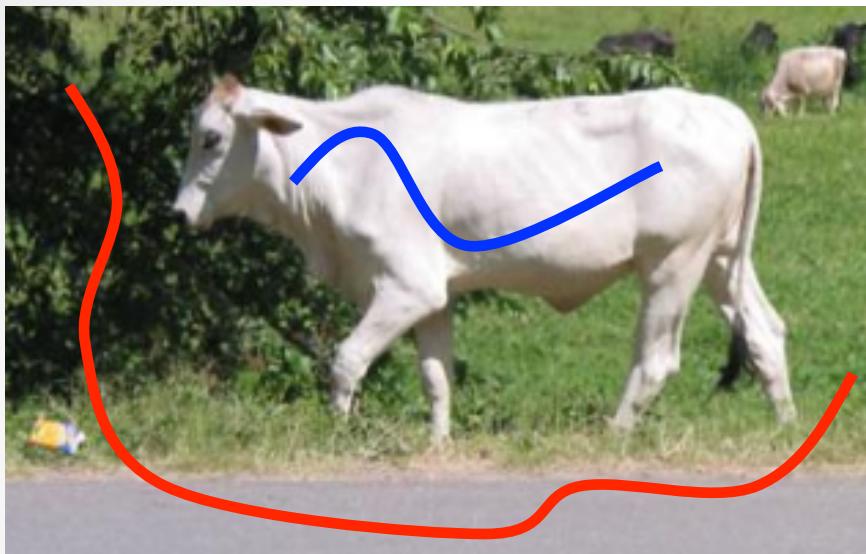
How ?

Cost function Models *our* knowledge about natural images

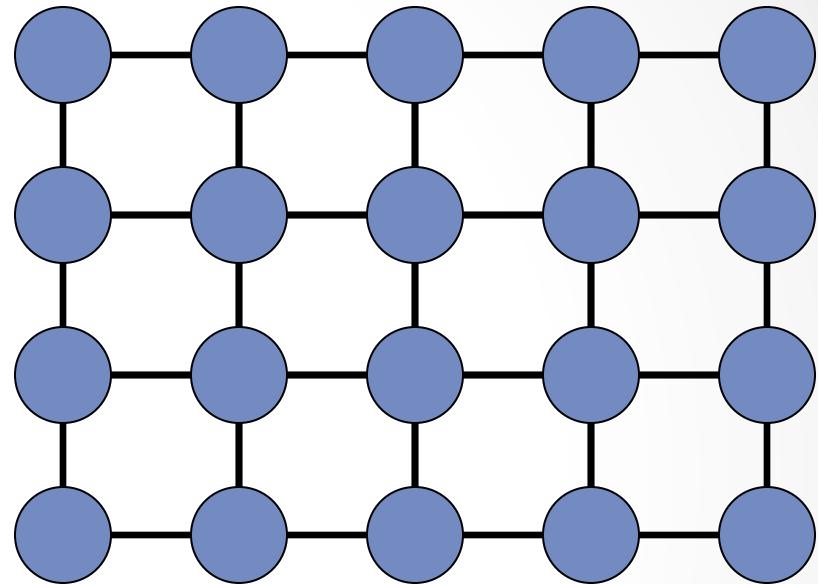
Optimize cost function to obtain the segmentation



Binary Image Segmentation



Object - white, Background - green/grey



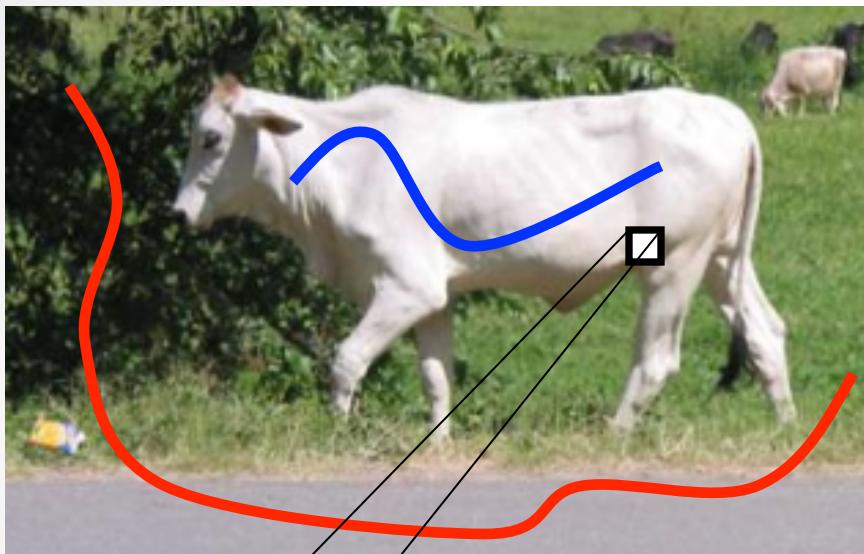
Graph $G = (V, E)$

Each vertex corresponds to a pixel

Edges define a 4-neighbourhood *grid* graph

Assign a label to each vertex from $L = \{\text{obj}, \text{bkg}\}$

Binary Image Segmentation

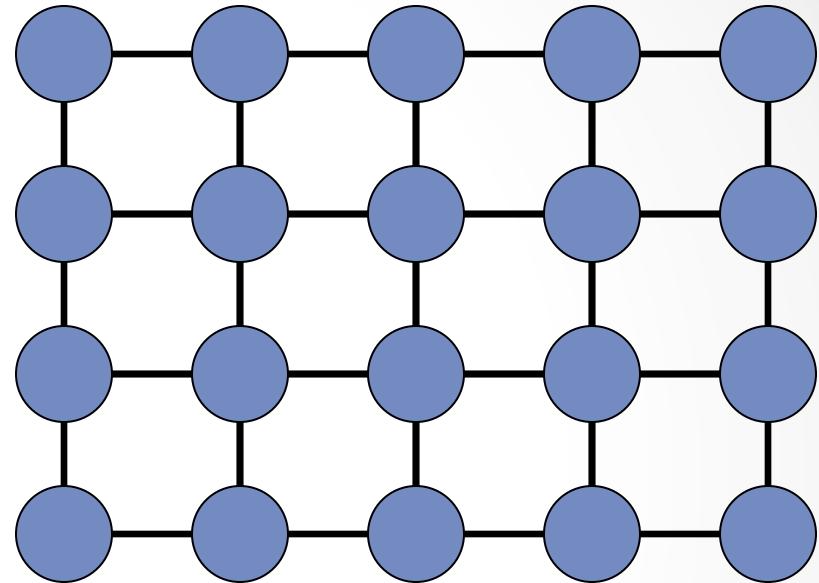


Object - white, Background - green/grey

Cost of a labelling $f : V \rightarrow L$



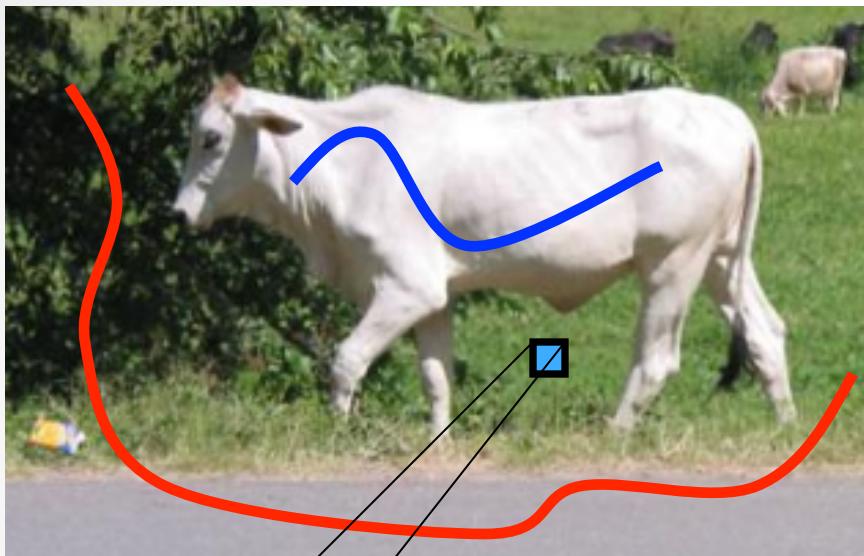
Cost of label ‘obj’ low Cost of label ‘bkg’ high



Graph $G = (V, E)$

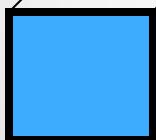
Per Vertex Cost

Binary Image Segmentation

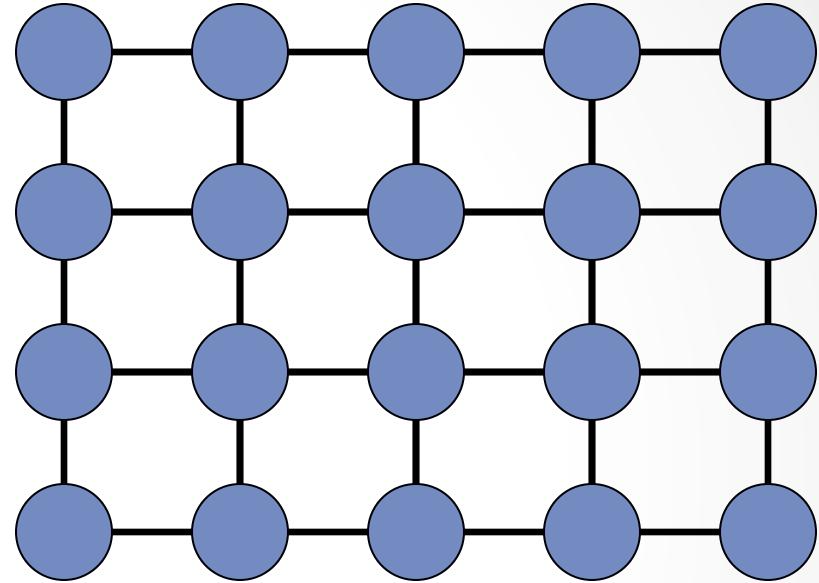


Object - white, Background - green/grey

Cost of a labelling $f : V \rightarrow L$



Cost of label ‘obj’ high Cost of label ‘bkg’ low

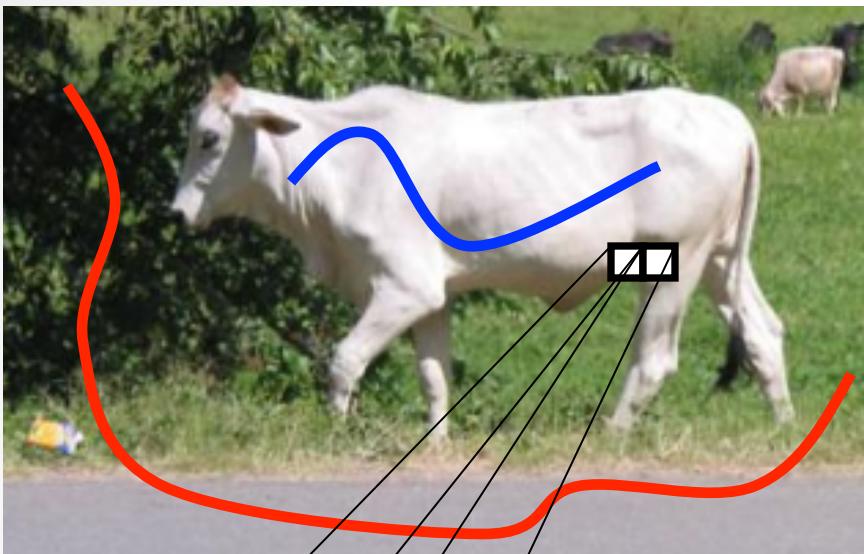


Graph $G = (V, E)$

Per Vertex Cost

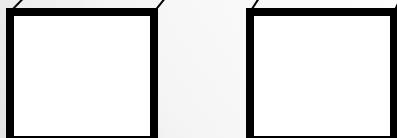
UNARY COST

Binary Image Segmentation



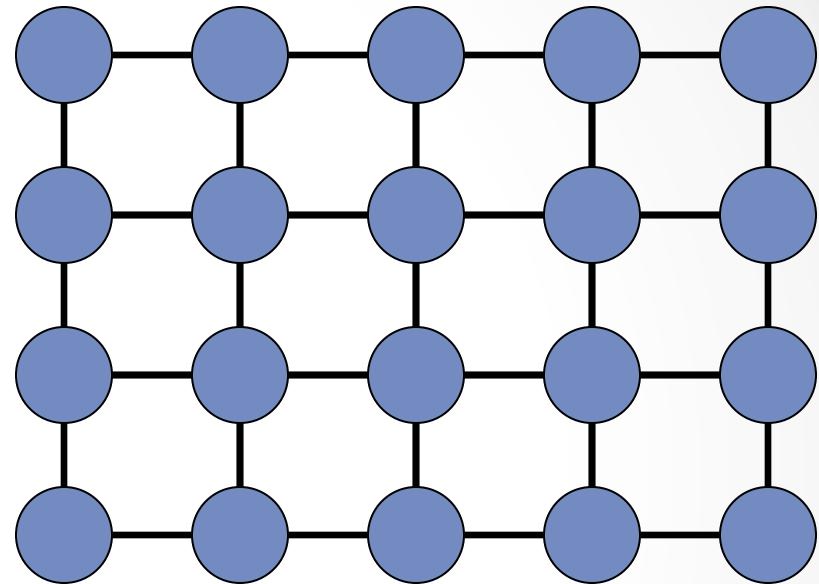
Object - white, Background - green/grey

Cost of a labelling $f : V \rightarrow L$



Cost of same label low

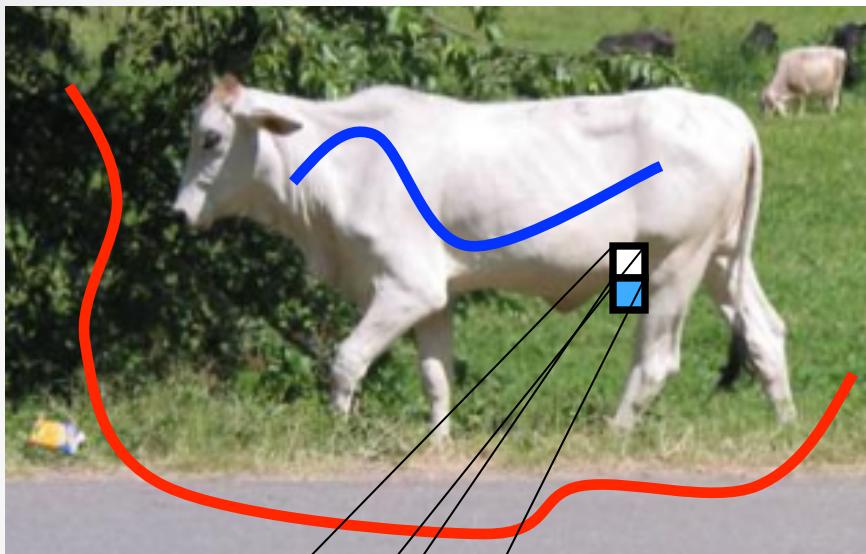
Cost of different labels high



Graph $G = (V, E)$

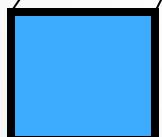
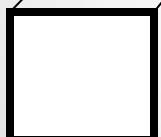
Per Edge Cost

Binary Image Segmentation



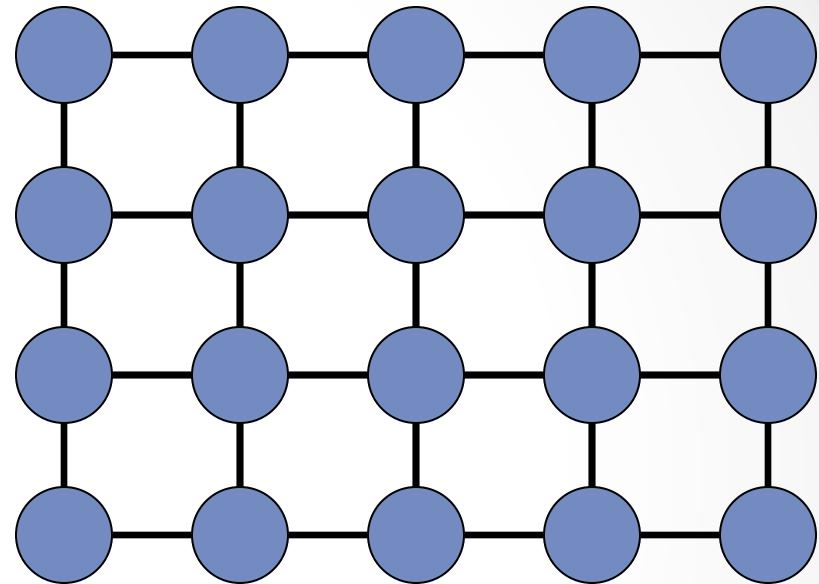
Object - white, Background - green/grey

Cost of a labelling $f : V \rightarrow L$



Cost of same label high

Cost of different labels low

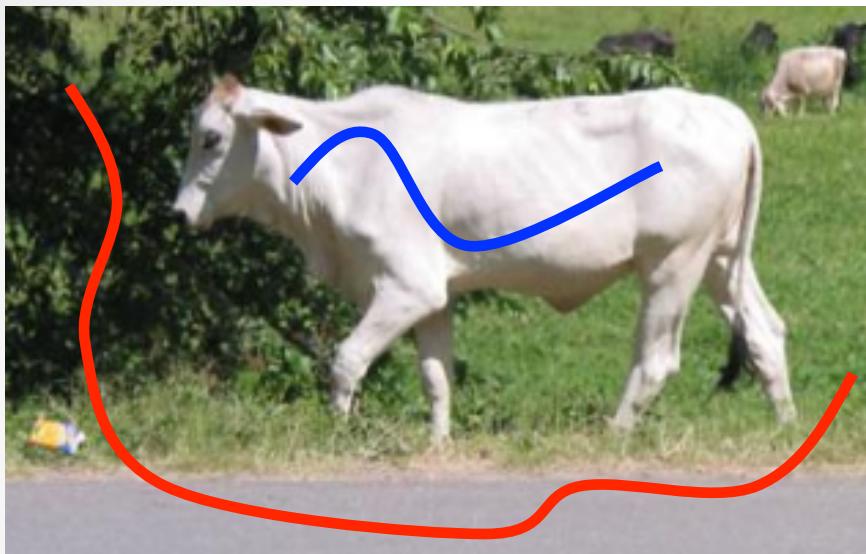


Graph $G = (V, E)$

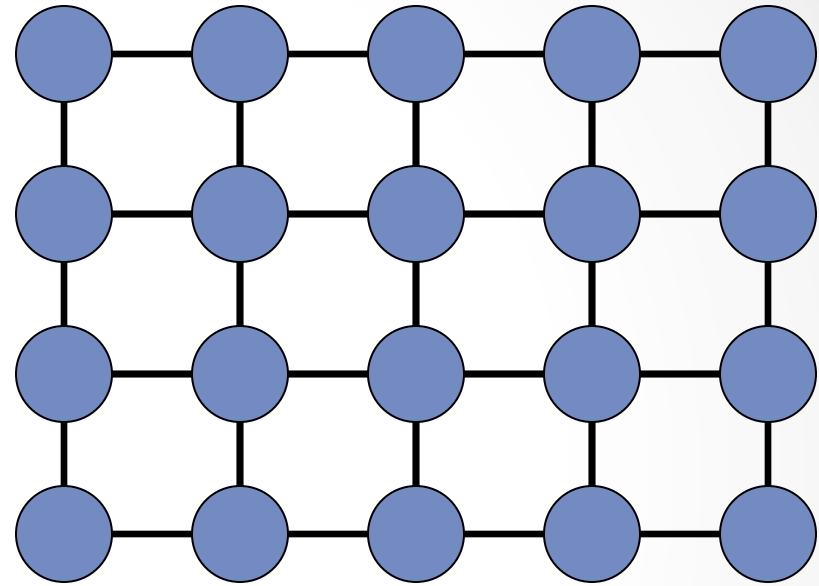
Per Edge Cost

PAIRWISE
COST

Binary Image Segmentation



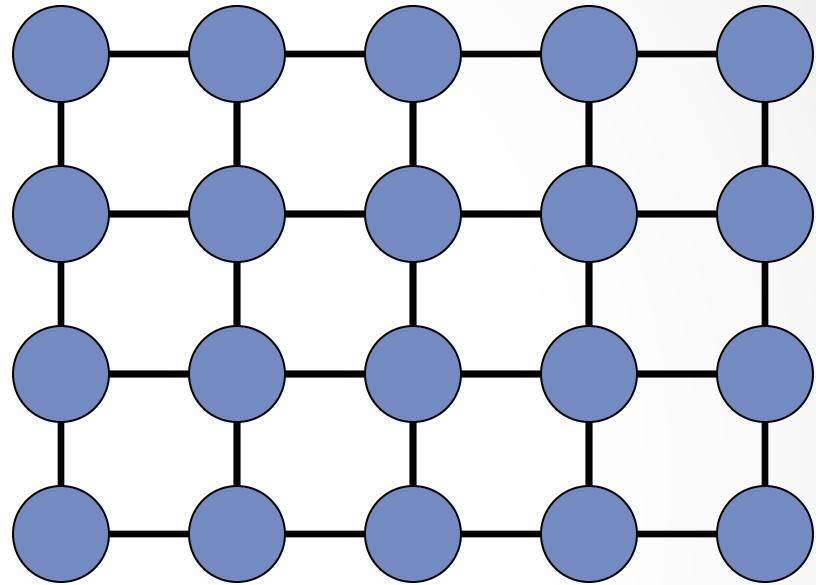
Object - white, Background - green/grey



Graph $G = (V, E)$

Problem: Find the labeling with minimum cost f^*

Binary Image Segmentation



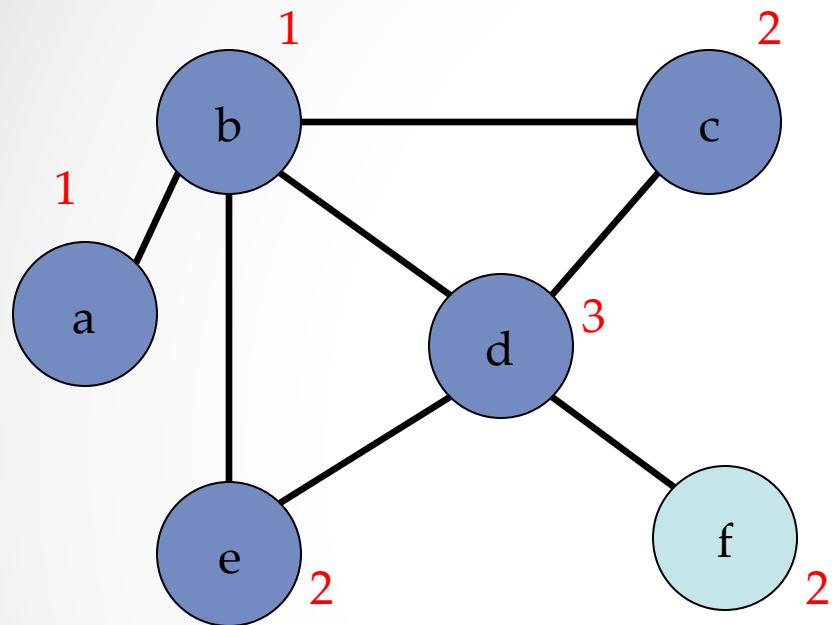
$$\text{Graph } G = (V, E)$$

$$L = \{fg, bg\}$$

Vertex corresponds to a pixel
Edges define grid graph

Problem: Find the labeling with minimum cost f^*

The General Problem



Graph $G = (V, E)$

Discrete label set $L = \{1, 2, \dots, h\}$

Assign a label to each vertex
 $f: V \rightarrow L$

Cost of a labelling $Q(f)$

Unary Cost

Pairwise Cost

Find $f^* = \arg \min Q(f)$

Formulation: Energy Function

Label l_1



Label l_0



V_a

D_a



V_b

D_b



V_c

D_c



V_d

D_d

Random Variables $V = \{V_a, V_b, \dots\}$

Labels $L = \{l_0, l_1, \dots\}$ Data D

Labelling $f: \{a, b, \dots\} \rightarrow \{0, 1, \dots\}$

Energy Function

Label l_1

2

4

6

3

Label l_0

5

2

3

7

V_a

V_b

V_c

V_d

D_a

D_b

D_c

D_d

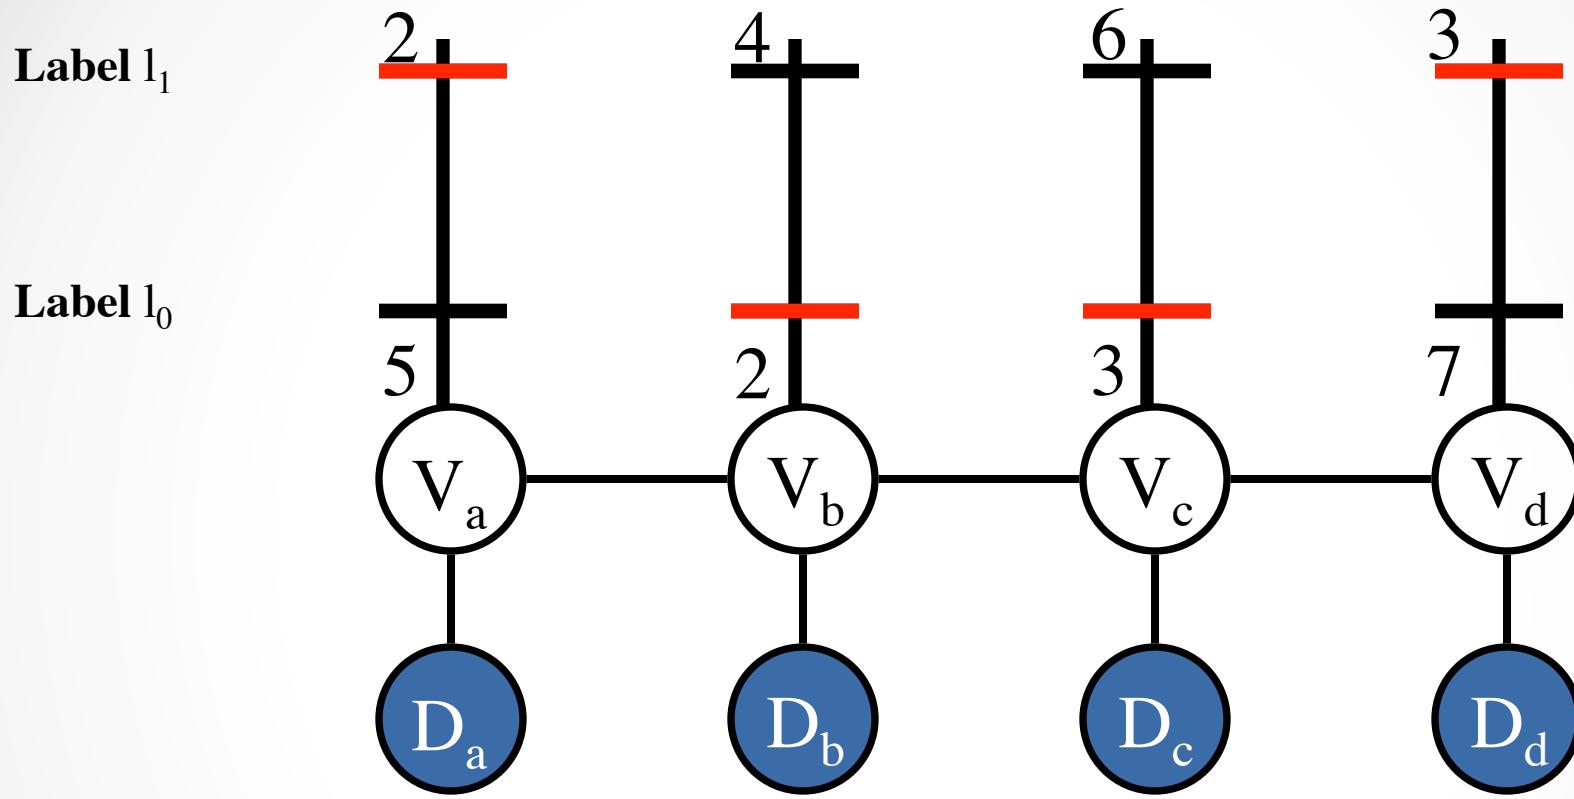
$$Q(f) = \sum_a \theta_{a;f(a)}$$

Unary Potential

Easy to minimize

Neighbourhood

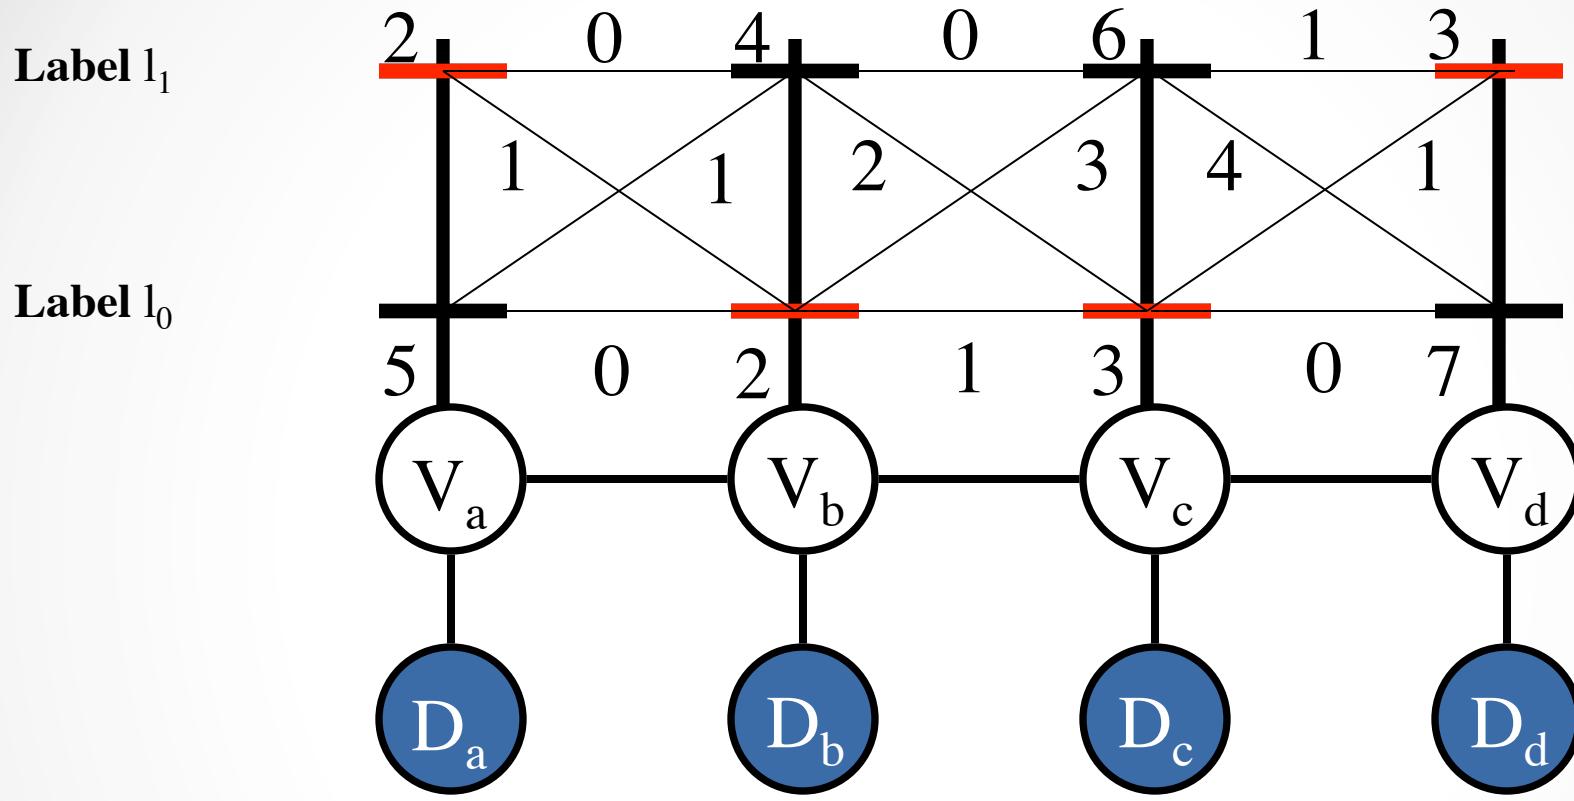
Energy Function



$E : (a,b) \in E \text{ iff } V_a \text{ and } V_b \text{ are neighbours}$

$$E = \{ (a,b), (b,c), (c,d) \}$$

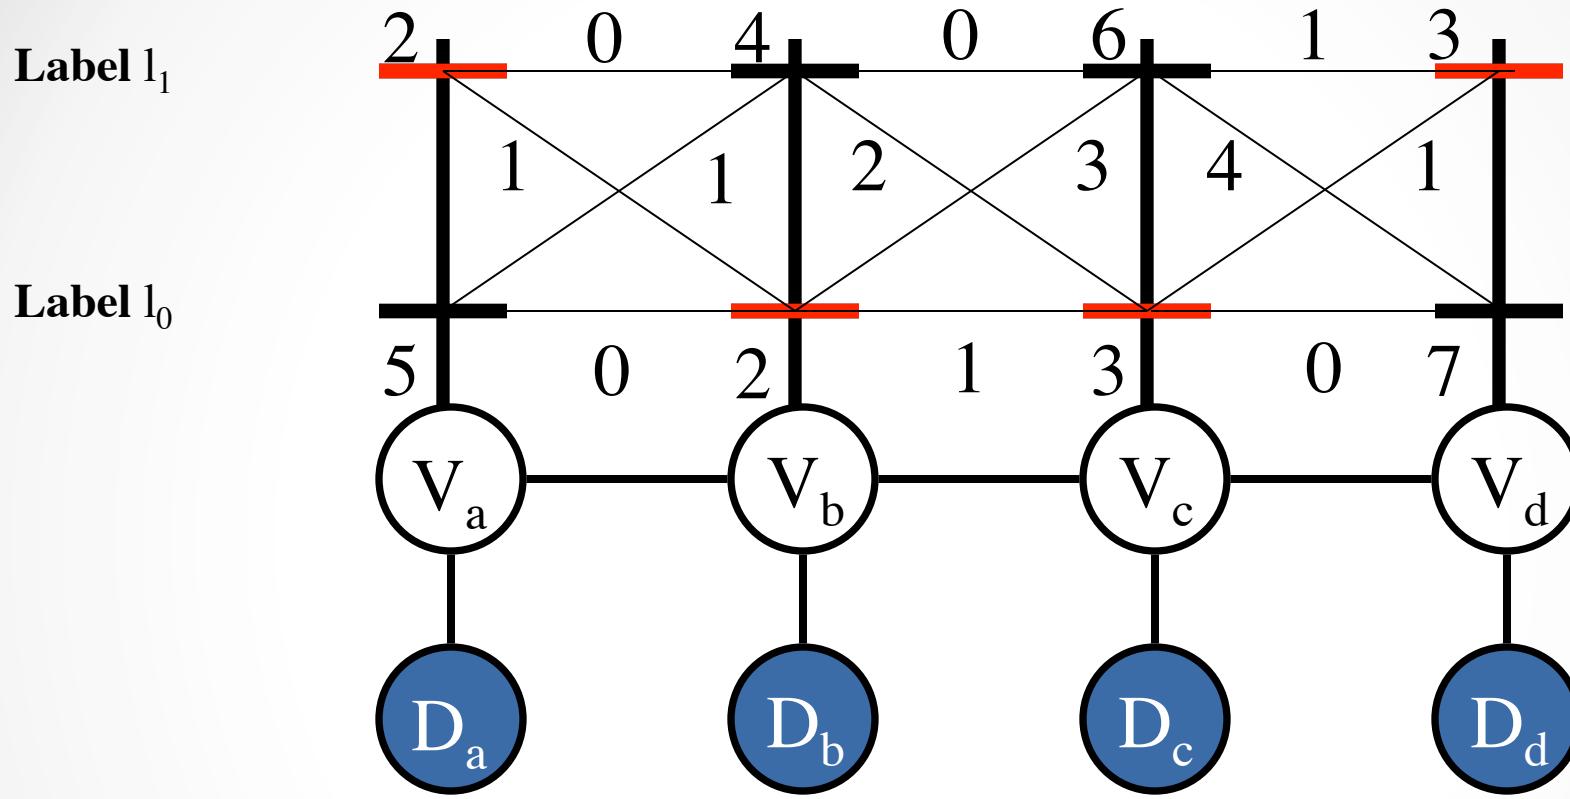
Energy Function



Pairwise Potential

$$Q(f) = \sum_a \theta_{a;f(a)} + \sum_{(a,b)} \theta_{ab;f(a)f(b)}$$

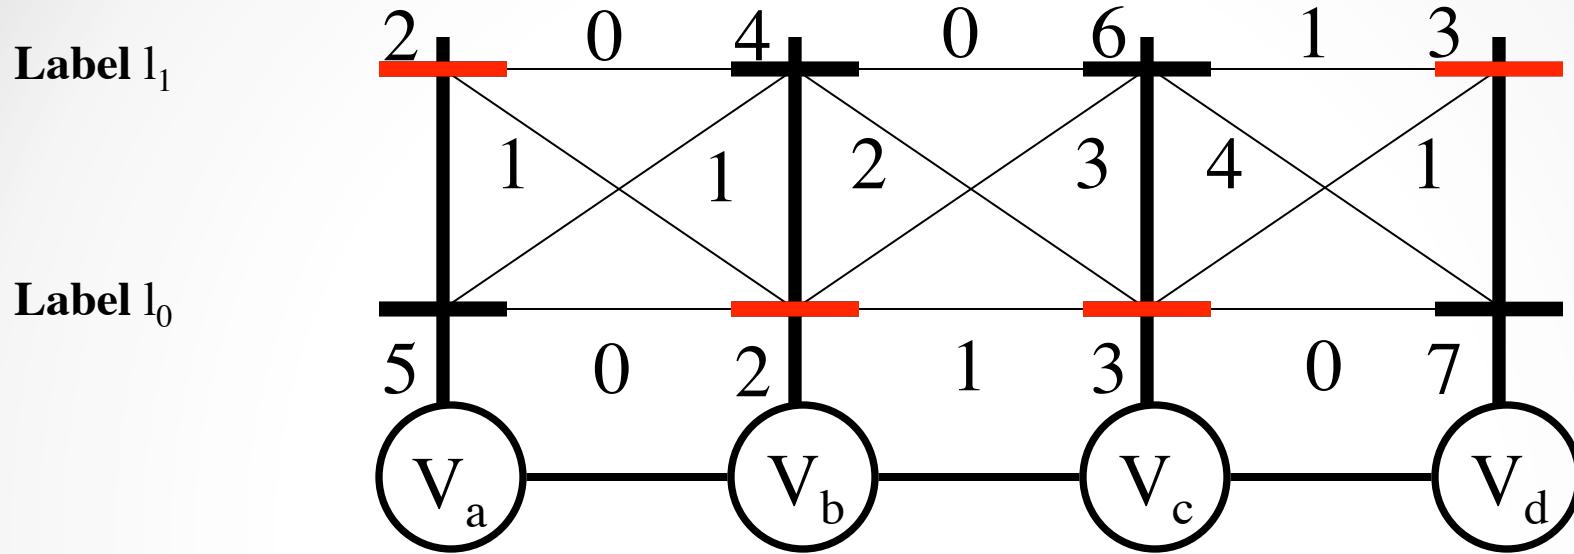
Energy Function



$$Q(f; \theta) = \sum_a \theta_{a;f(a)} + \sum_{(a,b)} \theta_{ab;f(a)f(b)}$$

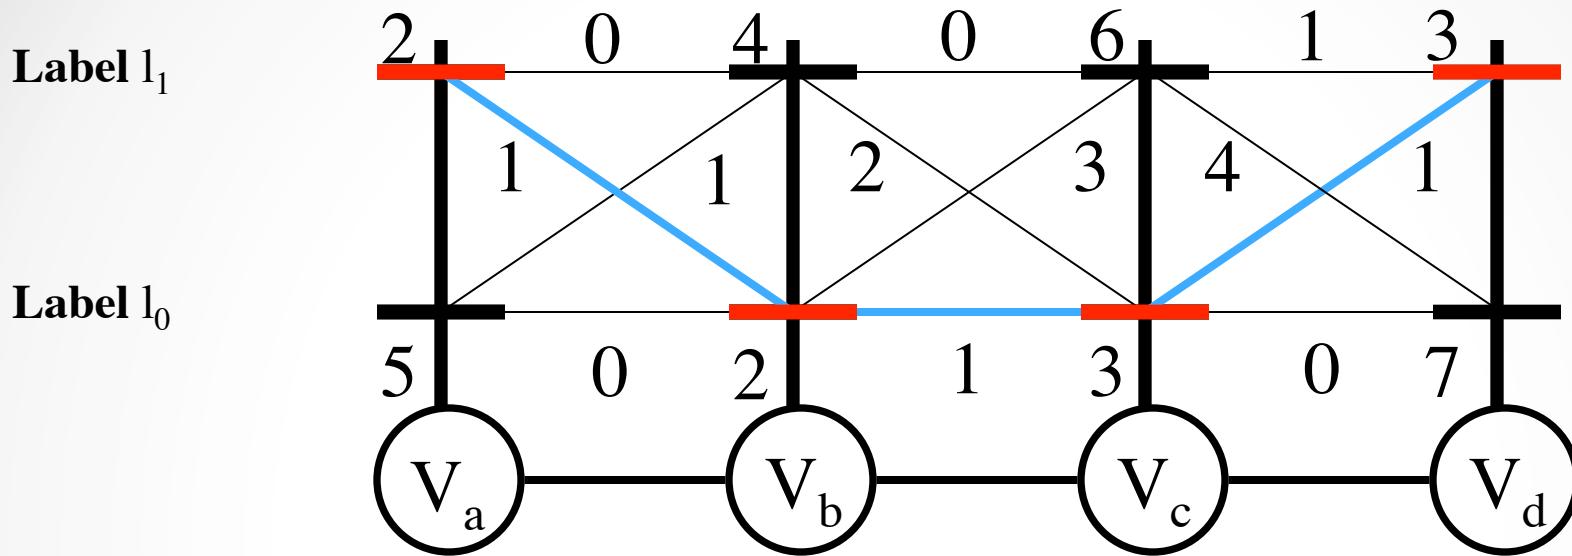
Parameter

MAP Estimation



$$Q(f; \theta) = \sum_a \theta_{a;f(a)} + \sum_{(a,b)} \theta_{ab;f(a)f(b)}$$

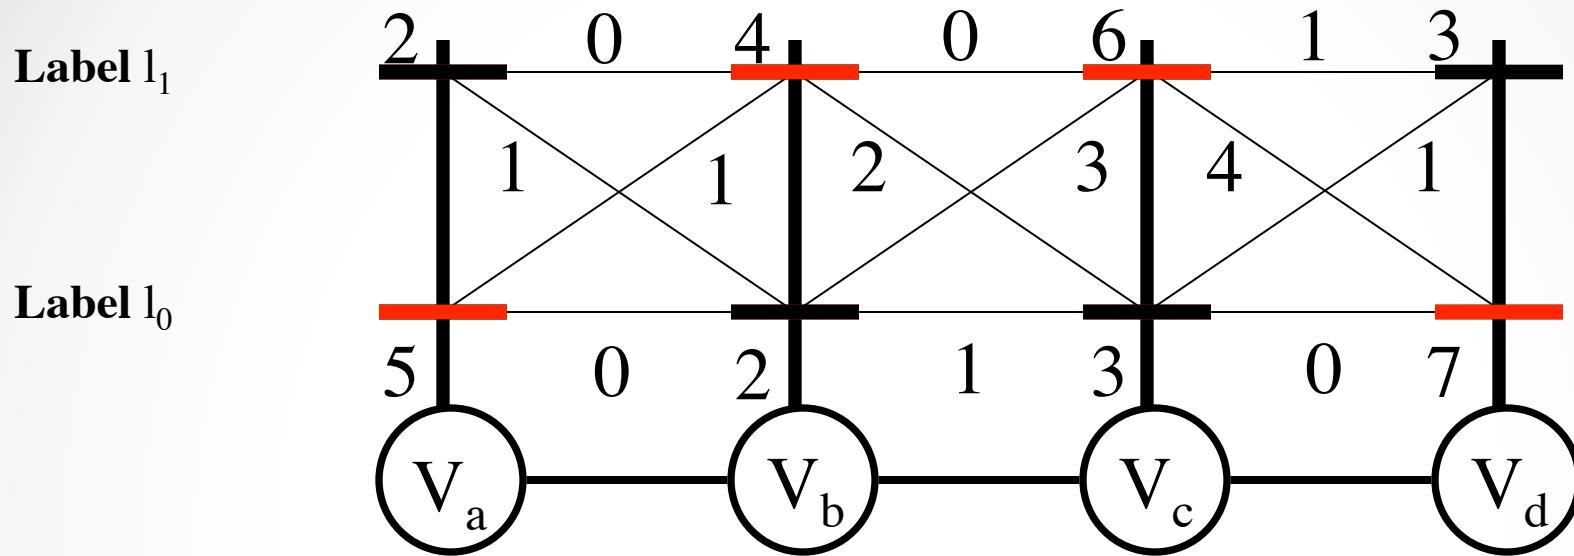
MAP Estimation



$$Q(f; \theta) = \sum_a \theta_{a;f(a)} + \sum_{(a,b)} \theta_{ab;f(a)f(b)}$$

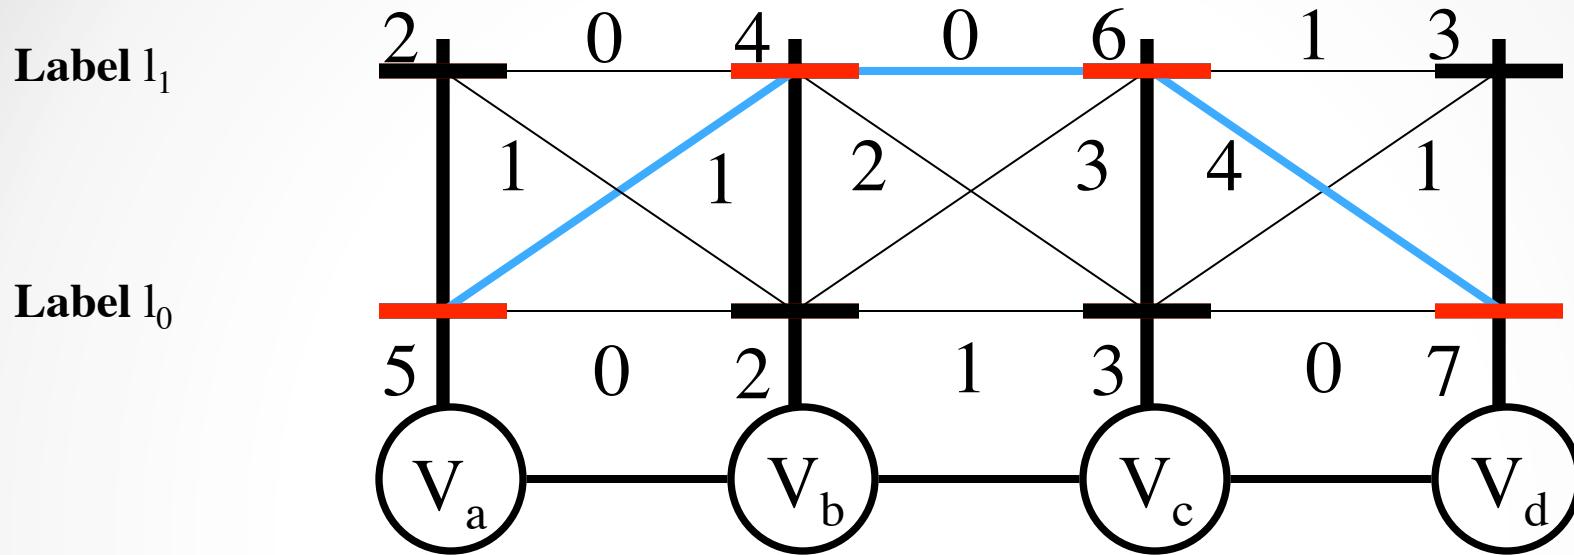
$$2 + 1 + 2 + 1 + 3 + 1 + 3 = 13$$

MAP Estimation



$$Q(f; \theta) = \sum_a \theta_{a;f(a)} + \sum_{(a,b)} \theta_{ab;f(a)f(b)}$$

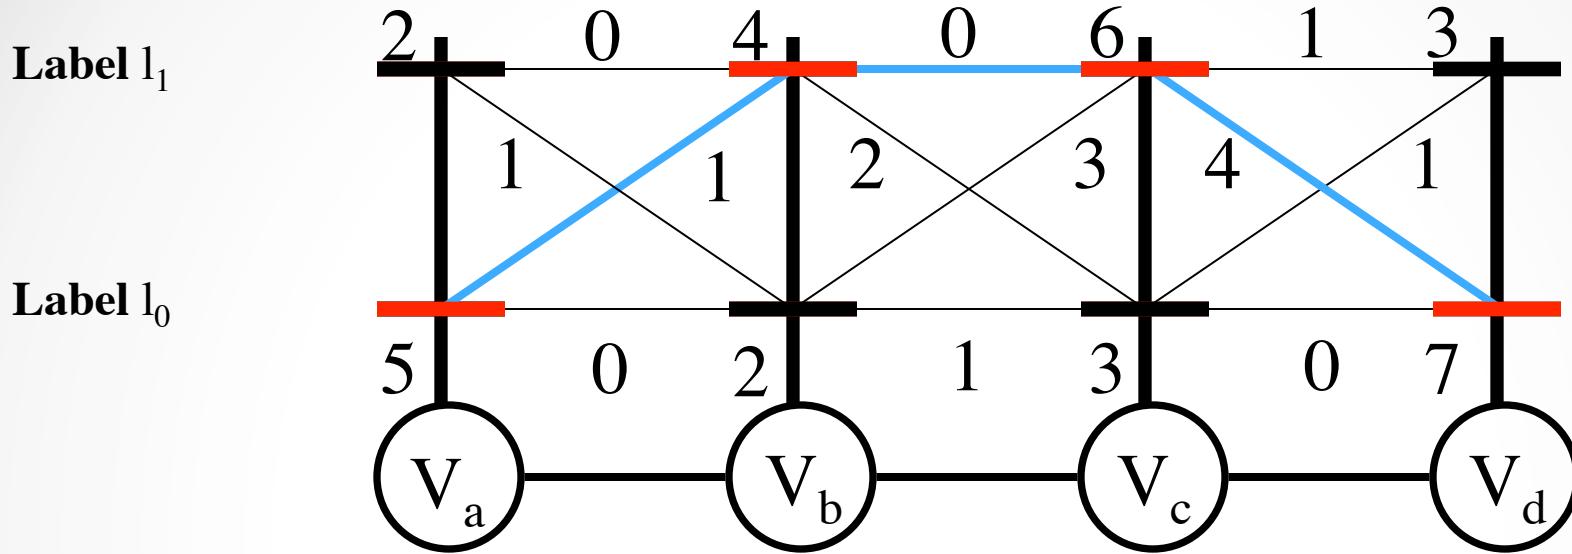
MAP Estimation



$$Q(f; \theta) = \sum_a \theta_{a;f(a)} + \sum_{(a,b)} \theta_{ab;f(a)f(b)}$$

$$5 + 1 + 4 + 0 + 6 + 4 + 7 = 27$$

MAP Estimation



$$q^* = \min Q(f; \theta) = Q(f^*; \theta)$$

$$Q(f; \theta) = \sum_a \theta_{a; f(a)} + \sum_{(a,b)} \theta_{ab; f(a)f(b)}$$

$$f^* = \arg \min Q(f; \theta)$$

MAP Estimation

16 possible labellings

$$f^* = \{1, 0, 0, 1\}$$

$$q^* = 13$$

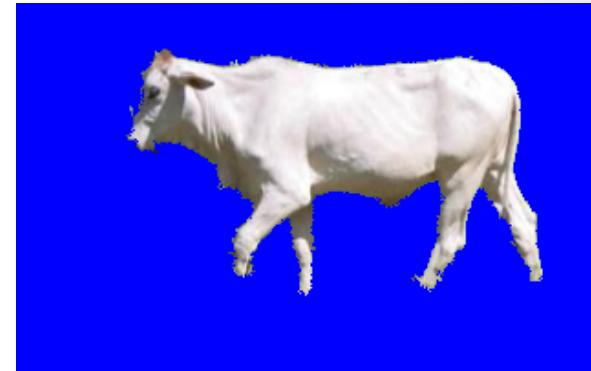
$f(a)$	$f(b)$	$f(c)$	$f(d)$	$Q(f; \theta)$
0	0	0	0	18
0	0	0	1	15
0	0	1	0	27
0	0	1	1	20
0	1	0	0	22
0	1	0	1	19
0	1	1	0	27
0.	1	1	1	20

$f(a)$	$f(b)$	$f(c)$	$f(d)$	$Q(f; \theta)$
1	0	0	0	16
1	0	0	1	13
1	0	1	0	25
1	0	1	1	18
1	1	0	0	18
1	1	0	1	15
1	1	1	0	23
1	1	1	1	16

Computational Complexity

Segmentation

$$2^{|V|}$$



$$|V| = \text{number of pixels} \approx 320 * 480 = 153600$$

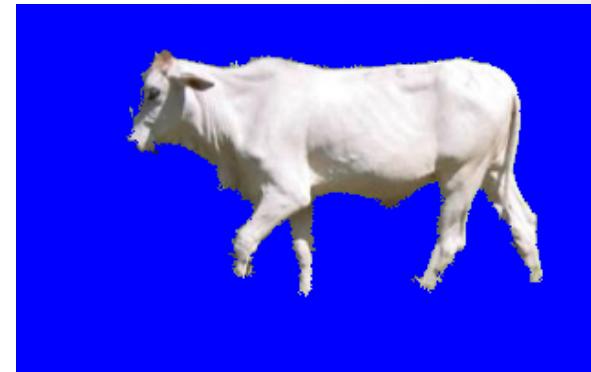
Can we do better than brute-force?

MAP Estimation is NP-hard !!

Computational Complexity

Segmentation

$$2^{|V|}$$

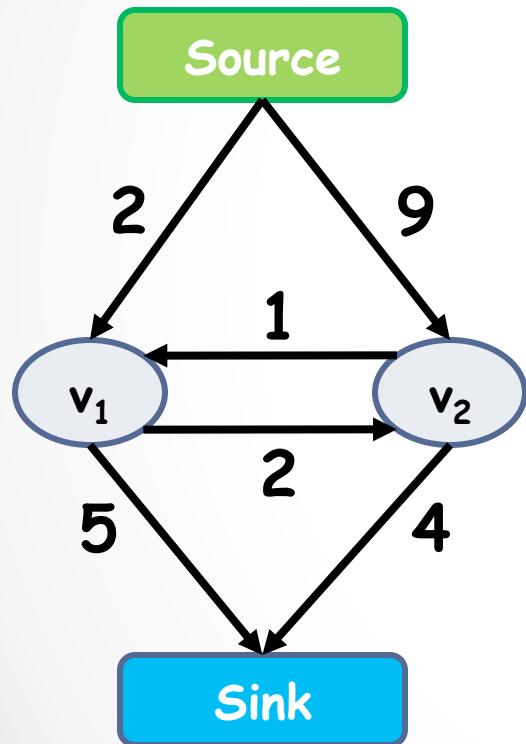


$$|V| = \text{number of pixels} \approx 320 * 480 = 153600$$

Exact algorithms do exist for special cases

Good approximate algorithms for general case

The st-Mincut Problem



Graph (V, E, C)

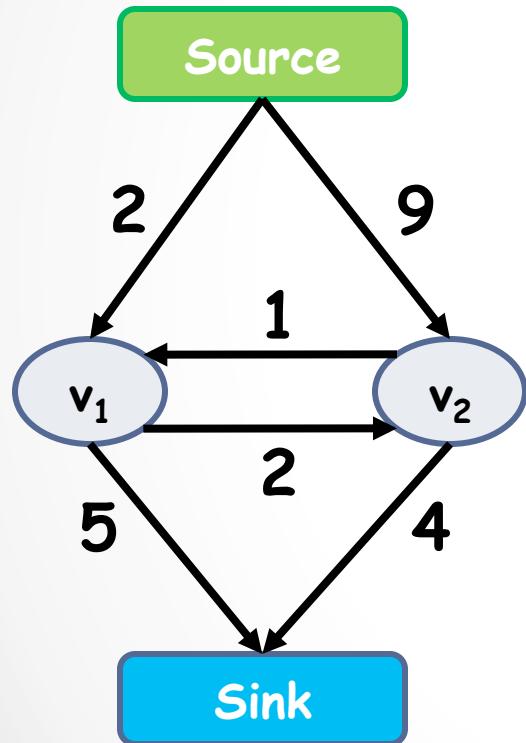
Vertices $V = \{v_1, v_2 \dots v_n\}$

Edges $E = \{(v_1, v_2) \dots\}$

Costs $C = \{c_{(1, 2)} \dots\}$

The Solution: st-Mincut

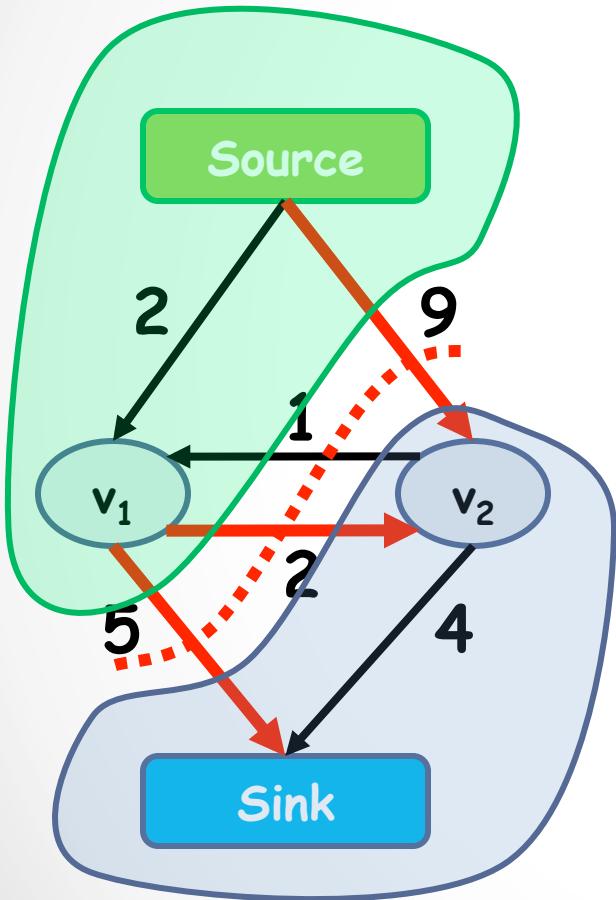
What is an st-cut?



The st-Mincut Problem

What is an st-cut?

An st-cut (S, T) divides the nodes between source and sink.



What is the cost of an st-cut?

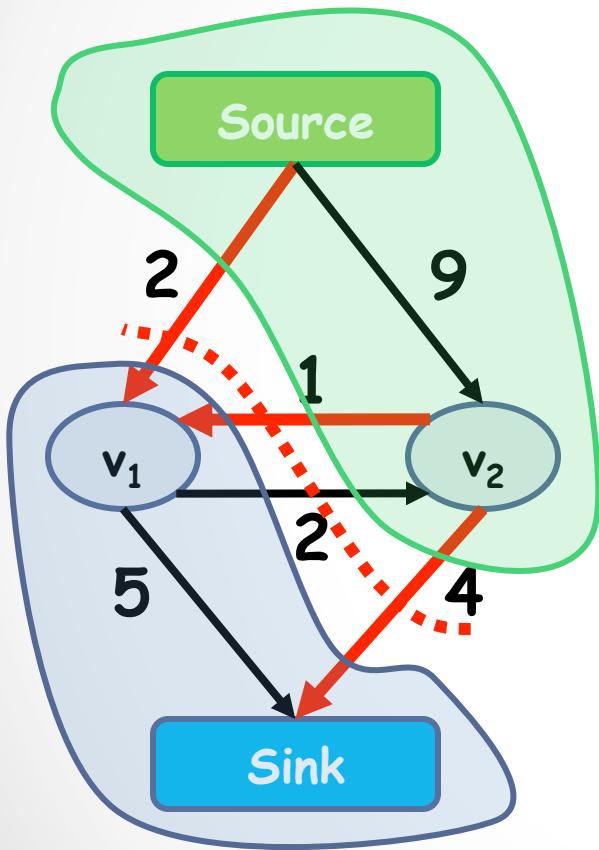
Sum of cost of all edges going from S to T

$$5 + 2 + 9 = 16$$

The st-Mincut Problem

What is an st-cut?

An st-cut (S, T) divides the nodes between source and sink.



What is the cost of an st-cut?

Sum of cost of all edges going from S to T

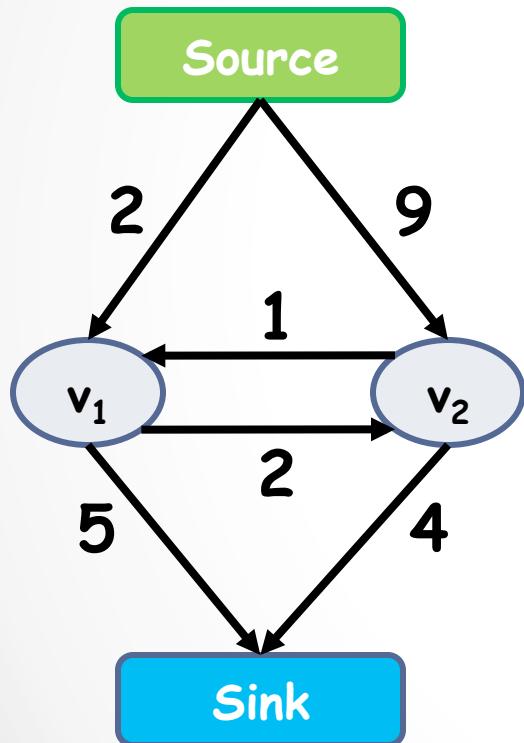
What is the st-mincut?

st-cut with the minimum cost

$$2 + 1 + 4 = 7$$

How to Compute the st-Mincut?

Solve the dual maximum flow problem
Compute the maximum flow between
Source and Sink



Constraints

Edges: Flow < Capacity

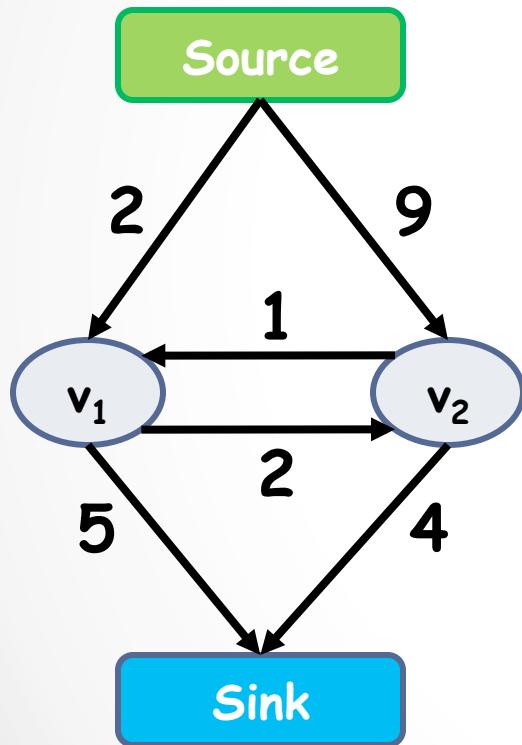
Nodes: Flow in = Flow out

Min-cut/Maxflow Theorem

In every network, the maximum flow equals the cost of the st-mincut

Maxflow Algorithms

Flow = 0



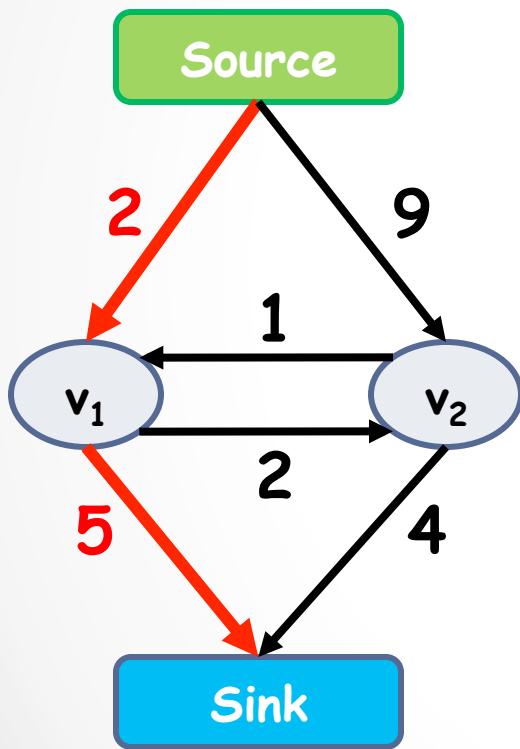
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 0



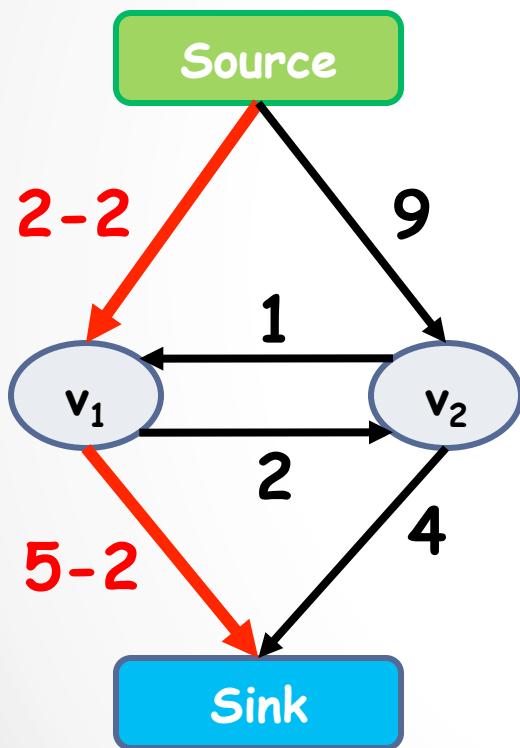
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 0 + 2



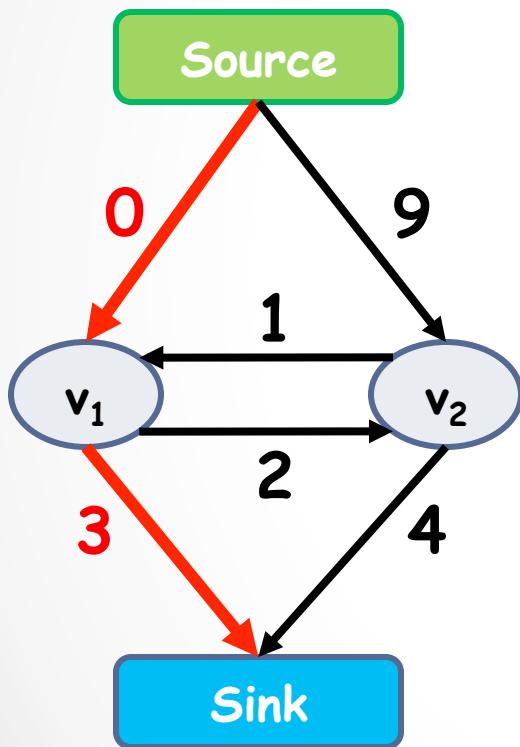
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 2



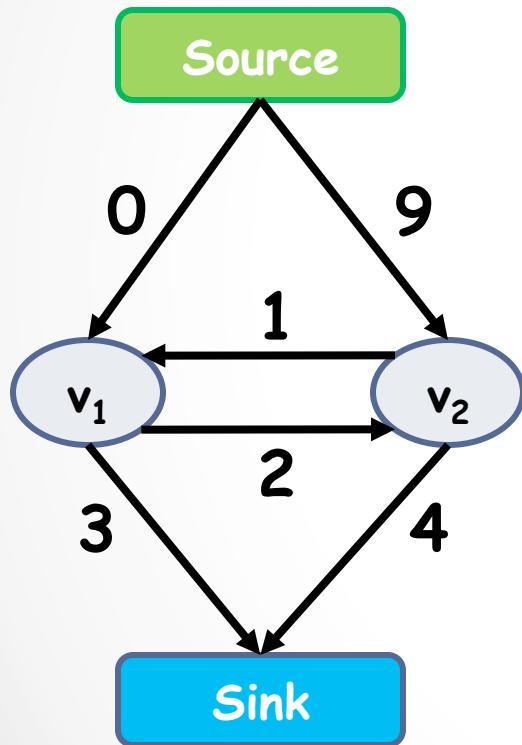
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 2



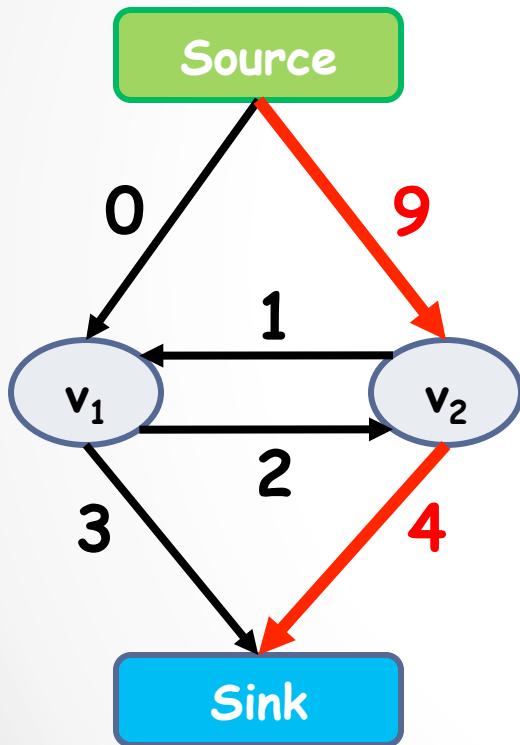
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 2



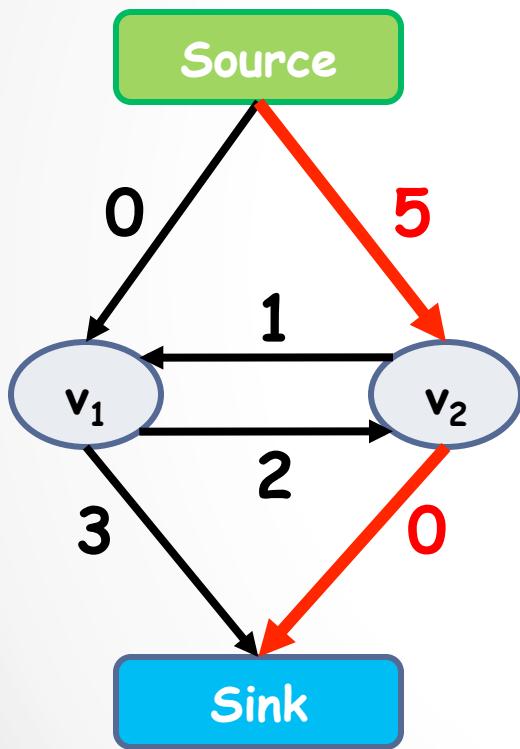
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 2 + 4



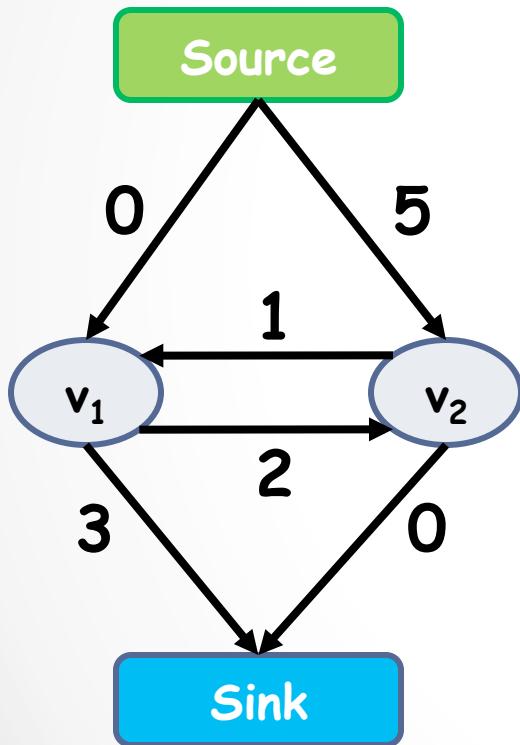
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 6



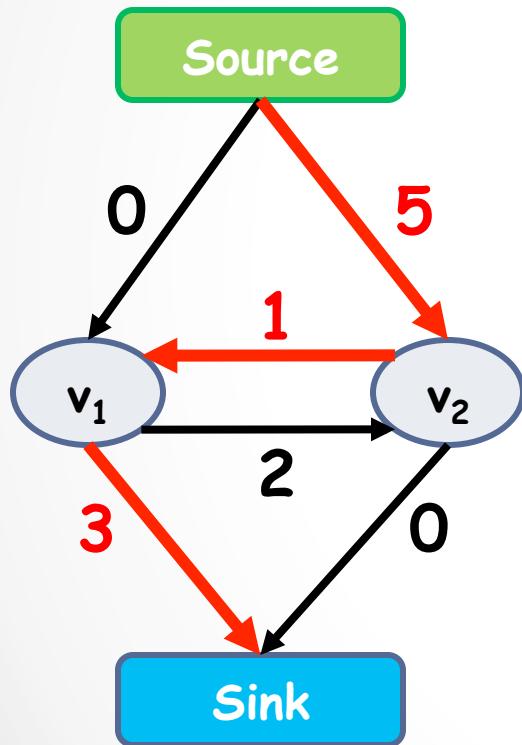
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 6



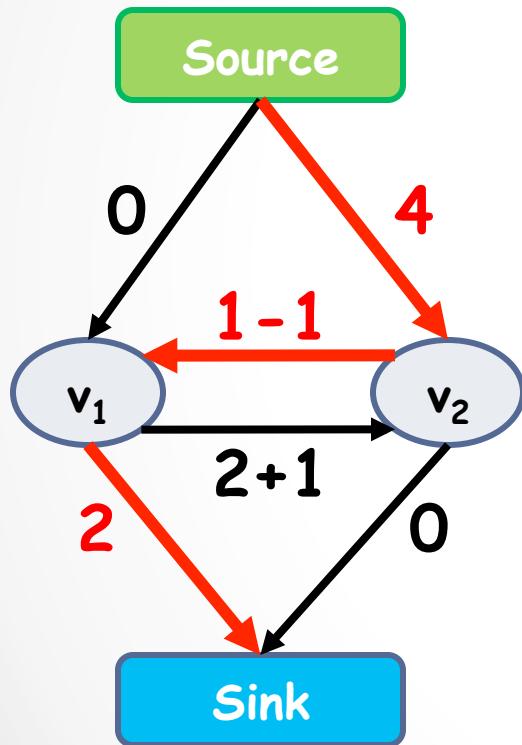
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 6 + 1



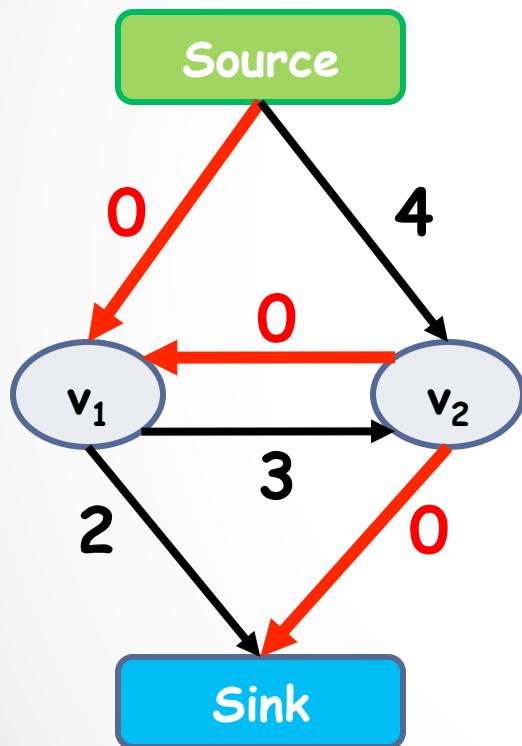
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 7



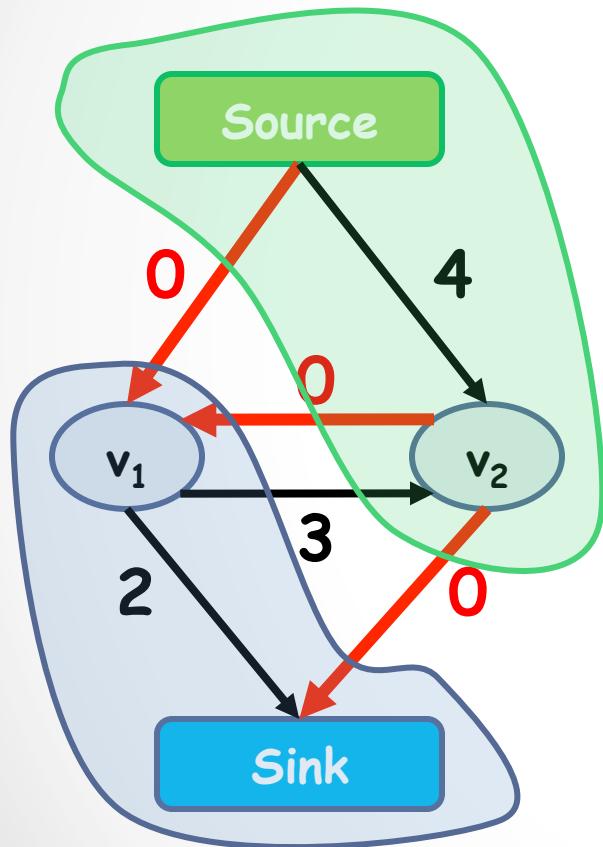
Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 7



Augmenting Path Based Algorithms

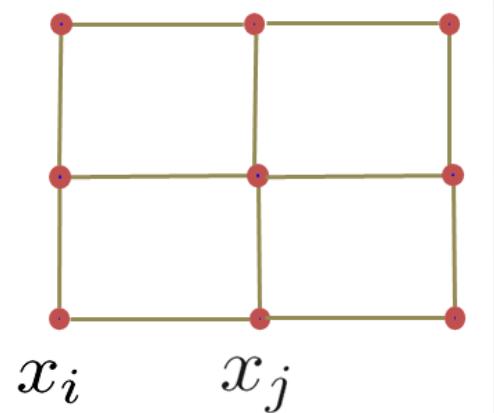
1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow in Computer Vision

- Specialized algorithms for vision problems

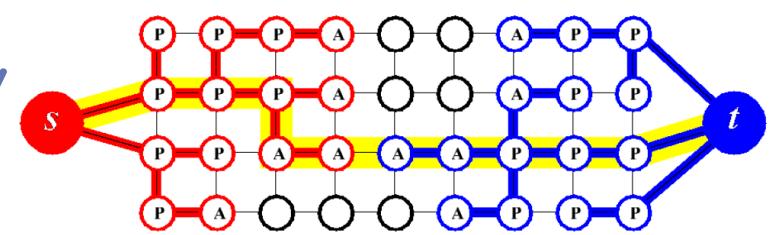
- Grid graphs
- Low connectivity ($m \sim O(n)$)



- Dual search tree augmenting path algorithm

[Boykov and Kolmogorov PAMI 2004]

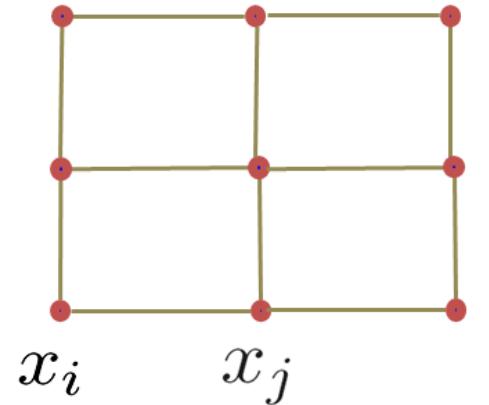
- Finds approximate shortest augmenting paths efficiently
- High worst-case time complexity
- Empirically outperforms other algorithms on vision problems



Maxflow in Computer Vision

- Specialized algorithms for vision problems

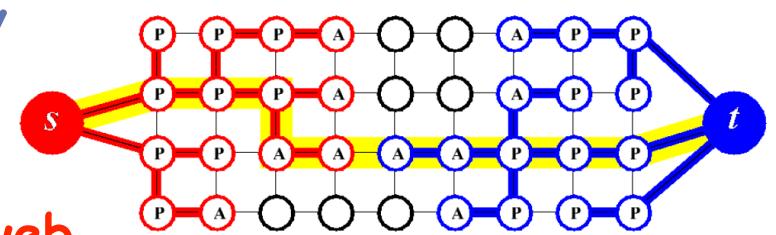
- Grid graphs
- Low connectivity ($m \sim O(n)$)



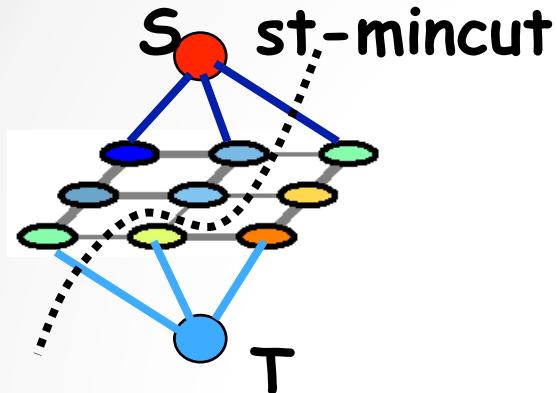
- Dual search tree augmenting path algorithm

[Boykov and Kolmogorov PAMI 2004]

- Finds approximate shortest augmenting paths efficiently
 - High worst-case time complexity
 - Empirically outperforms other algorithms on vision problems
 - **Efficient code available on the web**
- <http://pub.ist.ac.at/~vnk/software.html>



St-mincut and Energy Minimization



Minimizing a Quadratic Pseudoboolean function $E(x)$

Functions of boolean variables

$$E: \{0, 1\}^n \rightarrow \mathbb{R}$$

$$E(x) = \sum_i c_i x_i + \sum_{i,j} c_{ij} x_i (1 - x_j)$$

$$c_{ij} \geq 0$$

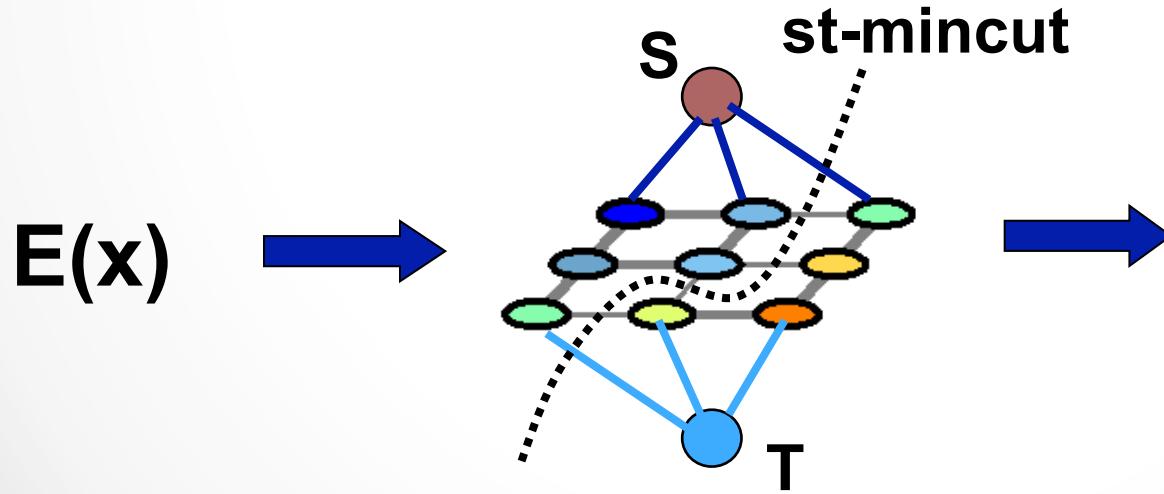


Polynomial time st-mincut algorithms require non-negative edge weights

How does this Work?

Construct a graph such that:

1. Any st-cut corresponds to an assignment of x
2. The cost of the cut is equal to the energy of $x : E(x)$



Solution

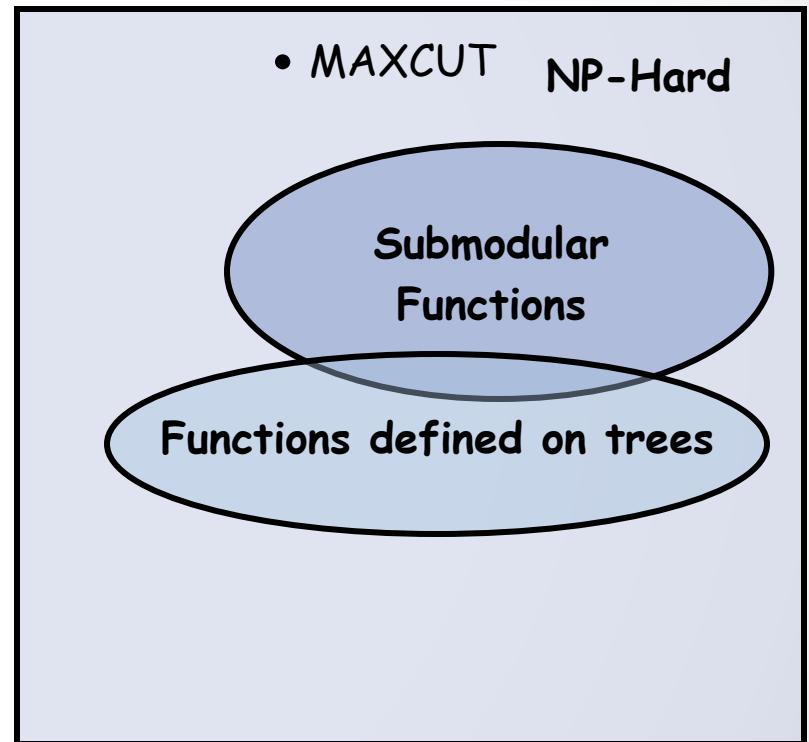
Minimizing Energy Functions

- **General Energy Functions**

- NP-hard to minimize
- Only approximate minimization possible

- **Easy energy functions**

- Solvable in polynomial time
- Submodular $\sim O(n^6)$



Space of Function
Minimization Problems

GrabCut

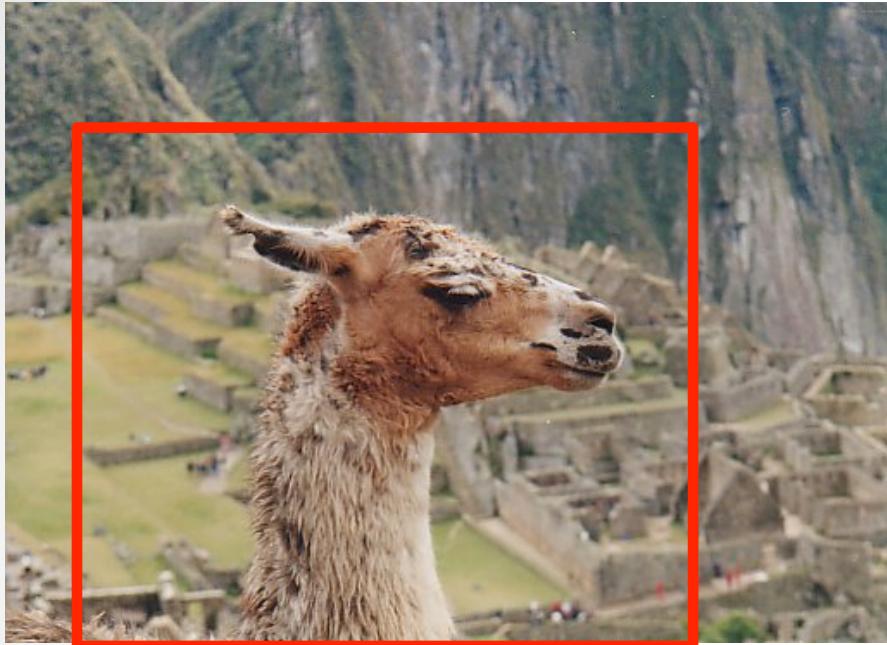
Interactive Foreground Extraction using Iterated Graph Cuts

Carsten Rother

Vladimir Kolmogorov

Andrew Blake

The Problem

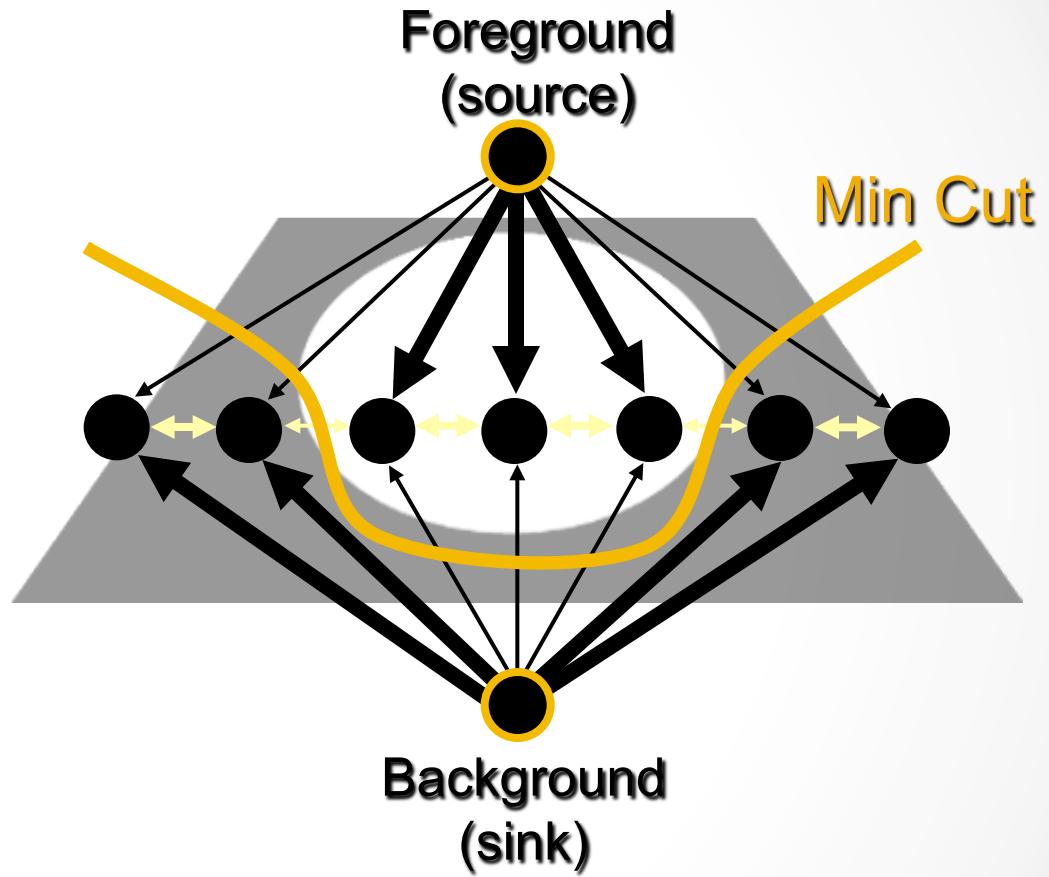
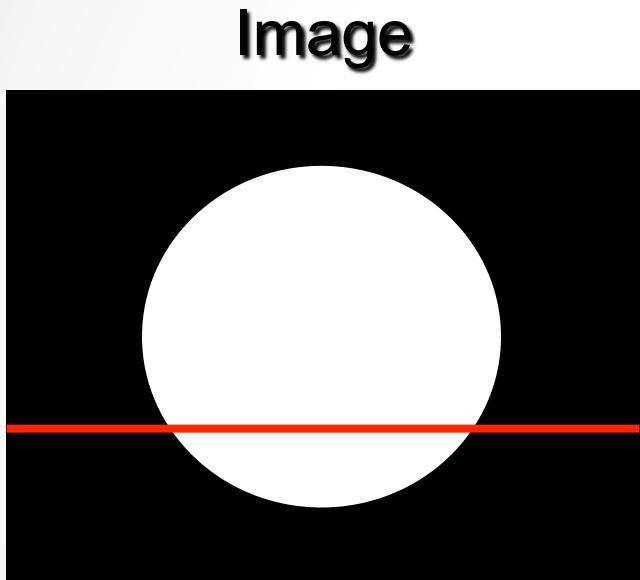


Fast &
Accurate ?



- Less user input: only rectangle
- Handle color
- Extract matte as post-process

Use Basic Graph Cut for Segmentation



Cut: separating source and sink; Energy: collection of edges

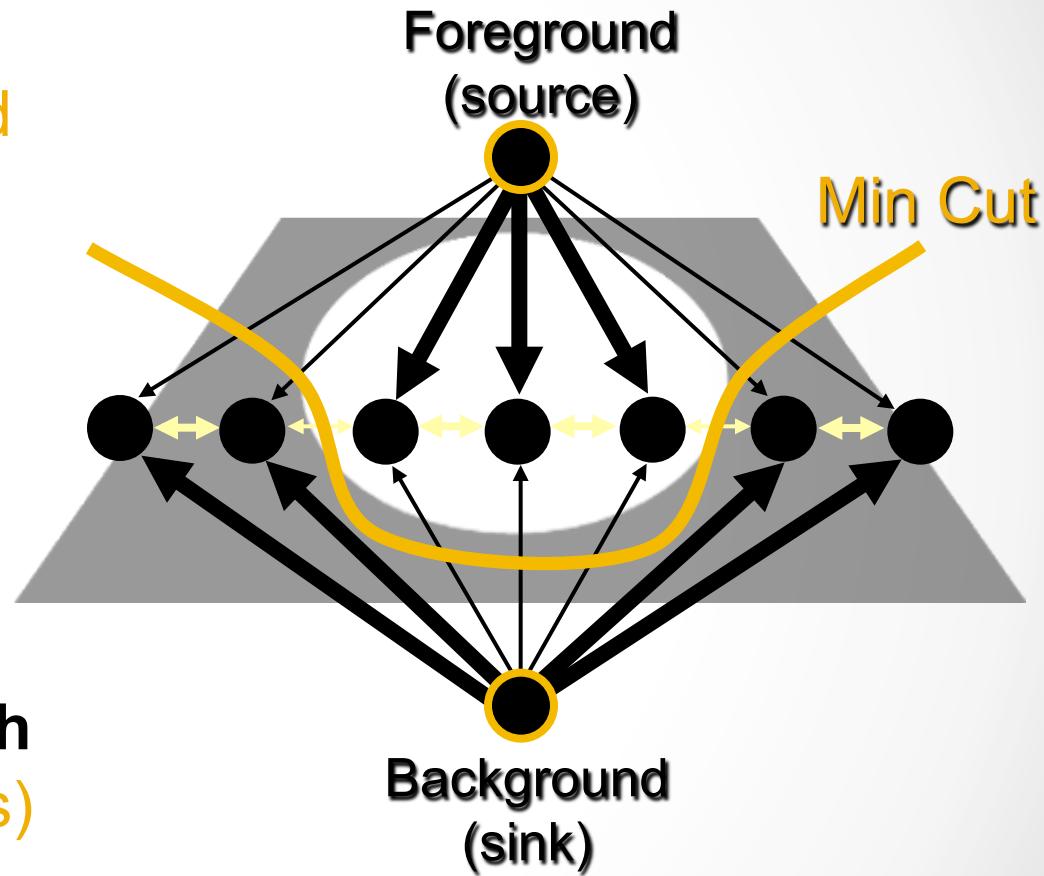
Min Cut: Global minimal energy in polynomial time

Graph Cuts for Foreground Extraction

Assume we know foreground is **white** and background is **black**

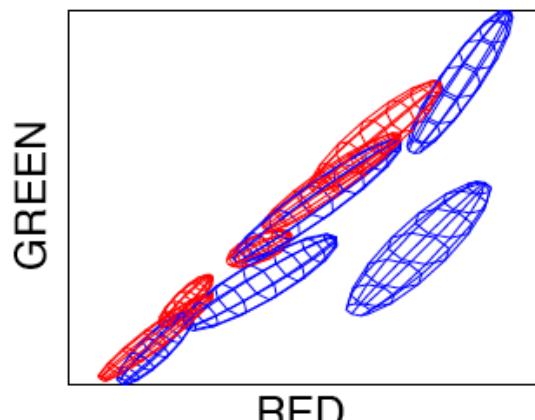
Data term = **whiteness**
(cost of assigning label)

Regularization = **color match**
(cost of separating neighbors)

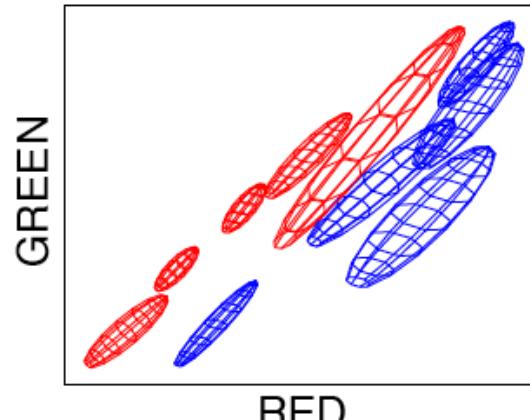


Color Data Term

- Model 3D color histogram with Gaussians
 - Because brute force histogram would be sparse
 - Gaussian Mixture Model (GMM)
 - Just means histogram = sum of Gaussians
 - They advise 5 Gaussians



(b)



(c)

Iterated Graph Cuts



User Initialisation



Learn foreground
color model



Graph cuts to
infer the
foreground

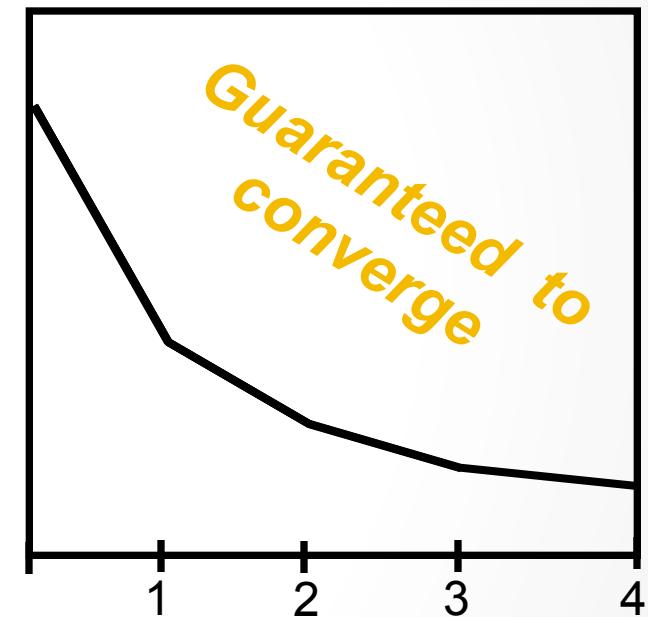
Grabcut: Iterative approach

- Initialize
 - Background with rectangle boundary pixels
 - Foreground with the interior of rectangle
- Iterate until convergence
 - Compute color probabilities (GMM) of each region
 - Perform graphcut segmentation
- Apply matting at boundary
- Potentially, user edits to correct mistakes

Iterated Graph Cuts

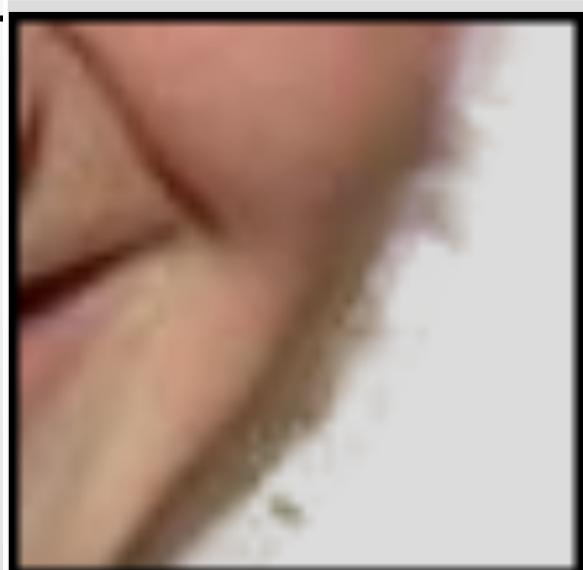
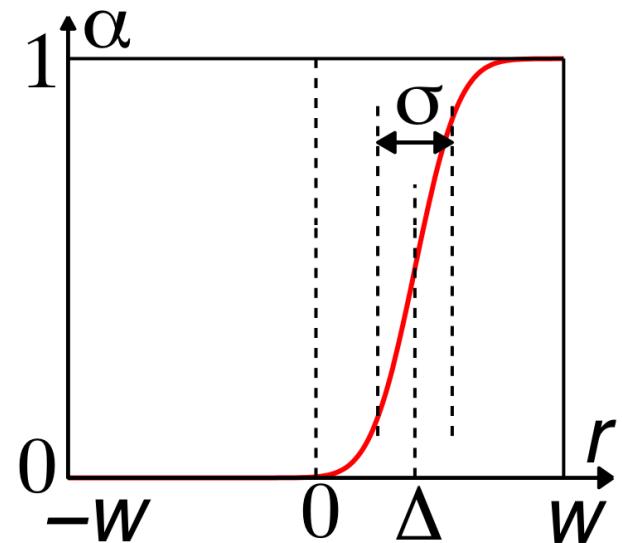
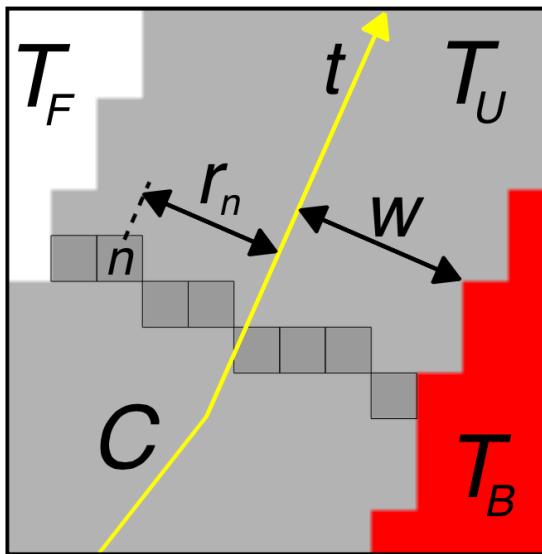
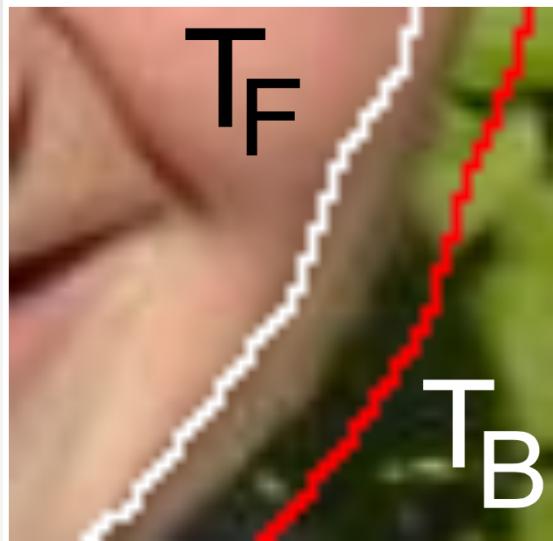


Result

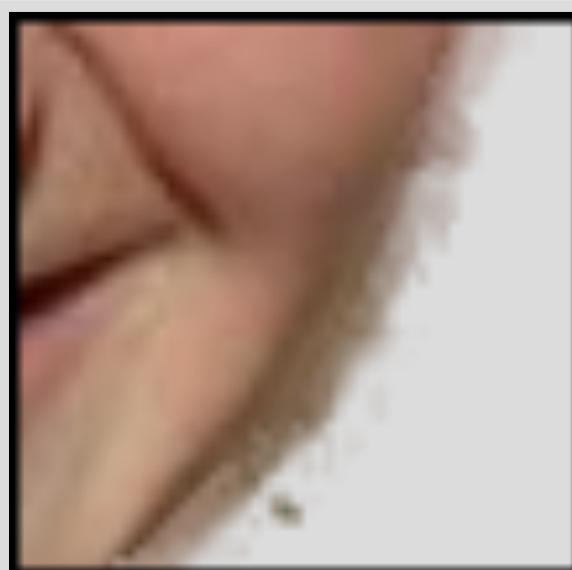


Energy after each Iteration

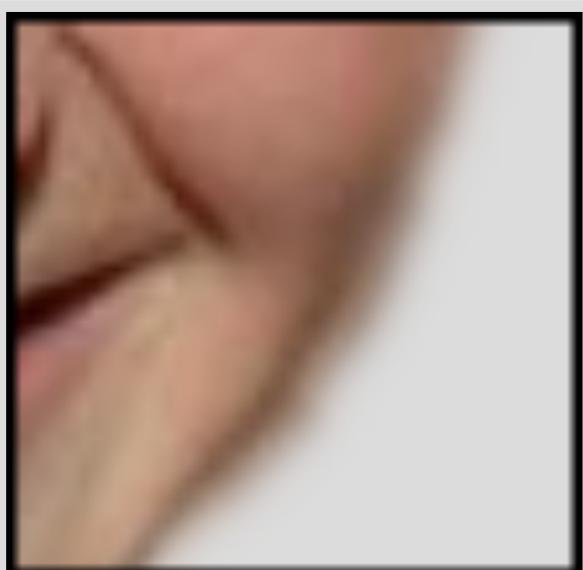
Border Matting



Knockout 2

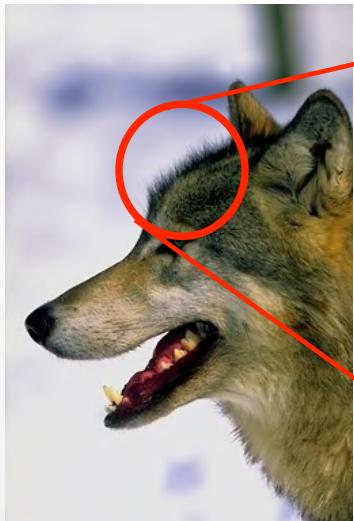
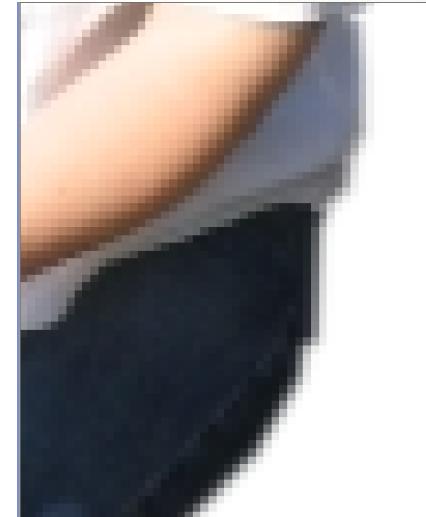


Bayes Matte



GrabCut

Border Matting Results



Examples: Moderate Complexity



... GrabCut completes automatically

User Interaction



Automatic



Segmentation



User



Interaction



Automatic



Segmentation



Examples: Difficult Ones

Camouflage &
Low Contrast

Initial
Rectangle



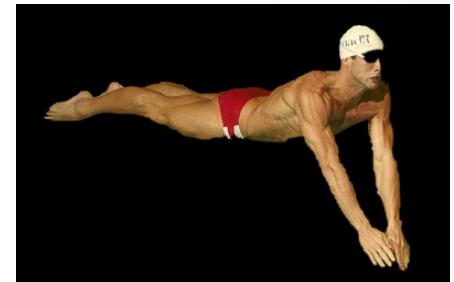
Initial
Result



Fine structure



No telepathy



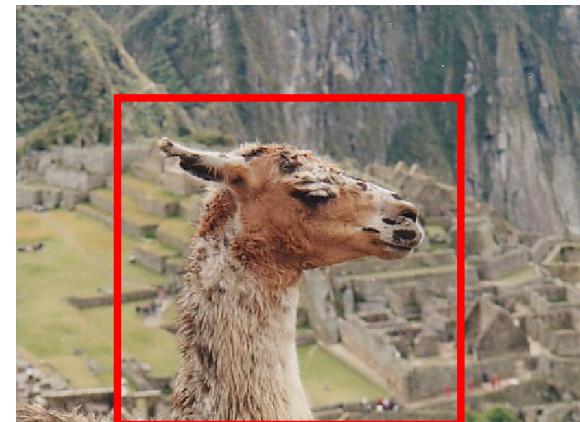
GraphCut vs. GrabCut

Boykov and Jolly (2001)

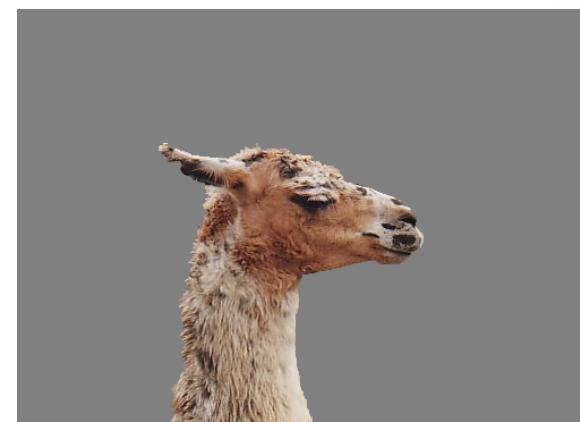
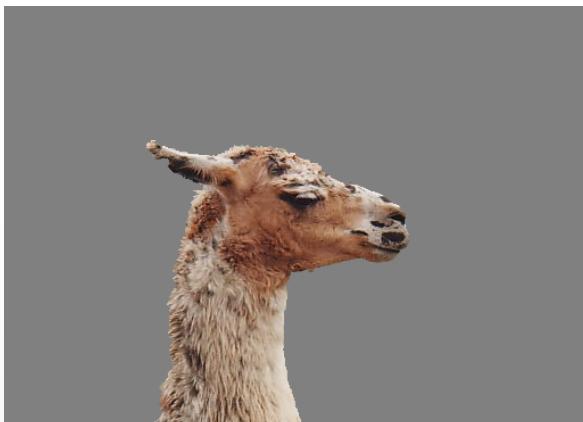
User
Input



GrabCut



Result



Error Rate: 0.72%

Error Rate: 0.72%

Reading Material

- Required
 - “Normalized Cuts and Image Segmentation”, by Jianbo Shi and Jitendra Malik, TPAMI, Aug 2000
 - “Grabcut: Interactive Foreground Extraction using Iterated Graph Cuts”, by Rother, Kolmogorov and Blake, Siggraph 2004.
- Additional
 - “OBJ CUT”, by M. Pawan Kumar P.H.S. Torr and A. Zisserman, CVPR 2005.

Usage of MRF: An Example

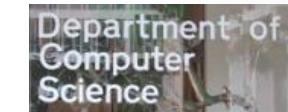
Binarizing Natural Scene Text

Anand Mishra, Karteeck Alahari and C. V. Jawahar

“An MRF Model for Binarization of Natural Scene Text”,

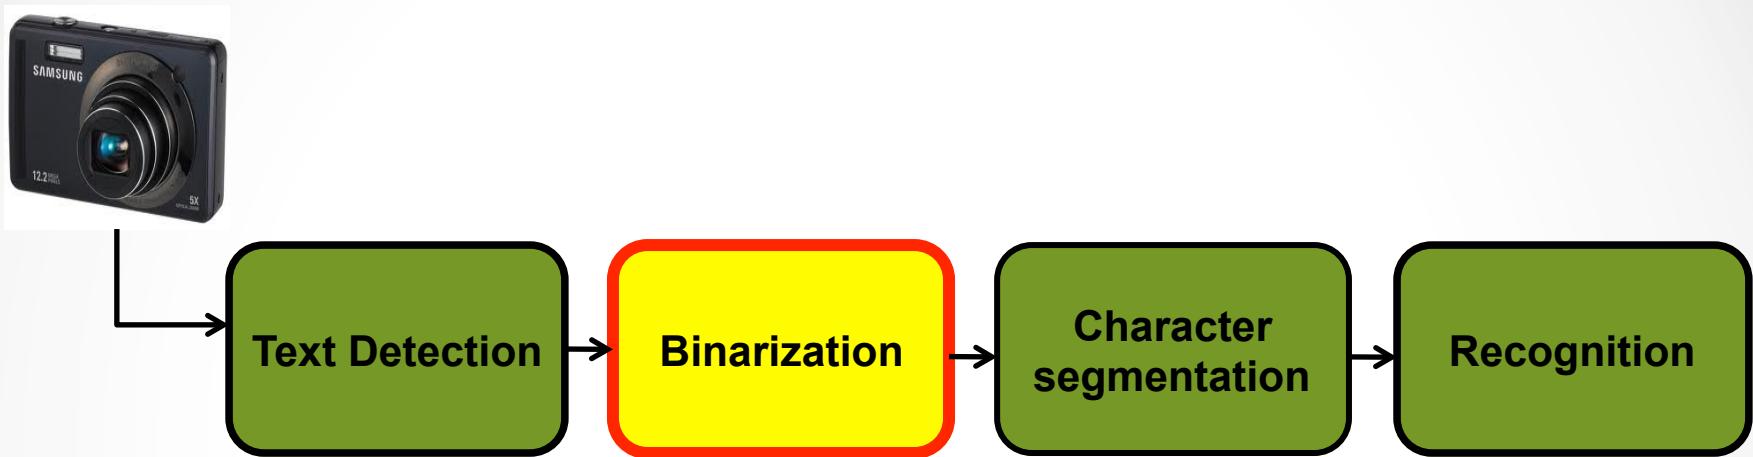
Proc. ICDAR, Sept. 2011, Beijing, China

Natural Scene Text



- Challenging ICDAR 2003 and Street View Text (SVT) datasets

Natural Scene Text Recognition

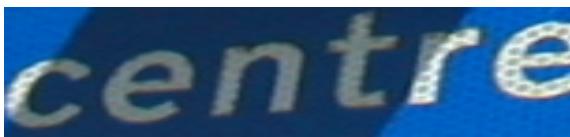


LITTER



REDBACK

Scene Text Binarization: Challenges



- Reflections/
Transparency
- Illumination/Shadows
- Specularity
- Low Contrast
- Complex BG
- Noise
- Blur

Failure of Existing Methods

Original



Otsu



Kittler



Niblack



Sauvola



- ICDAR 2003 Robust Word Recognition Competition
 - A.Mishra et al. (ICDAR 2011): State-of-the-art
 - Y. Zhou et al. (ICDAR 2013): +2%, spotting using a small lexicon

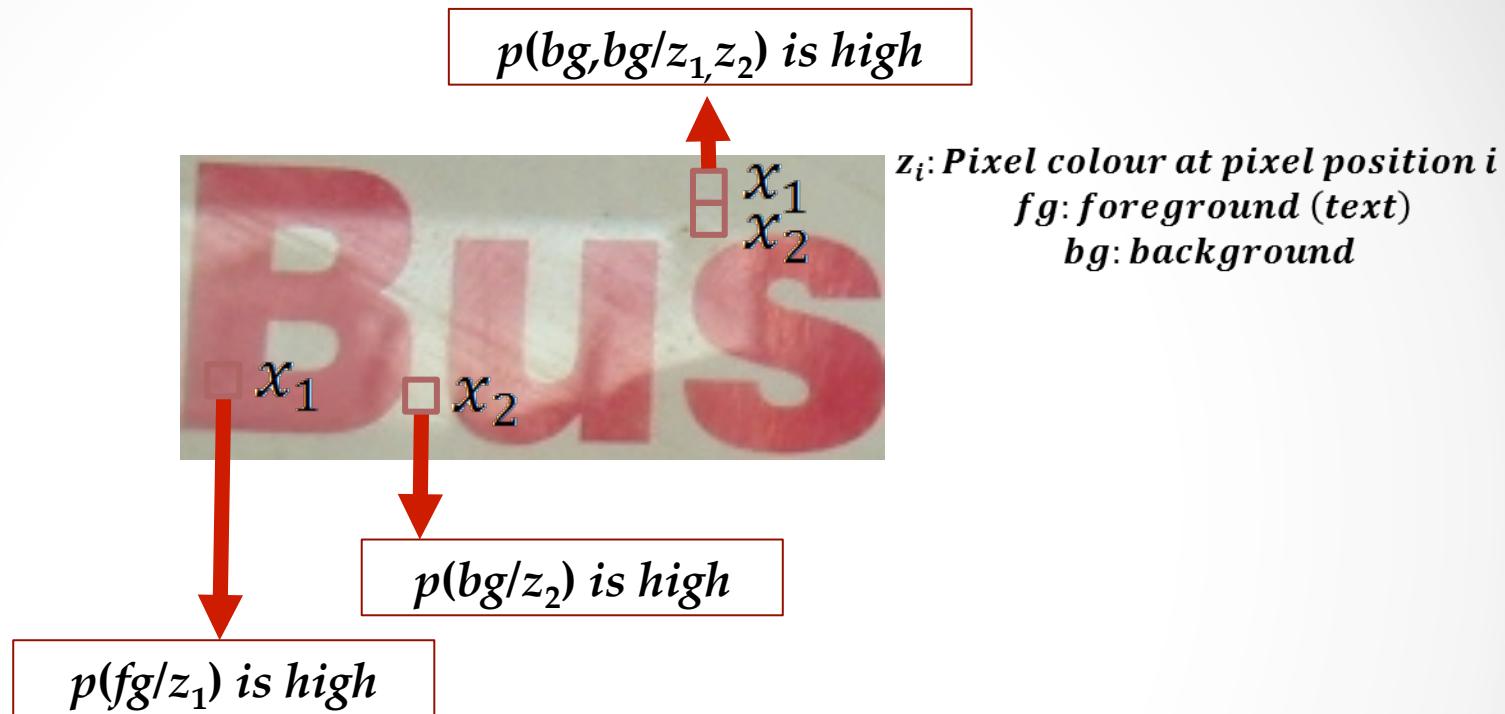
An MRF based Binarization



Assign a label to each pixel from $L = \{\text{Text} (0), \text{Background}(1)\}$

Many labelings possible, we are interested in
“the optimal” one

MRF Formulation: Potentials



Minimize

$$E(x) = -\sum_i \log p(x_i|z_i) + \lambda_1 \sum_{i,j \in N} \exp(-\beta \|z_i - z_j\|^2)$$

Unary (data) Term

Pair wise (smoothness) Term

MRF Formulation: Gradient Potential



Gradient magnitude at pixel position i

$$\text{Pair wise term} = \lambda_1 \sum_{i,j \in N} \exp(-\beta \|z_i - z_j\|^2) + \lambda_2 \sum_{i,j \in N} \exp(-\beta \|w_i - w_j\|^2)$$

Edginess Term

A vertical arrow points from the text "Gradient magnitude at pixel position i" down towards the second term in the equation, specifically pointing to the variable w_i .

An MRF based Binarization

The problem is to minimize following energy (MRF energy):

$$E(x) = \textit{Unary term} + \textit{Pairwise term}$$

Two questions:

- 1) How to learn the probabilities $p(x_i|z_i)$ used to compute the unary term?**
- 2) How to find the minima of above energy?**

Learning Probabilities



Canny Edge operator



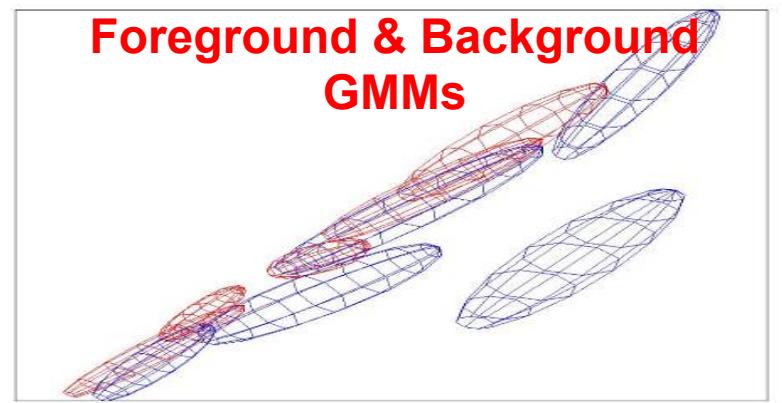
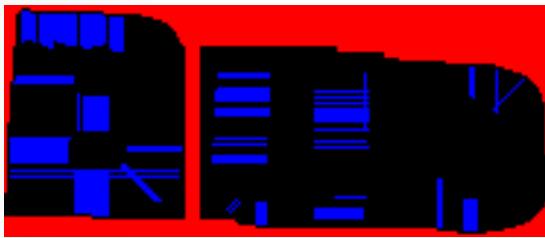
Find foreground -
background seeds



Blue colour:
Foreground

Red colour:
Background

Color Modeling through GMMs



Unary term is calculated based on the probability of a pixel colour belonging to one of the GMM components

An MRF based Binarization

The problem is to minimize following energy (MRF energy):

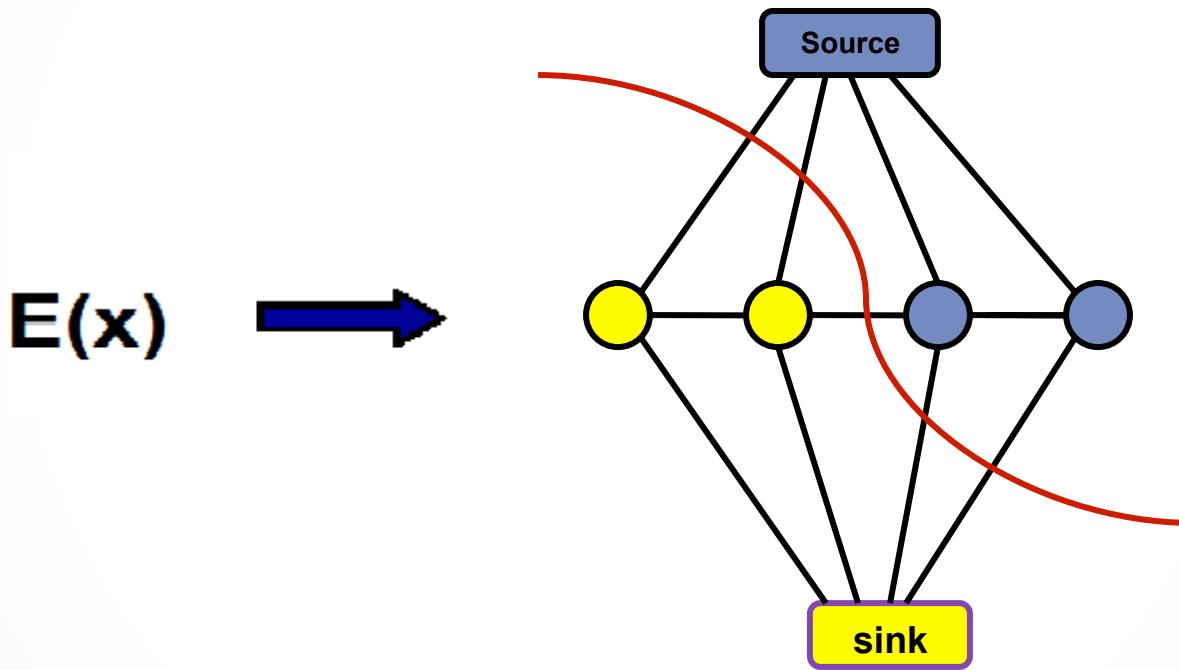
$$E(x) = \textit{Unary term} + \textit{Pair wise term}$$

Two questions:

- 1) How to learn the probabilities $p(x_i|z_i)$ used to compute the unary term?
- 2) How to find the minima of above energy?

Graph Cut

Minimum of MRF energy = min cut of graph

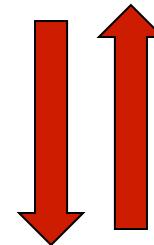


Efficient codes available to compute min cut of such graph

An Iterative Graph Cut based Approach



Learn GMMs to model foreground and background colours



Graph cuts to refine binarization

Qualitative Results

Bus

Life

Howard

Memorex

Bus

Life

Howard

Memorex

Qualitative Results

1600

22

BOROUGH

CD-R

1600

22

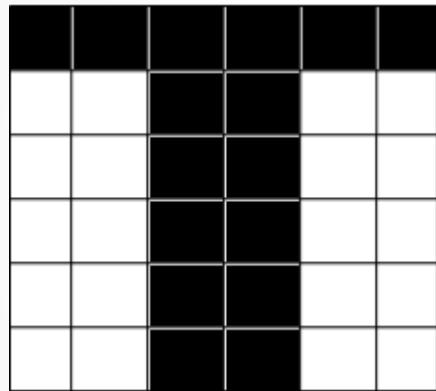
BOROUGH

CD-R

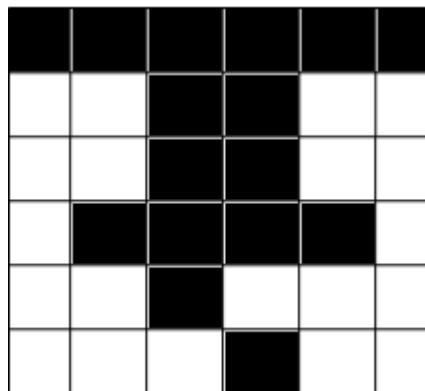
Quantitative Results

- OCR accuracy
- Pixel level accuracy

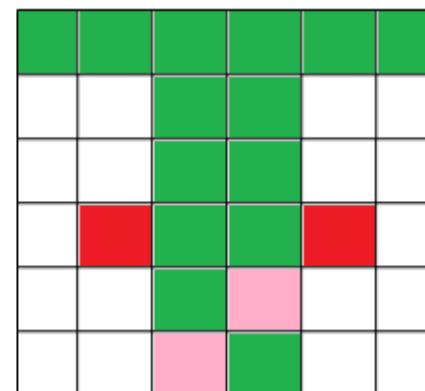
ABBY



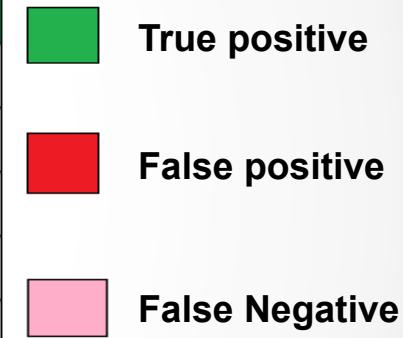
Ground truth



Binarization result



Evaluation result



$$precision = \frac{\text{Total number of green boxes}}{\text{Total number of green boxes} + \text{Total number of red boxes}}$$

$$recall = \frac{\text{Total number of green boxes}}{\text{Total number of green boxes} + \text{Total number of pink boxes}}$$

$$f-score = \frac{2 \times precision \times recall}{precision + recall} \times 100$$

Results (ABBYY OCR Accuracy)

Method	Word Accuracy (%)	Character Accuracy (%)
Otsu	41.52	51.74
Sauvola	39.77	51.63
Niblack	39.18	42.31
Kittler	41.12	49.88
Otsu + CT	45.03	51.98
MRF (without edginess diff.)	49.12	55.94
MRF (with edginess diff.)	52.04	60.14

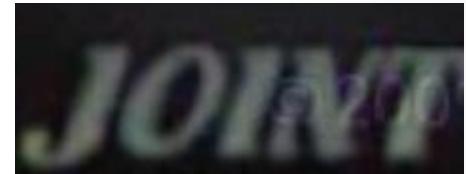
Results (Pixel Level Accuracy)

Method	f-score (%)
Otsu	79.32
Sauvola	73.87
Niblack	76.86
Kittler	72.89
Otsu + CT	78.12
MRF (without edginess diff.)	87.84
MRF (with edginess diff.)	88.64

More Results

Results based on Street View Text Dataset

Method	Word Recognition accuracy (%)
ABBYY	32.61%
MRF Binarization + ABBYY	42.81%



Kai Wang and Serge Belongie (ECCV 2010) have introduced a challenging Street View Text (SVT) dataset

When can it Fail?

- Colour may not characterize the character.
- Failing to learn text-BG probabilities



Reading Material

- Required
 - “Normalized Cuts and Image Segmentation”, by Jianbo Shi and Jitendra Malik, TPAMI, Aug 2000
 - “Grabcut: Interactive Foreground Extraction using Iterated Graph Cuts”, by Rother, Kolmogorov and Blake, Siggraph 2004.
 - “An MRF Model for Binarization of Natural Scene Text”, by Anand Mishra, Karteek Alahari and C.V. Jawahar, ICDAR 2011.
- Additional
 - “OBJ CUT”, by M. Pawan Kumar P.H.S. Torr and A. Zisserman, CVPR 2005.